




Memory Segmentation and Free-Space Management

Inefficiency of Base and Bound Approach


- Big chunk of “free” space
- “free” space takes up physical memory
- Hard to run when an address space does not fit into physical memory



Segmentation

- Segment is just **a contiguous portion** of the address space of a particular length
 - Logically-different segment: code, stack, heap, and etc.
- Each segment can be **placed** in **different part of physical memory**
 - **Base** and **bounds** exist **per each segment**

Placing Segment In Physical Memory

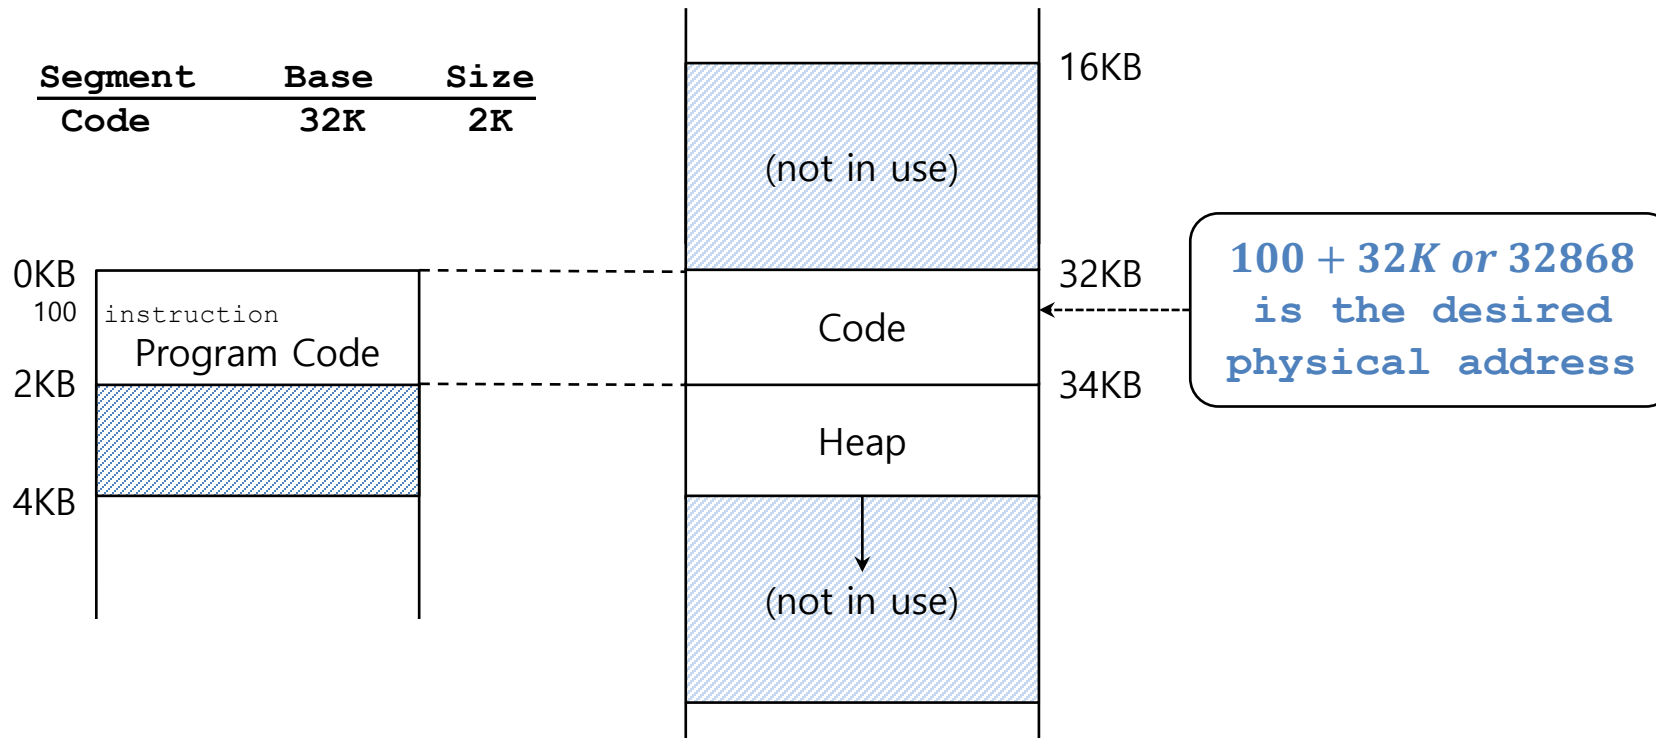


Segment	Base	Size
Code	32K	2K
Heap	34K	2K
Stack	28K	2K

Address Translation on Segmentation

$$\text{physical address} = \text{offset} + \text{base}$$

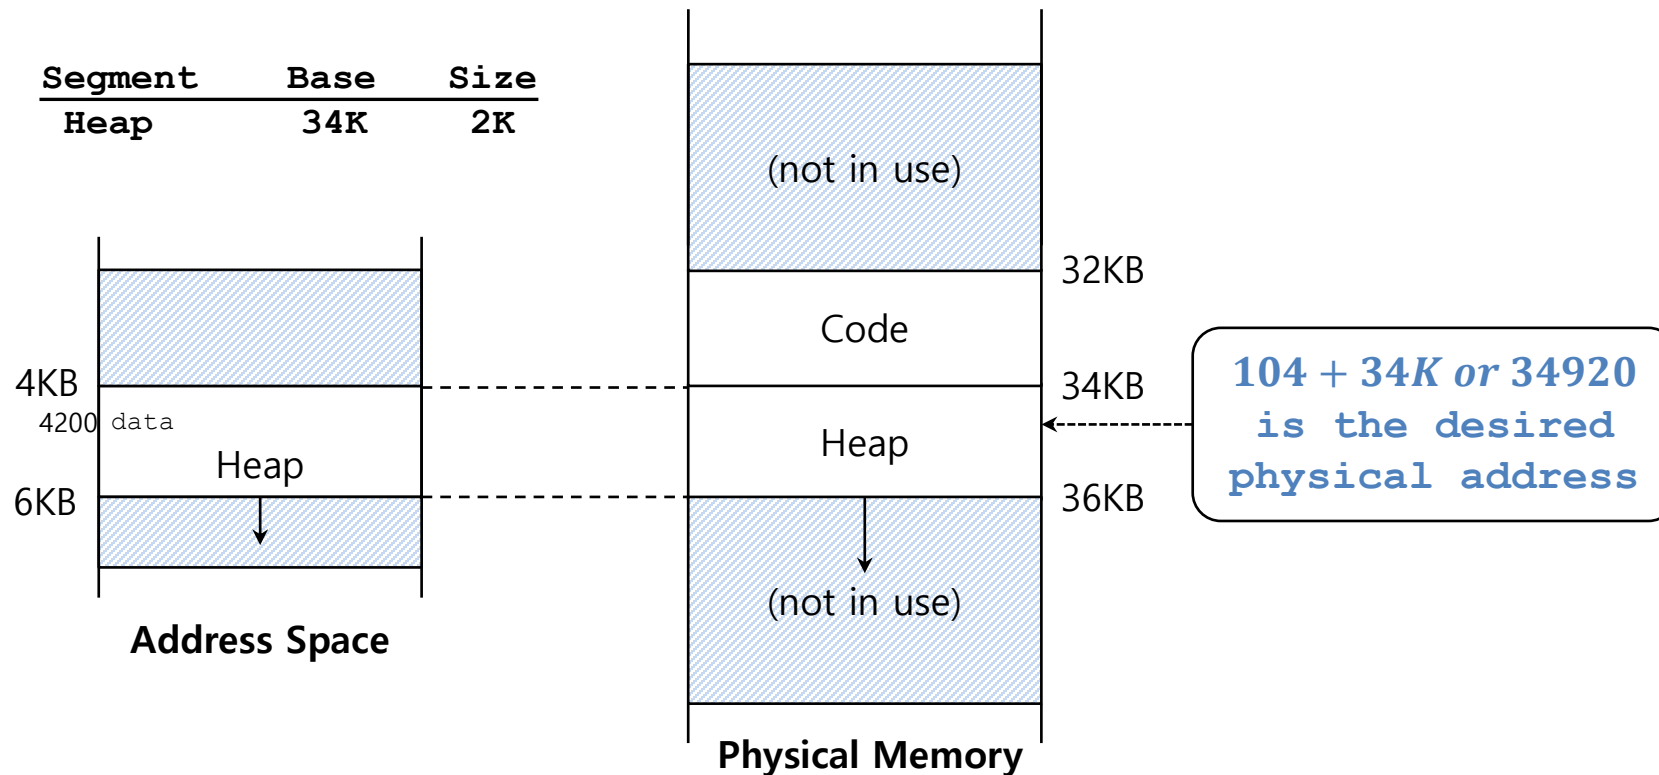
- The `offset` of virtual address 100 is 100
 - The code segment **starts at virtual address 0** in address space



Address Translation on Segmentation


Virtual address + base is not the correct physical address

- The offset of virtual address 4200 is 104
 - The heap segment starts at virtual address 4096 in address space



Segmentation Fault or Violation


- If an **illegal address** such as 7KB which is beyond the end of heap is referenced, the OS occurs **segmentation fault**
 - The hardware detects that address is **out of bounds**



Referring to Segment


- **Explicit approach**

- Chop up the address space into segments based on the **top few bits** of virtual address



- Example: virtual address 4200 (01000001101000)

Segment	bits
Code	00
Heap	01
Stack	10
-	11




Referring to Segment

```
1  // get top 2 bits of 14-bit VA
2  Segment = (VirtualAddress & SEG_MASK) >> SEG_SHIFT
3  // now get offset
4  Offset = VirtualAddress & OFFSET_MASK
5  if (Offset >= Bounds[Segment])
6      RaiseException(PROTECTION_FAULT)
7  else
8      PhysAddr = Base[Segment] + Offset
9      Register = AccessMemory(PhysAddr)
```

- SEG_MASK = 0x3000 (1100000000000000)
- SEG_SHIFT = 12
- OFFSET_MASK = 0xFFF (0011111111111111)

Referring to Stack Segment

- Stack grows **backward**
- **Extra hardware support** is need
 - The hardware checks which way the segment grows
 - 1: positive direction, 0: negative direction



Segment Register(with Negative-Growth Support)

Segment	Base	Size	Grows Positive?
Code	32K	2K	1
Heap	34K	2K	1
Stack	28K	2K	0

Support for Sharing

- Segment can be **shared between address space**
 - **Code sharing** is still in use in systems today
 - by extra hardware support
- Extra hardware support is need for form of **Protection bits**
 - **A few more bits** per segment to indicate **permissions** of **read**, **write** and **execute**

Segment Register Values(with Protection)

Segment	Base	Size	Grows	Positive?	Protection
Code	32K	2K		1	Read-Execute
Heap	34K	2K		1	Read-Write
Stack	28K	2K		0	Read-Write

Fine-Grained and Coarse-Grained

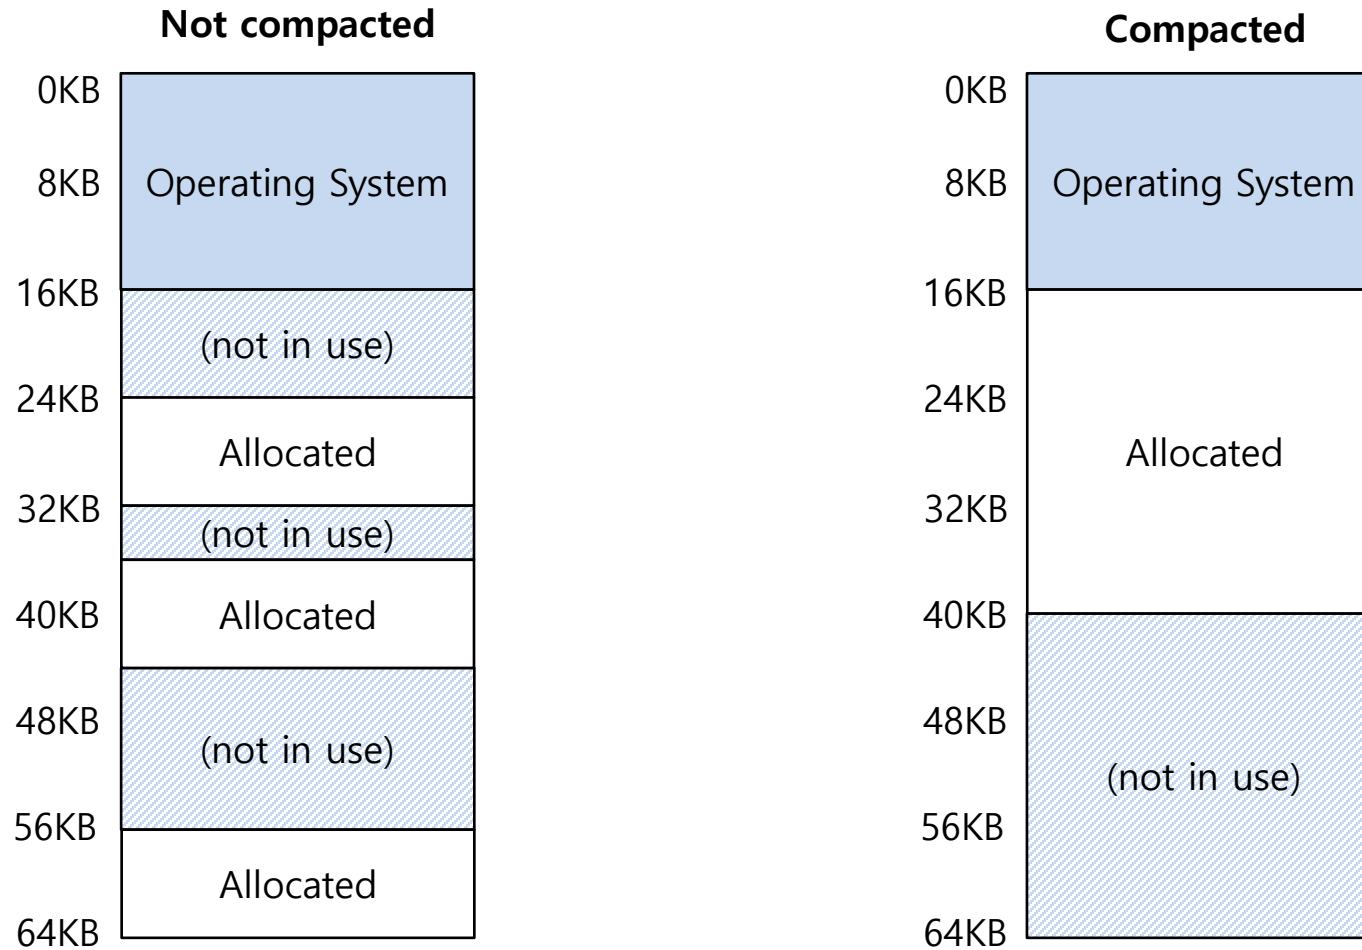


- **Coarse-Grained** means segmentation in a small number
 - e.g., code, heap, stack
- **Fine-Grained** segmentation allows **more flexibility** for address space in some early system
 - To support many segments, Hardware support with a **segment table** is required

OS support: Fragmentation


- **External Fragmentation:** little holes of **free space** in physical memory that make difficulty to allocate new segments
 - There is **24KB free**, but **not in one contiguous** segment
 - The OS **cannot** satisfy the **20KB request**
- **Compaction: rearranging** the exiting segments in physical memory
 - Compaction is **costly**
 - **Stop** running process
 - **Copy** data to somewhere
 - **Change** segment register value

Memory Compaction



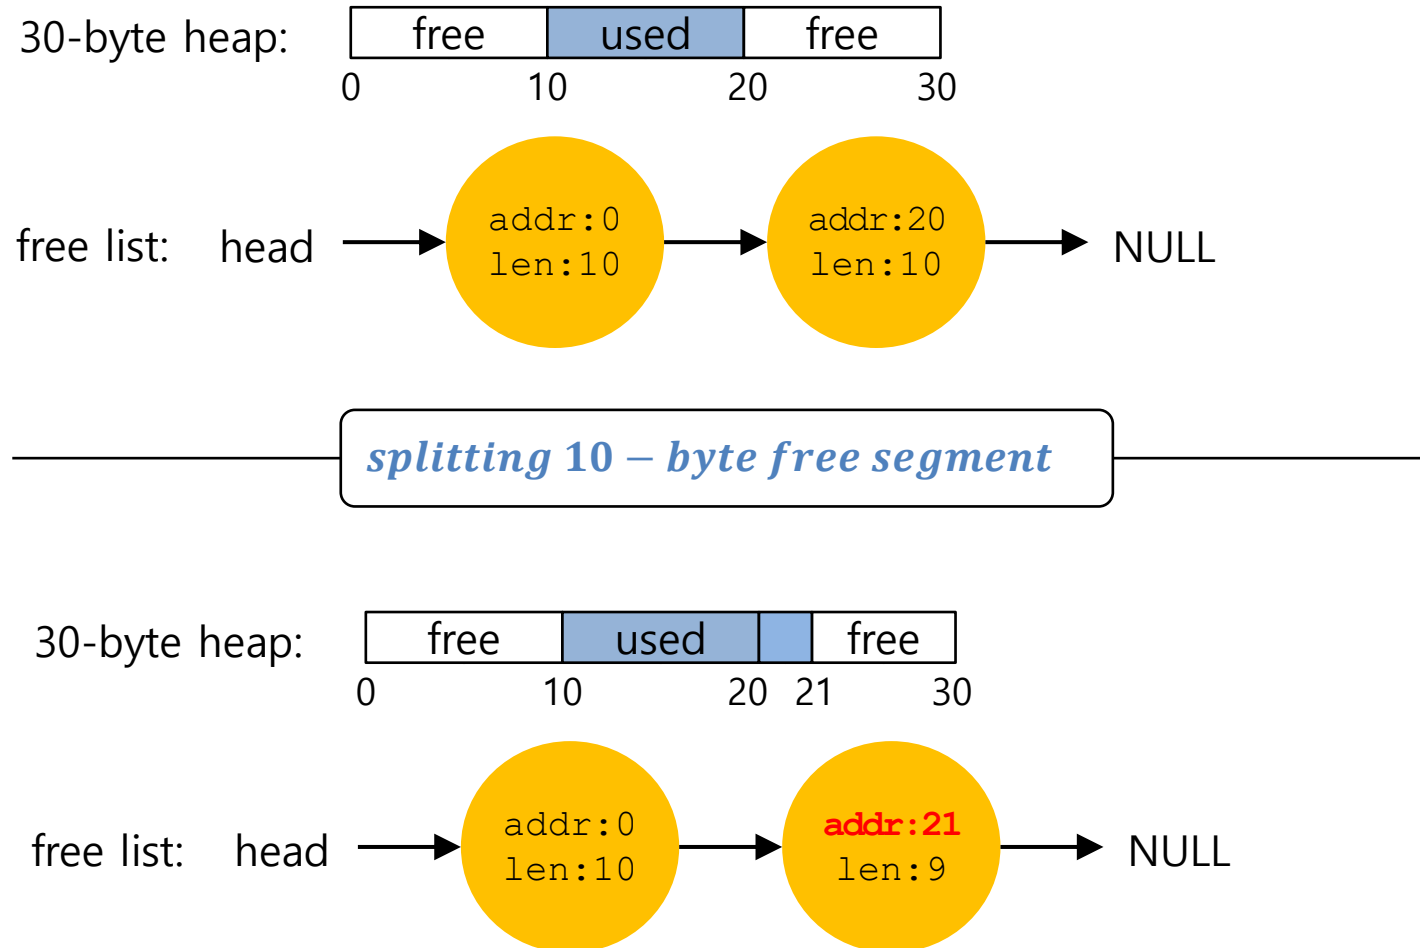
Splitting

- Finding a free chunk of memory that can satisfy the request and splitting it into two
 - When request for memory allocation is **smaller** than the size of free chunks




Splitting

- Two 10-bytes free segment with **1-byte request**



Coalescing


- If a user requests memory that is **bigger than free chunk size**, the list will **not find** such a free chunk
- Coalescing: **Merge** returning a free chunk with existing chunks into a large single free chunk if **addresses** of them are **nearby**



Tracking The Size of Allocated Regions

- The interface to `free(void *ptr)` does not take a size parameter
 - How does the library know the size of memory region that will be back into free list?
- Most allocators store extra information in a header block


```
ptr = malloc(20);
```



An Allocated Region Plus Header

The Header of Allocated Memory Chunk

- The header minimally contains the size of the allocated memory region
- The header may also contain
 - Additional pointers to speed up deallocation
 - A magic number for integrity checking



```
typedef struct __header_t {  
    int size;  
    int magic;  
} header_t;
```

A Simple Header

The Header of Allocated Memory Chunk

- The size for free region is the size of the header plus the size of the space allocated to the user
- If a user request N bytes, the library searches for a free chunk of size N plus the size of the header
- Simple pointer arithmetic to find the header pointer

```
void free(void *ptr) {  
    header_t *hptr = (void *)ptr - sizeof(header_t);  
}
```

Embedding A Free List

- The memory-allocation library **initializes** the heap and **puts** the first element of **the free list** in the **free space**
 - The library **can't use** `malloc()` to build a list **within itself**

Embedding A Free List

- Description of a node of the list


```
typedef struct __node_t {  
    int size;  
    struct __node_t *next;  
} node_t;
```

- Building heap and putting a free list
 - Assume that the heap is built via mmap() system call

```
// mmap() returns a pointer to a chunk of free space  
node_t *head = mmap(NULL, 4096, PROT_READ|PROT_WRITE,  
                    MAP_ANON|MAP_PRIVATE, -1, 0);  
head->size = 4096 - sizeof(node_t);  
head->next = NULL;
```

A Heap With One Free Chunk

```
// mmap() returns a pointer to a chunk of free space
node_t *head = mmap(NULL, 4096, PROT_READ|PROT_WRITE,
                    MAP_ANON|MAP_PRIVATE, -1, 0);
head->size = 4096 - sizeof(node_t);
head->next = NULL;
```

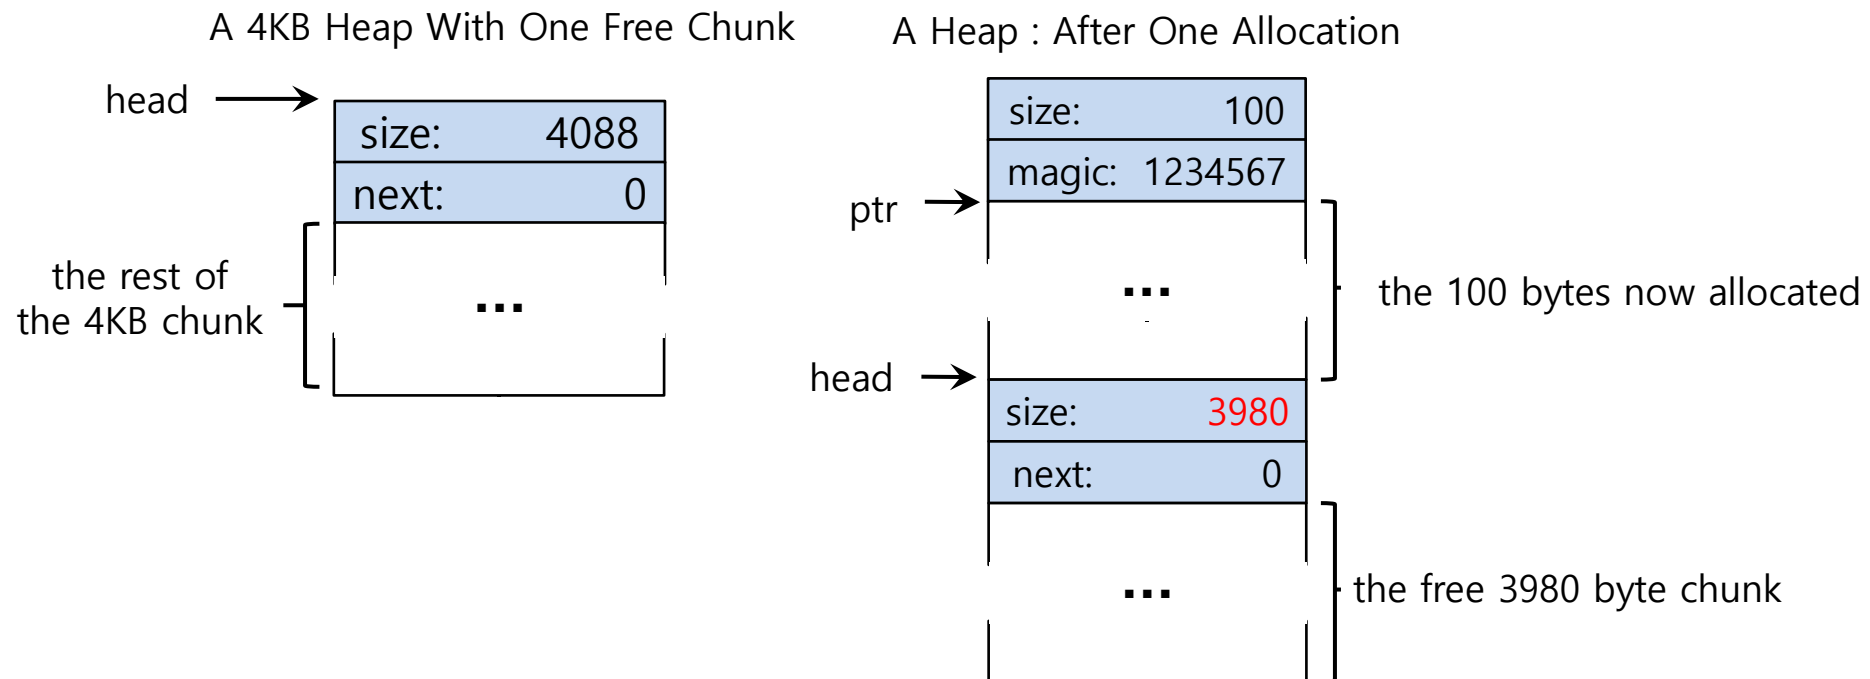


Embedding A Free List: Allocation

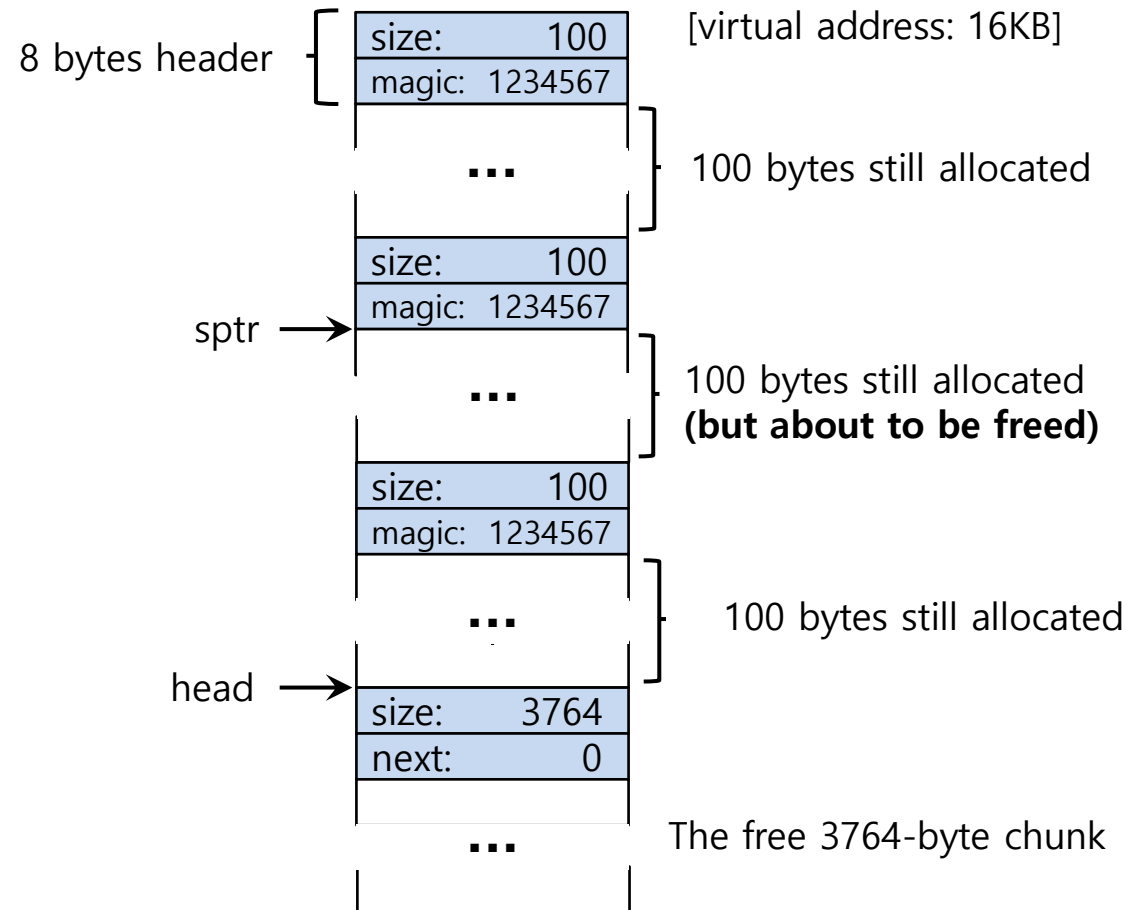
- If a chunk of memory is requested, the library will first find a chunk that is large enough to accommodate the request
- The library will
 - Split the large free chunk into two
 - One for the request and the remaining free chunk
 - Shrink the size of free chunk in the list

Embedding A Free List: Allocation

- Example: a request for 100 bytes by `ptr = malloc(100)`
 - Allocating 108 bytes out of the existing one free chunk
 - shrinking the one free chunk to 3980(4088 minus 108)



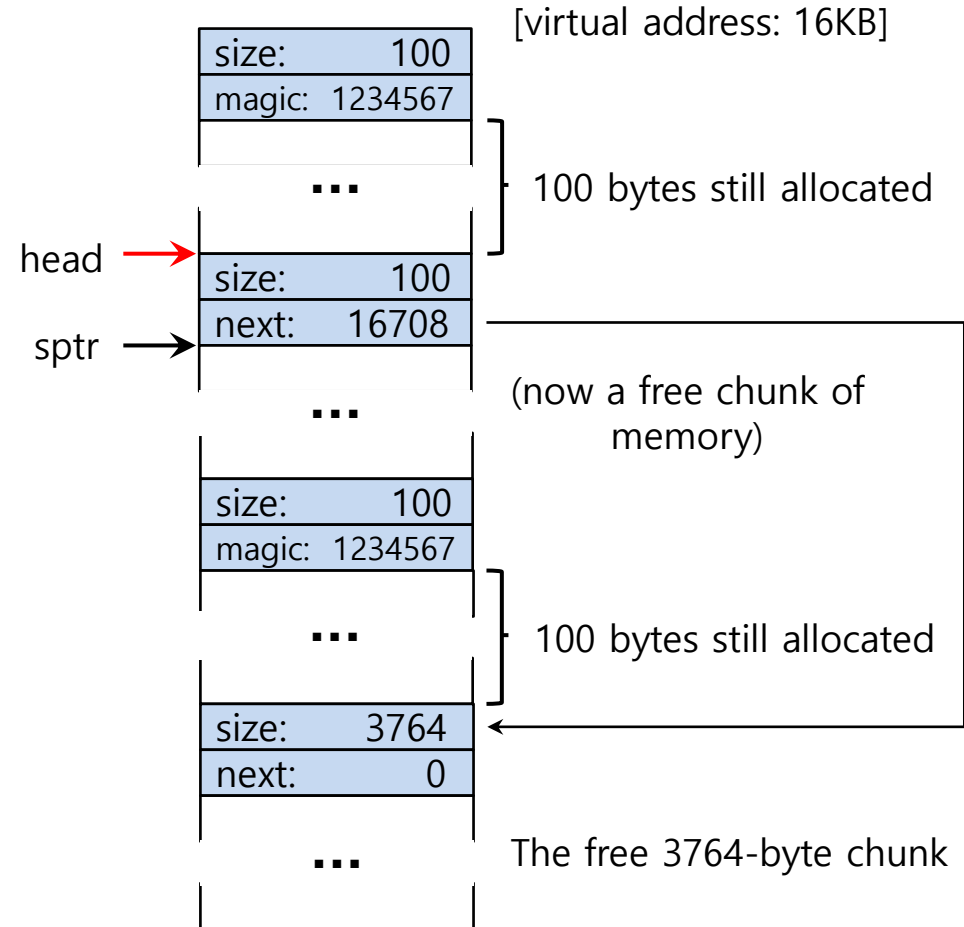
Free Space With Chunks Allocated



Free Space With Three Chunks Allocated

Free Space With `free()`

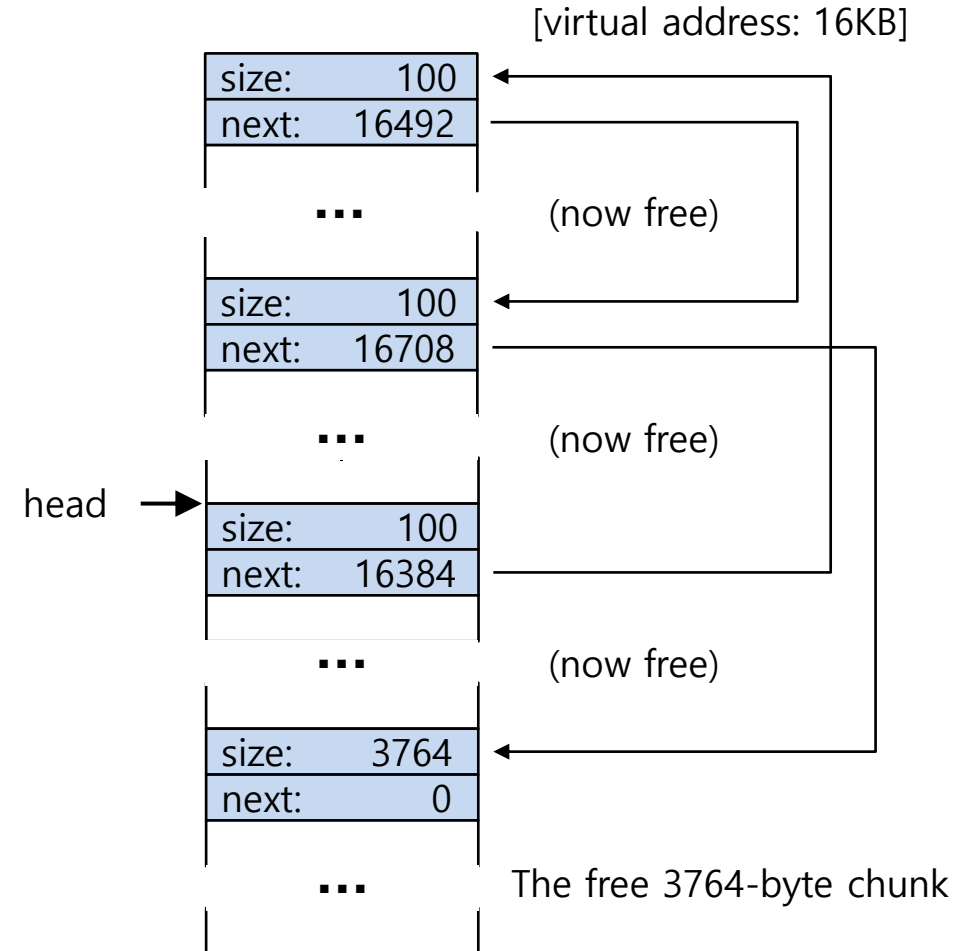
- Example: `free(sptr)`
 - The 100 bytes chunks is **back into** the free list
 - The free list will **start** with a **small chunk**
 - The list header will point the small chunk



Free Space With Freed Chunks


- Let's assume that the last two in-use chunks are freed

- External Fragmentation** occurs
 - Coalescing** is needed in the list



Growing The Heap

- Most allocators **start** with a **small-sized heap** and then **request more** memory from the OS when they run out
 - e.g., `sbrk()`, `brk()` in most UNIX systems



Managing Free Space: Basic Strategies

- Best Fit:
 - Finding free chunks that are **big or bigger than the request**
 - Returning the **one of smallest** in the chunks **in the group** of candidates
- Worst Fit:
 - Finding the **largest free chunks** and allocation the amount of the request
 - **Keeping the remaining chunk** on the free list


Managing Free Space: Basic Strategies




- First Fit:
 - Finding the **first chunk** that is **big enough** for the request
 - Returning the requested amount and remaining the rest of the chunk
- Next Fit:
 - Finding the first chunk that is big enough for the request
 - Searching at **where one was looking** at instead of the beginning of the list

Examples of Basic Strategies


- Allocation Request Size 15



- Result of Best-fit



- Result of Worst-fit



Other Approaches: Segregated List

- Segregated List:
 - Keeping free chunks in different size in a separate list for the size of popular request
 - New Complication:
 - **How much** memory should dedicate to **the pool of memory** that serves **specialized requests** of a given size?
 - **Slab allocator** handles this issue

Other Approaches: Segregated List


■ Slab Allocator

- Allocate a number of object caches
 - The objects are likely to be requested frequently
 - e.g., locks, file-system inodes, etc
- **Request some memory** from a more general memory allocator when **a given cache is running low** on free space

Other Approaches: Buddy Allocation

■ Binary Buddy Allocation

- The allocator **divides free space** by two **until a block** that is big enough to accommodate the request is **found**



64KB free space for 7KB request

Other Approaches: Buddy Allocation

- Buddy allocation can suffer from **internal fragmentation**
- Buddy system makes **coalescing** simple
 - **Coalescing** two blocks into the next level of blocks