



Memory Paging

Concept of Paging

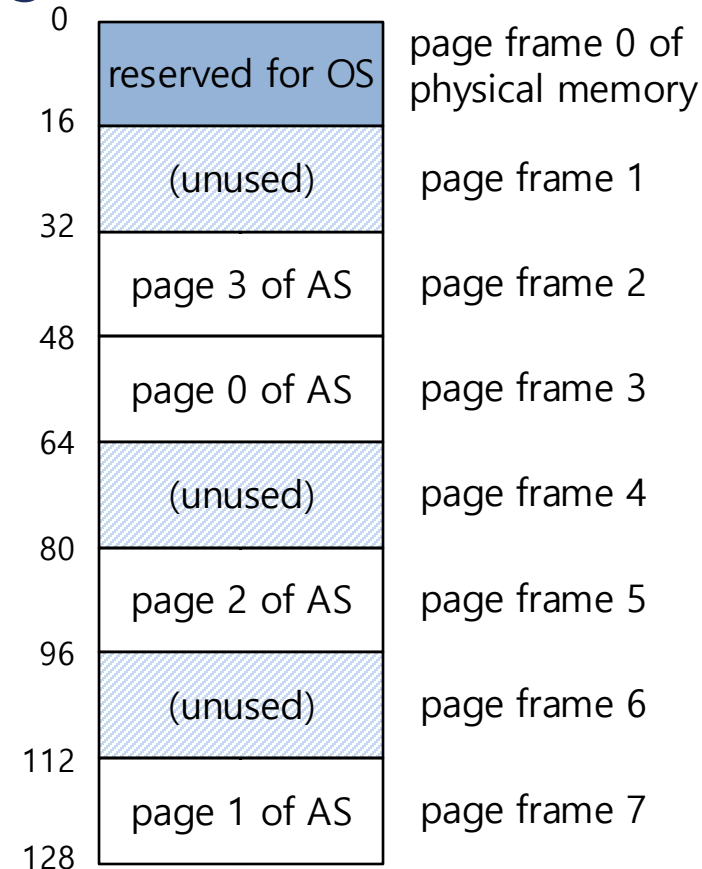

- Paging **splits up** address space into **fixed-sized** unit called a **page**
 - Segmentation: variable size of logical segments(code, stack, heap, etc.)
- With paging, **physical memory** is also **split** into some number of pages called a **page frame**
- **Page table** per process is needed **to translate** the virtual address to physical address

Advantages Of Paging

- **Flexibility:** Supporting the abstraction of address space effectively
 - Don't need assumption how heap and stack grow and are used
- **Simplicity:** ease of free-space management
 - The page in address space and the page frame are the same size
 - Easy to allocate and keep a free list

Example: A Simple Paging


- 128-byte physical memory with 16 bytes page frames
- 64-byte address space with 16 bytes pages




64-Byte Address Space Placed In Physical Memory

Address Translation

- Two components in the virtual address
 - VPN: virtual page number
 - Offset: offset within the page




- Example: virtual address 21 in 64-byte address space



Example: Address Translation

- The virtual address 21 in 64-byte address space

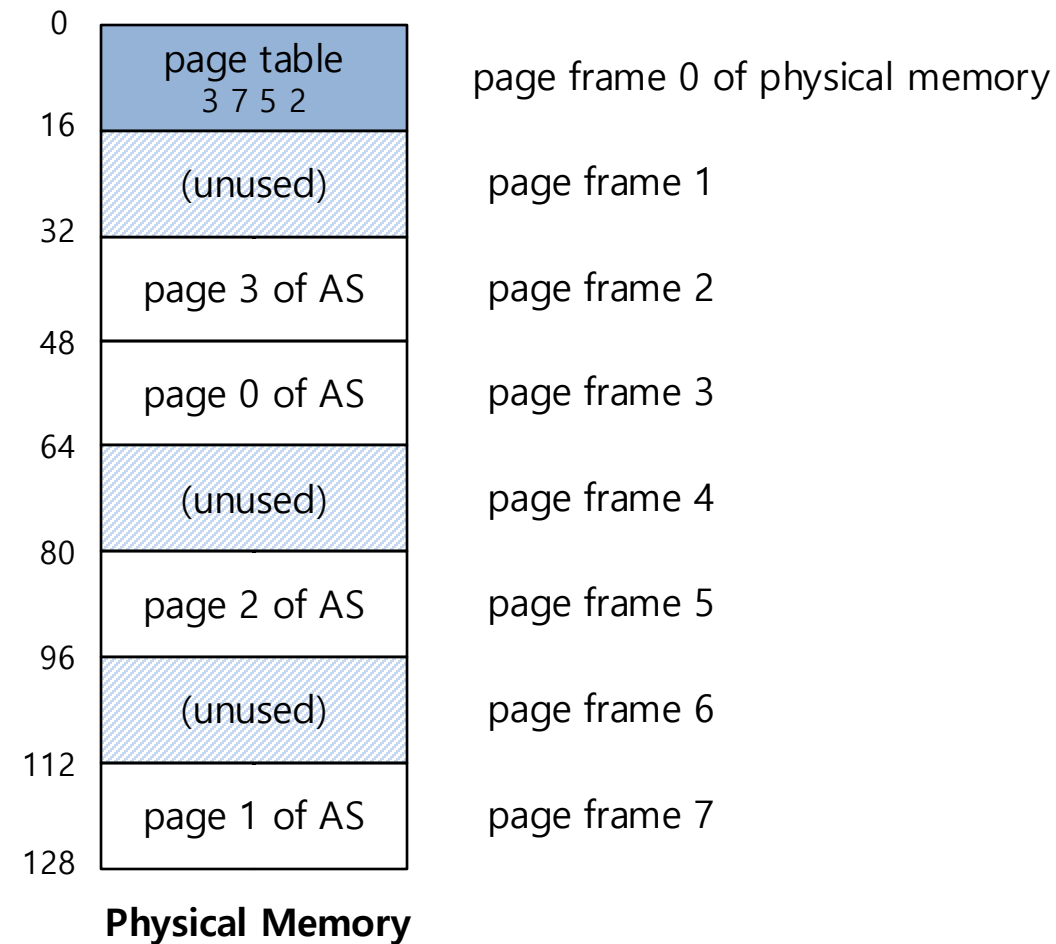


Where Are Page Tables Stored?



- Page tables can get awfully large
 - 32-bit address space with 4-KB pages, 20 bits for VPN
 - $4MB = 2^{20} \text{ entries} * 4 \text{ Bytes per page table entry}$
- Page tables for each process are stored in memory

Example: Page Table in Kernel Physical Memory




What Is In The Page Table?

- The page table is just a **data structure** that is used to map the virtual address to physical address
 - Simplest form: a linear page table, an array
- The OS **indexes** the array by VPN, and looks up the page-table entry

Common Flags Of Page Table Entry

- **Valid Bit:** Indicating whether the particular translation is valid
- **Protection Bit:** Indicating whether the page could be read from, written to, or executed from
- **Present Bit:** Indicating whether this page is in physical memory or on disk(swapped out)
- **Dirty Bit:** Indicating whether the page has been modified since it was brought into memory
- **Reference Bit(Accessed Bit):** Indicating that a page has been accessed

Example: x86 Page Table Entry



An x86 Page Table Entry(PTE)

- P: present
- R/W: read/write bit
- U/S: supervisor
- A: accessed bit
- D: dirty bit
- PFN: the page frame number

Paging: Too Slow

- To find a location of the desired PTE, the **starting location** of the page table is **needed**
- For every memory reference, paging requires the OS to perform one **extra memory reference**

Accessing Memory With Paging

```
1      // Extract the VPN from the virtual address
2      VPN = (VirtualAddress & VPN_MASK) >> SHIFT
3
4      // Form the address of the page-table entry (PTE)
5      PTEAddr = PTBR + (VPN * sizeof(PTE))
6
7      // Fetch the PTE
8      PTE = AccessMemory(PTEAddr)
9
10     // Check if process can access the page
11     if (PTE.Valid == False)
12         RaiseException(SEGMENTATION_FAULT)
13     else if (CanAccess(PTE.ProtectBits) == False)
14         RaiseException(PROTECTION_FAULT)
15     else
16         // Access is OK: form physical address and fetch it
17         offset = VirtualAddress & OFFSET_MASK
18         PhysAddr = (PTE.PFN << PFN_SHIFT) | offset
19         Register = AccessMemory(PhysAddr)
```

A Memory Trace

- Example: A Simple Memory Access

```
int array[1000];  
...  
for (i = 0; i < 1000; i++)  
    array[i] = 0;
```

- Compile and execute

```
prompt> gcc -o array array.c -Wall -o  
prompt> ./array
```

- Resulting Assembly code

```
0x1024 movl $0x0, (%edi,%eax,4)  
0x1028 incl %eax  
0x102c cmpl $0x03e8,%eax  
0x1030 jne 0x1024
```

A Virtual(And Physical) Memory Trace

