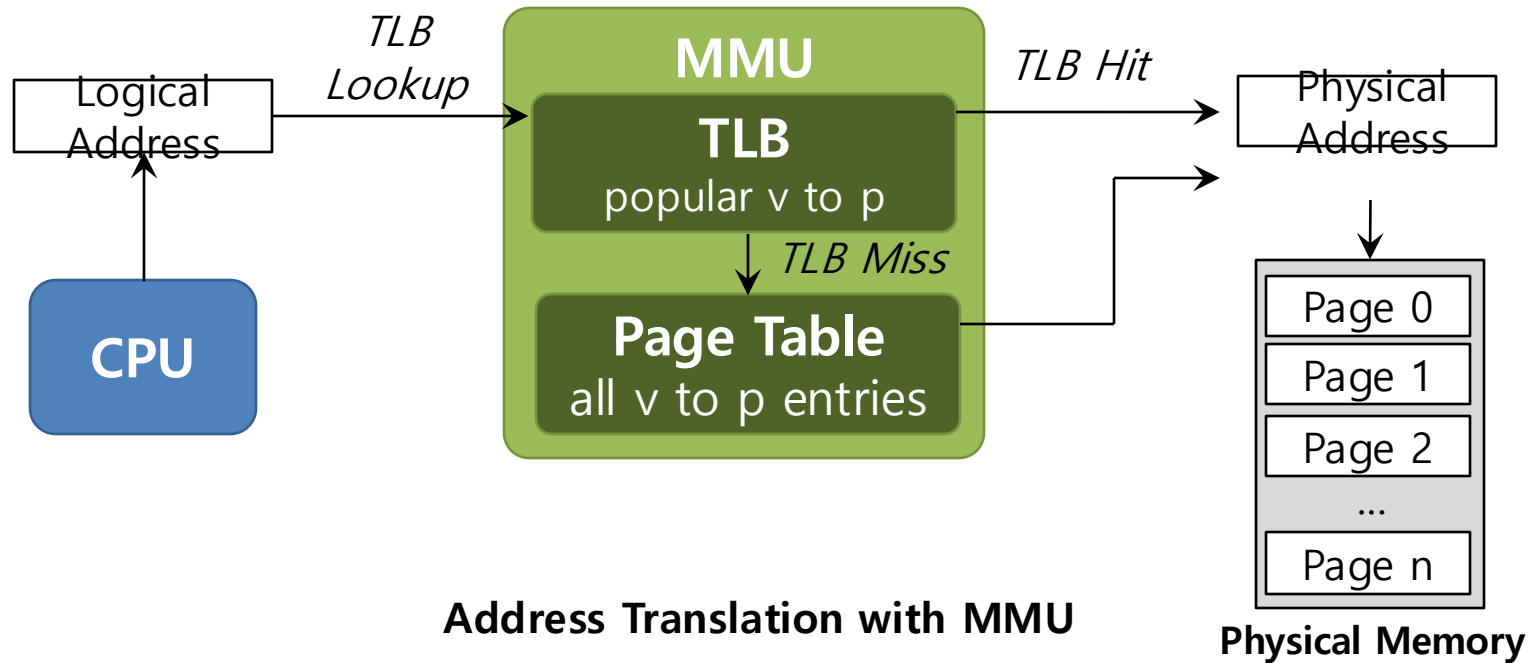




# Locality and Caching

# TLB (Translation Look-Aside Buffer)

- Part of the chip's memory-management unit(MMU)
- A hardware cache of **popular** virtual-to-physical address translation



# TLB Basic Algorithms

```
1: VPN = (VirtualAddress & VPN_MASK ) >> SHIFT
2: (Success , TlbEntry) = TLB_Lookup(VPN)
3:     if(Success == True){ // TLB Hit
4:         if(CanAccess(TlbEntry.ProtectBit) == True ){
5:             offset = VirtualAddress & OFFSET_MASK
6:             PhysAddr = (TlbEntry.PFN << SHIFT) | Offset
7:             AccessMemory( PhysAddr )
8:         }else RaiseException(PROTECTION_ERROR)
```

- (1 lines) extract the virtual page number(VPN)
- (2 lines) check if the TLB holds the translation for this VPN
- (5-8 lines) extract the page frame number from the relevant TLB entry, and form the desired physical address and access memory

# TLB Basic Algorithms

```
11:      }else{ //TLB Miss
12:          PTEAddr = PTBR + (VPN * sizeof(PTE))
13:          PTE = AccessMemory(PTEAddr)
14:          (...)
15:      }else{
16:          TLB_Insert( VPN , PTE.PFN , PTE.ProtectBits)
17:          RetryInstruction()
18:      }
19: }
```

- (11-12 lines) The hardware accesses the page table to find the translation
- (16 lines) updates the TLB with the translation

# Example: Accessing An Array

- How a TLB can improve its performance

	OFFSET				
	00	04	08	12	16
VPN = 00					
VPN = 01					
VPN = 03					
VPN = 04					
VPN = 05					
VPN = 06		a[0]	a[1]	a[2]	
VPN = 07	a[3]	a[4]	a[5]	a[6]	
VPN = 08	a[7]	a[8]	a[9]		
VPN = 09					
VPN = 10					
VPN = 11					
VPN = 12					
VPN = 13					
VPN = 14					
VPN = 15					

```
0:      int sum = 0 ;
1:      for( i=0; i<10; i++) {
2:                  sum+=a[i];
3:      }
```


**The TLB improves performance  
due to **spatial locality****

3 misses and 7 hits  
Thus **TLB hit rate** is 70%

# Locality


## ■ Temporal Locality

- An instruction or data item that has been recently accessed will likely be re-accessed soon in the future



## ■ Spatial Locality

- If a program accesses memory at address  $x$ , it will likely soon access memory near  $x$



# Who Handles The TLB Miss?

- Hardware handle the TLB miss entirely on CISC
  - The hardware has to know exactly where the page tables are located in memory
  - The hardware would “walk” the page table, find the correct page-table entry and extract the desired translation, update and retry instruction
  - hardware-managed TLB

# Who Handles The TLB Miss?

- RISC have what is known as a software-managed TLB
  - On a TLB miss, the hardware raises exception( trap handler )
    - Trap handler is code within the OS that is written with the express purpose of **handling TLB miss**



# TLB Control Flow Algorithm (OS Handled)

```
1:      VPN = (VirtualAddress & VPN_MASK) >> SHIFT
2:      (Success, TlbEntry) = TLB_Lookup(VPN)
3:      if (Success == True) // TLB Hit
4:          if (CanAccess(TlbEntry.ProtectBits) == True)
5:              Offset = VirtualAddress & OFFSET_MASK
6:              PhysAddr = (TlbEntry.PFN << SHIFT) | Offset
7:              Register = AccessMemory(PhysAddr)
8:          else
9:              RaiseException(PROTECTION_FAULT)
10:     else // TLB Miss
11:         RaiseException(TLB_MISS)
```

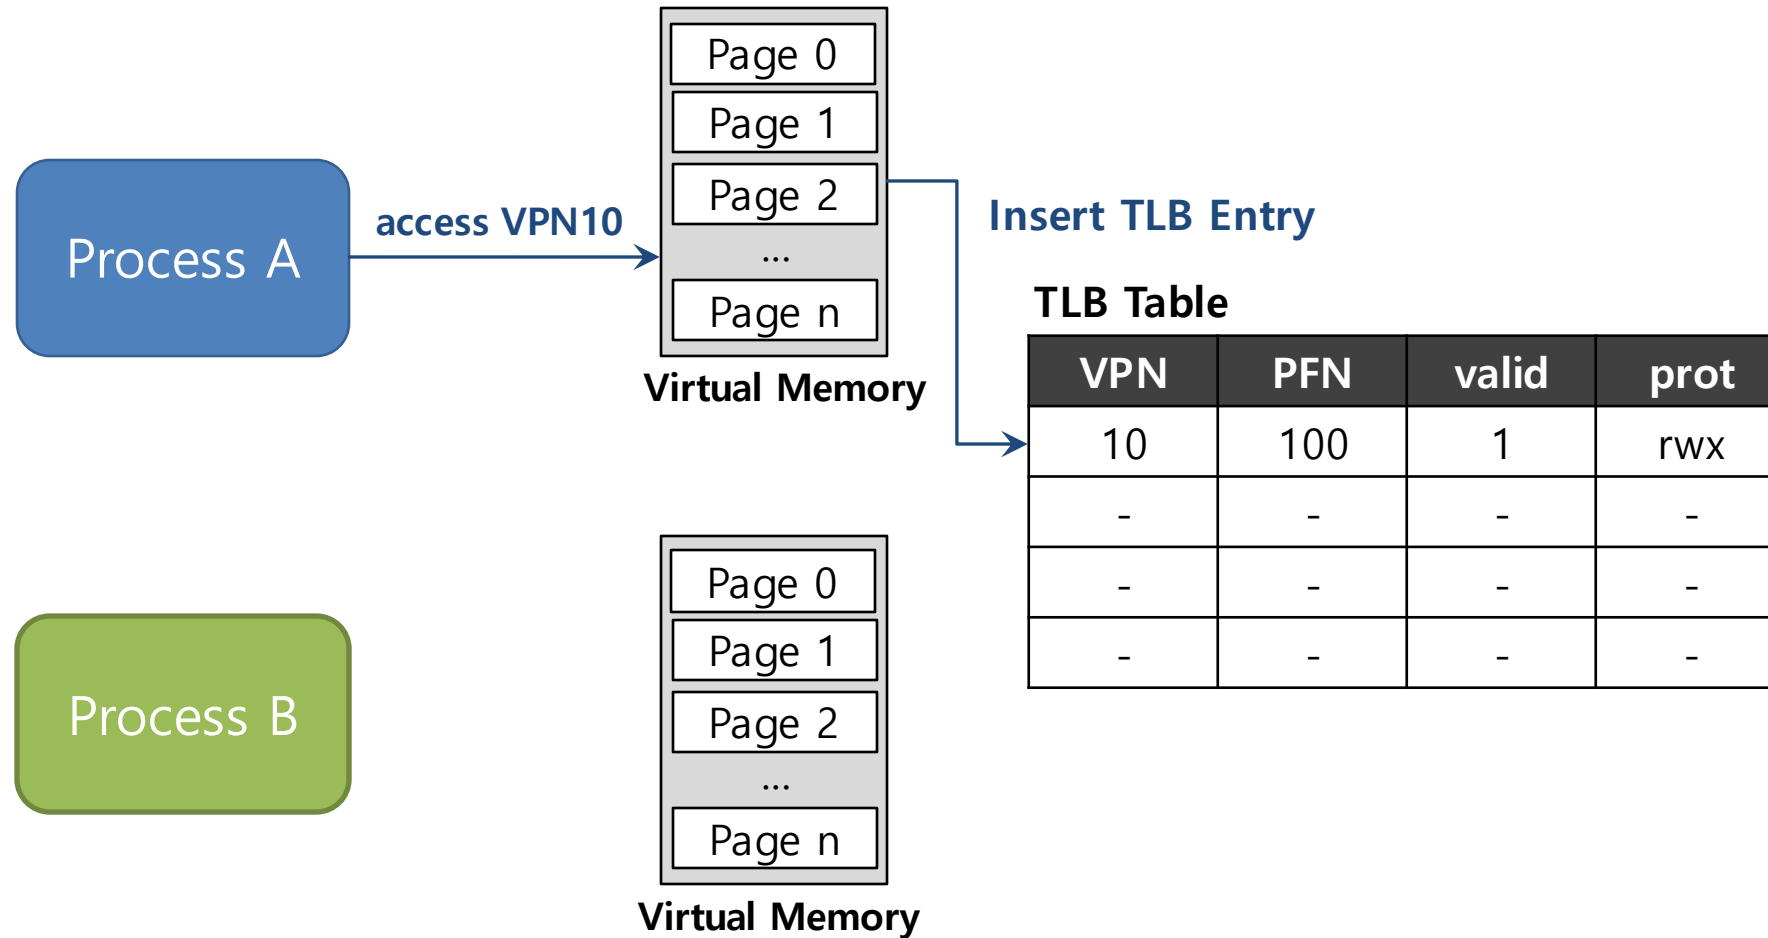
# TLB Entry

- TLB is managed by **Fully-Associative** method
  - A typical TLB might have 32,64, or 128 entries
  - Hardware search the entire TLB in parallel to find the desired translation
  - Other bits: valid bits , protection bits, address-space identifier, dirty bit

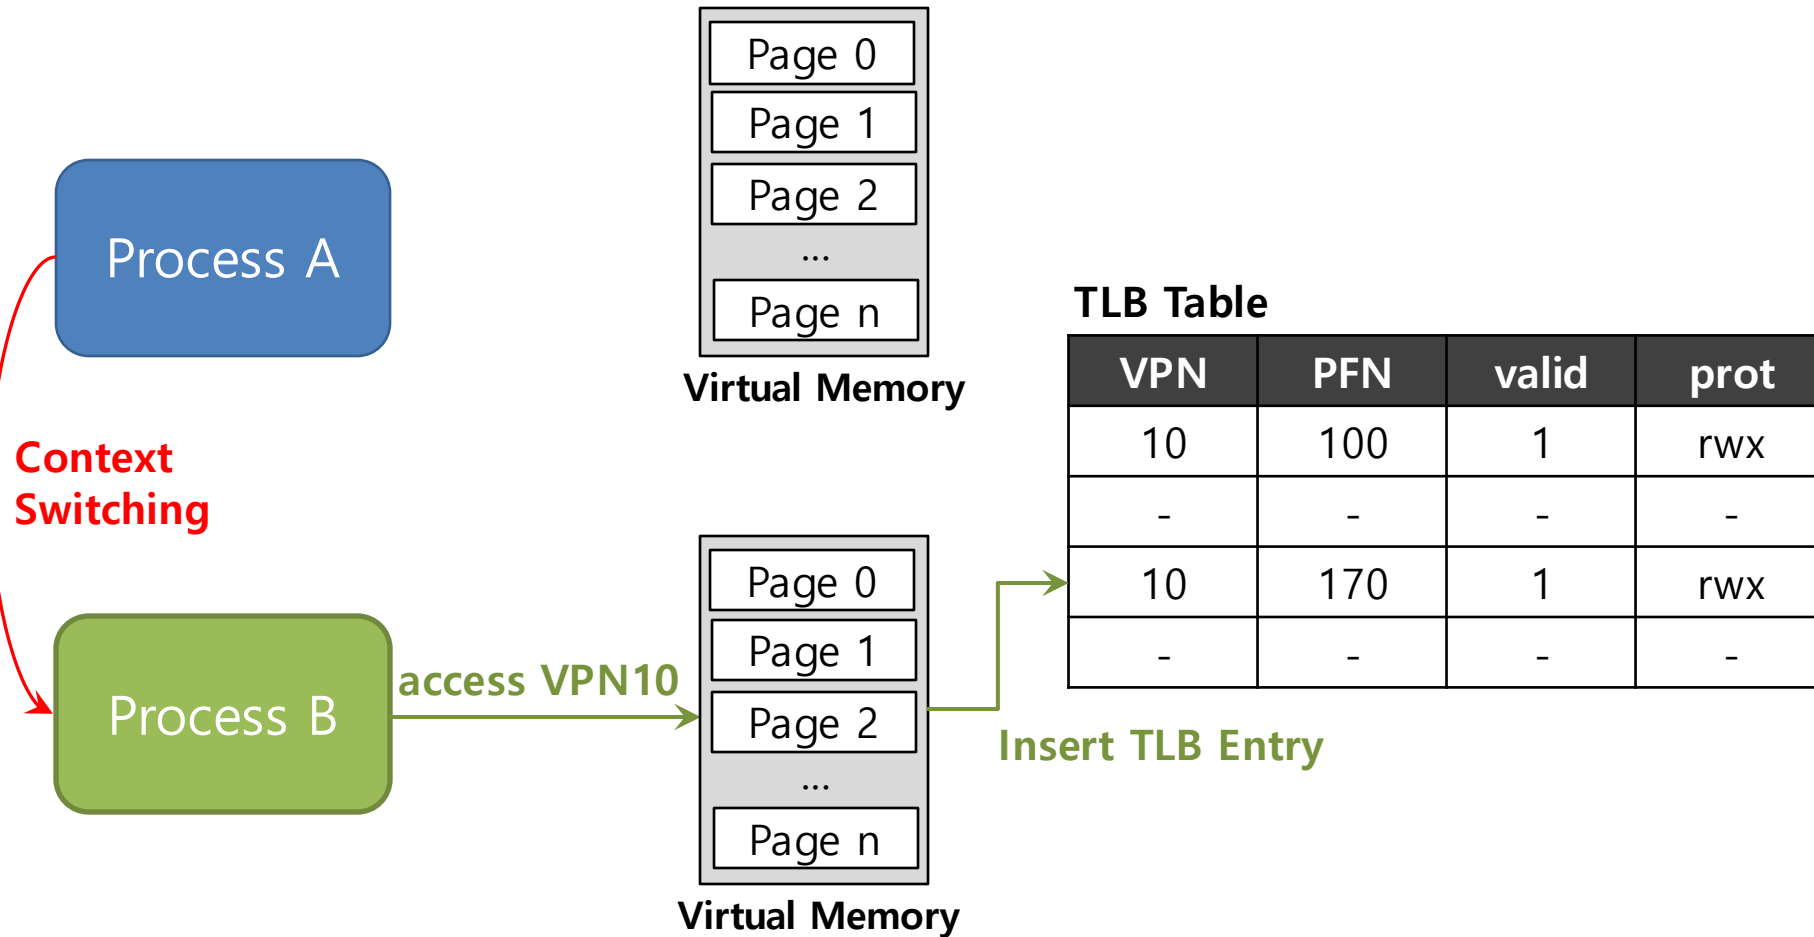


Typical TLB entry look like this

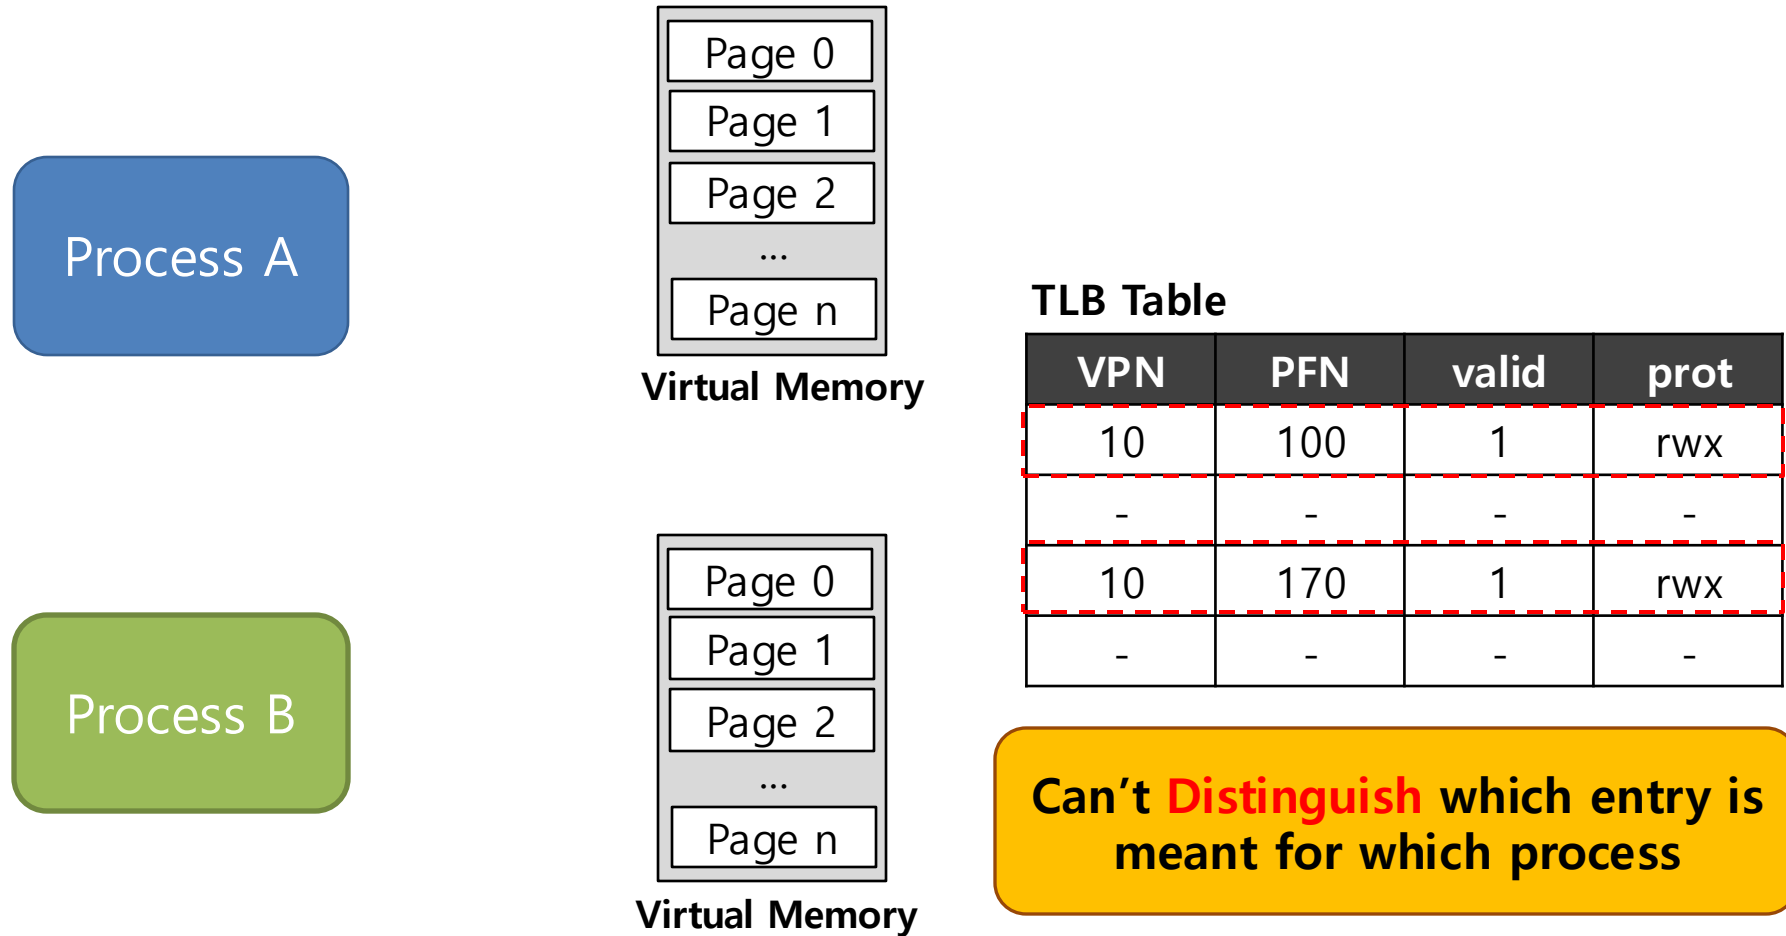
# TLB Issue: Context Switching



# TLB Issue: Context Switching

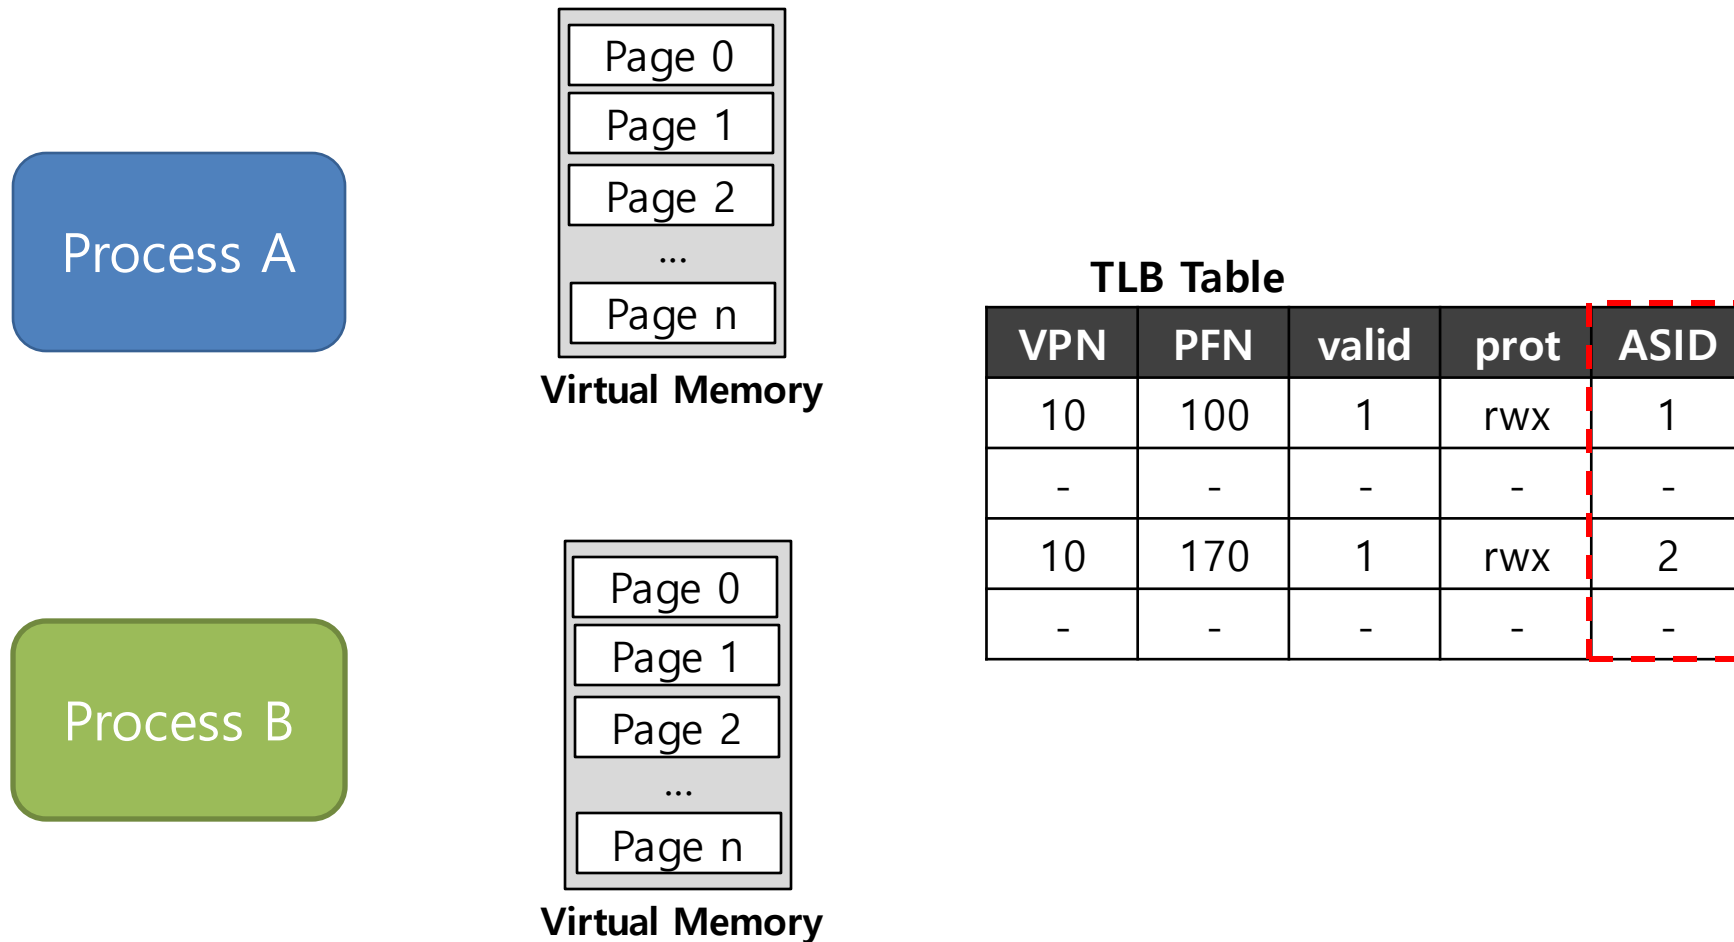


# TLB Issue: Context Switching



# To Solve Problem

- Provide an address space identifier(ASID) field in the TLB



# Another Case


- Two processes **share a page**
  - Process 1 is sharing physical page 101 with Process2
  - P1 maps this page into the 10<sup>th</sup> page of its address space
  - P2 maps this page to the 50<sup>th</sup> page of its address space

VPN	PFN	valid	prot	ASID
10	101	1	rwX	1
-	-	-	-	-
50	101	1	rwX	2
-	-	-	-	-

Sharing of pages is **useful** as it reduces the number of physical pages in use

# TLB Replacement Policy

- LRU(Least Recently Used)
  - Evict an entry that has not recently been used
  - Take advantage of *locality* in the memory-reference stream



Total 11 TLB miss



0	1	2	3	4	5	6	7	8	9	10	11	...		19	...		31
									VPN					G			ASID
										PFN							C D V

Flag	Content
19-bit VPN	The rest reserved for the kernel.
24-bit PFN	Systems can support with up to 64GB of main memory( $2^{24} * 4KB$ pages ).
Global bit(G)	Used for pages that are globally-shared among processes.
ASID	OS can use to distinguish between address spaces.
Coherence bit(C)	determine how a page is cached by the hardware.
Dirty bit(D)	marking when the page has been written.
Valid bit(V)	tells the hardware if there is a valid translation present in the entry.