




# Memory Management Policy

# Beyond Physical Memory: Mechanisms

- Require an additional level in the **memory hierarchy**
  - OS need a place to stash away portions of address space that currently aren't in great demand
  - In modern systems, this role is usually served by a **hard disk drive**



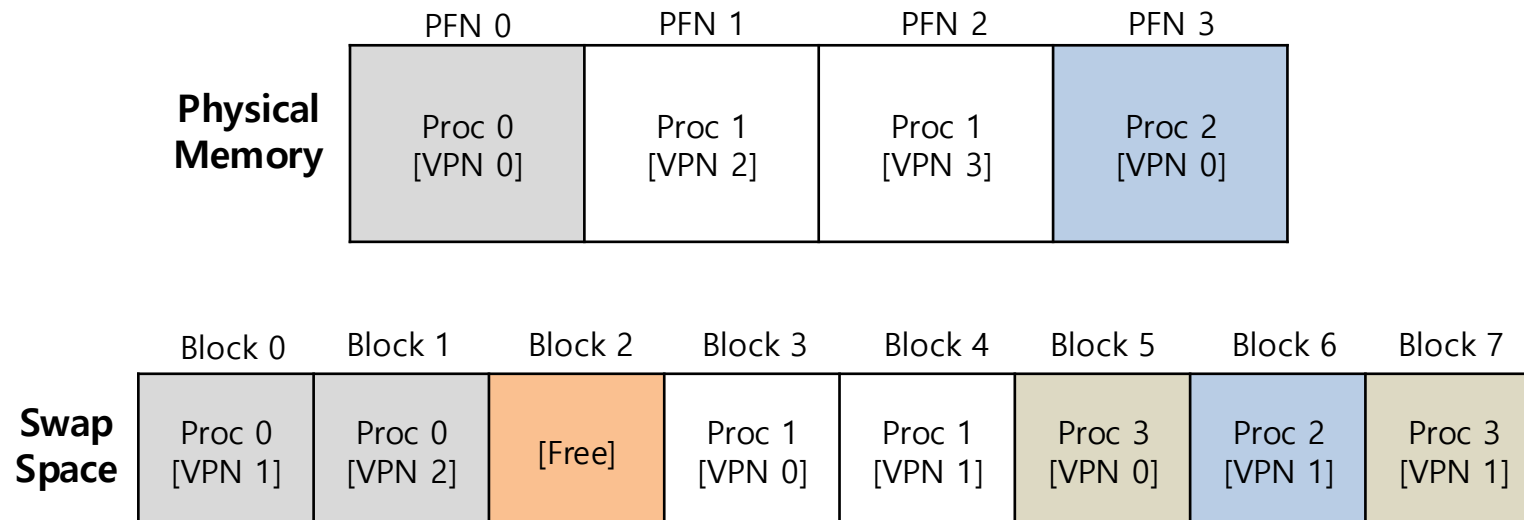
Memory Hierarchy in modern system

# Single Large Address for a Process

- Always need to first arrange for the code or data to be in memory when before calling a function or accessing data
- To Beyond just a single process
  - The addition of swap space allows the OS to support the illusion of a large virtual memory for multiple concurrently-running process

# Swap Space

- Reserve some space on the disk for moving pages back and forth
- OS need to remember to the swap space, in **page-sized unit**



Physical Memory and Swap Space

# Present Bit

- Add some machinery higher up in the system in order to support swapping pages to and from the disk
  - When the hardware looks in the PTE, it may find that the page is not present in physical memory

Value	Meaning
<b>1</b>	page is present in physical memory
<b>0</b>	The page is not in memory but rather on disk.

# What If Memory Is Full ?

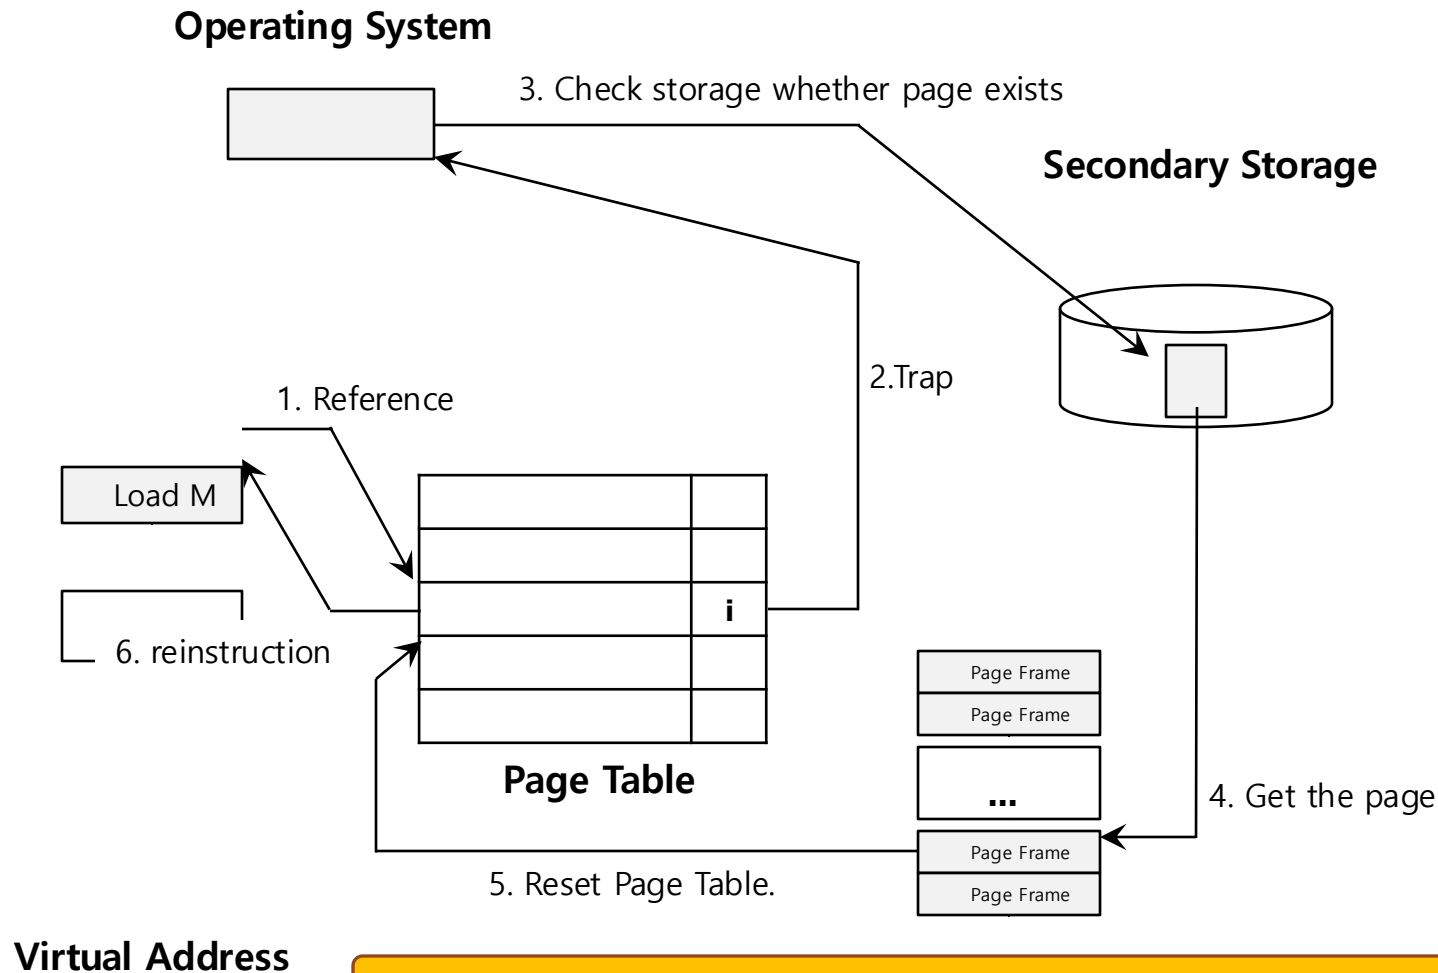
- The OS like to page out pages to make room for the new pages the OS is about to bring in
  - The process of picking a page to kick out, or replace is known as **page-replacement** policy

# The Page Fault

- Accessing page that is not in physical memory
  - If a page is not present and has been swapped disk, the OS need to swap the page into memory in order to service the page fault

# Page Fault Control Flow

- PTE used for data such as the PFN of the page for a disk addresses



When the OS receives a page fault, it looks in the PTE and issues the request to disk.



# Page Fault Control Flow – Hardware

```
1:      VPN = (VirtualAddress & VPN_MASK) >> SHIFT
2:      (Success, TlbEntry) = TLB_Lookup(VPN)
3:      if (Success == True) // TLB Hit
4:      if (CanAccess(TlbEntry.ProtectBits) == True)
5:          Offset = VirtualAddress & OFFSET_MASK
6:          PhysAddr = (TlbEntry.PFN << SHIFT) | Offset
7:          Register = AccessMemory(PhysAddr)
8:      else RaiseException(PROTECTION_FAULT)
```

# Page Fault Control Flow – Hardware

```
9:         else // TLB Miss
10:         PTEAddr = PTBR + (VPN * sizeof(PTE))
11:         PTE = AccessMemory(PTEAddr)
12:         if (PTE.Valid == False)
13:             RaiseException(SEGMENTATION_FAULT)
14:         else
15:             if (CanAccess(PTE.ProtectBits) == False)
16:                 RaiseException(PROTECTION_FAULT)
17:             else if (PTE.Present == True)
18:                 // assuming hardware-managed TLB
19:                 TLB_Insert(VPN, PTE.PFN, PTE.ProtectBits)
20:                 RetryInstruction()
21:             else if (PTE.Present == False)
22:                 RaiseException(PAGE_FAULT)
```

# Page Fault Control Flow – Software

```
1:      PFN = FindFreePhysicalPage()
2:      if (PFN == -1) // no free page found
3:          PFN = EvictPage() // run replacement algorithm
4:          DiskRead(PTE.DiskAddr, pfn) // sleep (waiting for I/O)
5:          PTE.present = True // update page table with present
6:          PTE.PFN = PFN // bit and translation (PFN)
7:          RetryInstruction() // retry instruction
```

- ◆ The OS must find a physical frame for the **soon-be-faulted-in page** to reside within
- ◆ If there is no such page, waiting for the **replacement algorithm** to run and kick some pages out of memory

# When Replacements Really Occur

- OS waits until memory is entirely full, and only then replaces a page to make room for some other page
  - This is a little bit unrealistic, and there are many reason for the OS to keep a small portion of memory free more proactively
- Swap Daemon, Page Daemon
  - There are fewer than **LW pages** available, a background thread that is responsible for freeing memory runs
  - The thread evicts pages until there are **HW pages** available

# Beyond Physical Memory: Policies

- Memory pressure forces the OS to start **paging out** pages to make room for actively-used pages
- Deciding which page to evict is encapsulated within the replacement policy of the OS

# Cache Management

- Goal in picking a replacement policy for this cache is to minimize the number of cache misses
- The number of cache hits and misses let us calculate the *average memory access time (AMAT)*

$$AMAT = (P_{Hit} * T_M) + (P_{Miss} * T_D)$$

Argument	Meaning
$T_M$	The cost of accessing memory
$T_D$	The cost of accessing disk
$P_{Hit}$	The probability of finding the data item in the cache(a hit)
$P_{Miss}$	The probability of not finding the data in the cache(a miss)

# The Optimal Replacement Policy



- Leads to the fewest number of misses overall
  - Replaces the page that will be accessed furthest in the future
  - Resulting in the **fewest-possible** cache misses
- Serve only as a comparison point, to know how close we are to **perfect**

# Tracing the Optimal Policy

## Reference Row

0 1 2 0 1 3 0 3 1 2 1

Access	Hit/Miss?	Evict	Resulting Cache State
0	Miss		0
1	Miss		0,1
2	Miss		0,1,2
0	Hit		0,1,2
1	Hit		0,1,2
3	Miss	2	0,1,3
0	Hit		0,1,3
3	Hit		0,1,3
1	Hit		0,1,3
2	Miss	3	0,1,2
1	Hit		0,1,2

Hit rate is  $\frac{\text{Hits}}{\text{Hits} + \text{Misses}} = 54.6\%$

Future is not known



# A Simple Policy: FIFO

- Pages were placed in a queue when they enter the system
- When a replacement occurs, the page on the tail of the queue(the “**First-in**” pages) is evicted
  - It is simple to implement, but can't determine the importance of blocks

# Tracing the FIFO Policy

Reference Row

0 1 2 0 1 3 0 3 1 2 1

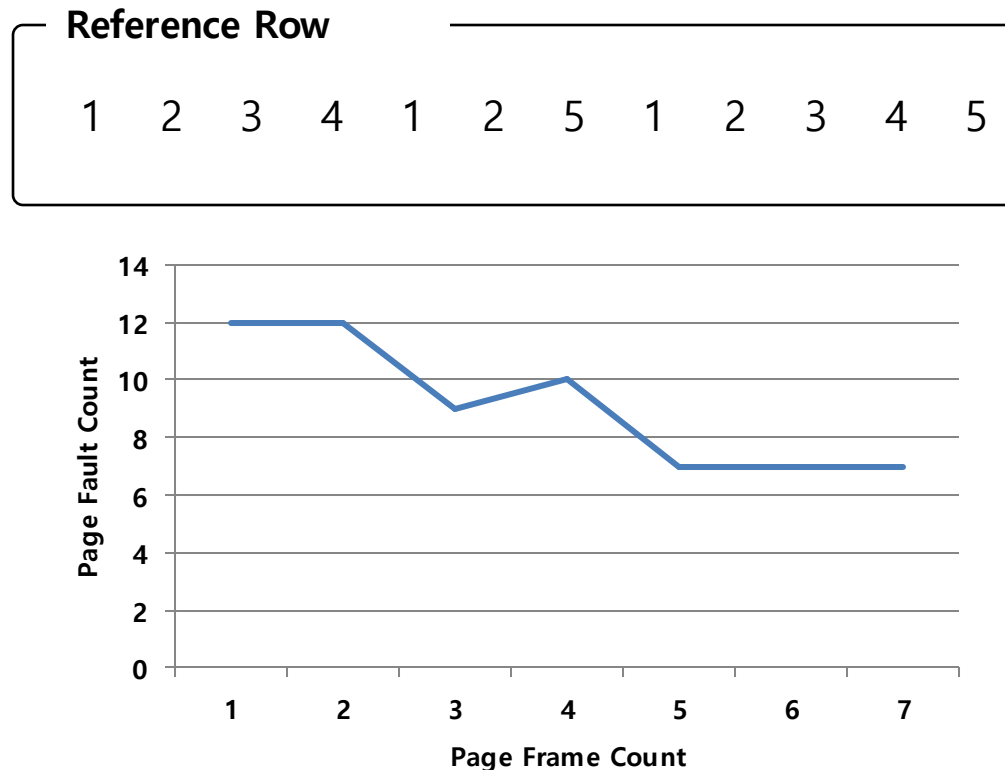
Access	Hit/Miss?	Evict	Resulting Cache State
0	Miss		0
1	Miss		0,1
2	Miss		0,1,2
0	Hit		0,1,2
1	Hit		0,1,2
3	Miss	0	1,2,3
0	Miss	1	2,3,0
3	Hit		2,3,0
1	Miss		3,0,1
2	Miss	3	0,1,2
1	Hit		0,1,2

Hit rate is  $\frac{\text{Hits}}{\text{Hits}+\text{Misses}} = 36.4\%$

Even though page 0 had been accessed a number of times, **FIFO still kicks it out**

# BELADY'S ANOMALY

- We would expect the cache hit rate to **increase** when the cache gets larger. But in this case, with FIFO, it gets worse




# Another Simple Policy: Random

- Picks a random page to replace under memory pressure.
  - It doesn't really try to be too intelligent in picking which blocks to evict
  - Random does depends entirely upon how lucky Random gets in its choice

Access	Hit/Miss?	Evict	Resulting Cache State
0	Miss		0
1	Miss		0,1
2	Miss		0,1,2
0	Hit		0,1,2
1	Hit		0,1,2
3	Miss	0	1,2,3
0	Miss	1	2,3,0
3	Hit		2,3,0
1	Miss	3	2,0,1
2	Hit		2,0,1
1	Hit		2,0,1

# Random Performance

- Sometimes, **Random is as good as optimal**, achieving 6 hits on the example trace



Random Performance over 10,000 Trials

# Using History

- Lean on the past and use history
  - Two type of historical information

Historical Information	Meaning	Algorithms
<b>recency</b>	The more recently a page has been accessed, the more likely it will be accessed again	LRU
<b>frequency</b>	If a page has been accessed many times, It should not be replcaed as it clearly has some value	LFU

# Using History : LRU

- Replaces the least-recently-used page


Reference Row

0 1 2 0 1 3 0 3 1 2 1

Access	Hit/Miss?	Evict	Resulting Cache State
0	Miss		0
1	Miss		0,1
2	Miss		0,1,2
0	Hit		1,2,0
1	Hit		2,0,1
3	Miss	2	0,1,3
0	Hit		1,3,0
3	Hit		1,0,3
1	Hit		0,3,1
2	Miss	0	3,1,2
1	Hit		3,2,1

# Workload Example : The No-Locality Workload

- Each reference is to a random page within the set of accessed pages
  - Workload accesses 100 unique pages over time
  - Choosing the next page to refer to at random




When the cache is large enough to fit the entire workload, it also **doesn't matter** which policy you use




# Workload Example : The 80-20 Workload

- Exhibits locality: 80% of the **reference** are made to 20% of the page
- The remaining 20% of the **reference** are made to the remaining 80% of the pages



# Workload Example : The Looping Sequential

- Refer to 50 pages in sequence.
  - Starting at 0, then 1, ... up to page 49, and then we Loop, repeating those accesses, for total of 10,000 accesses to 50 unique pages



# Implementing Historical Algorithms

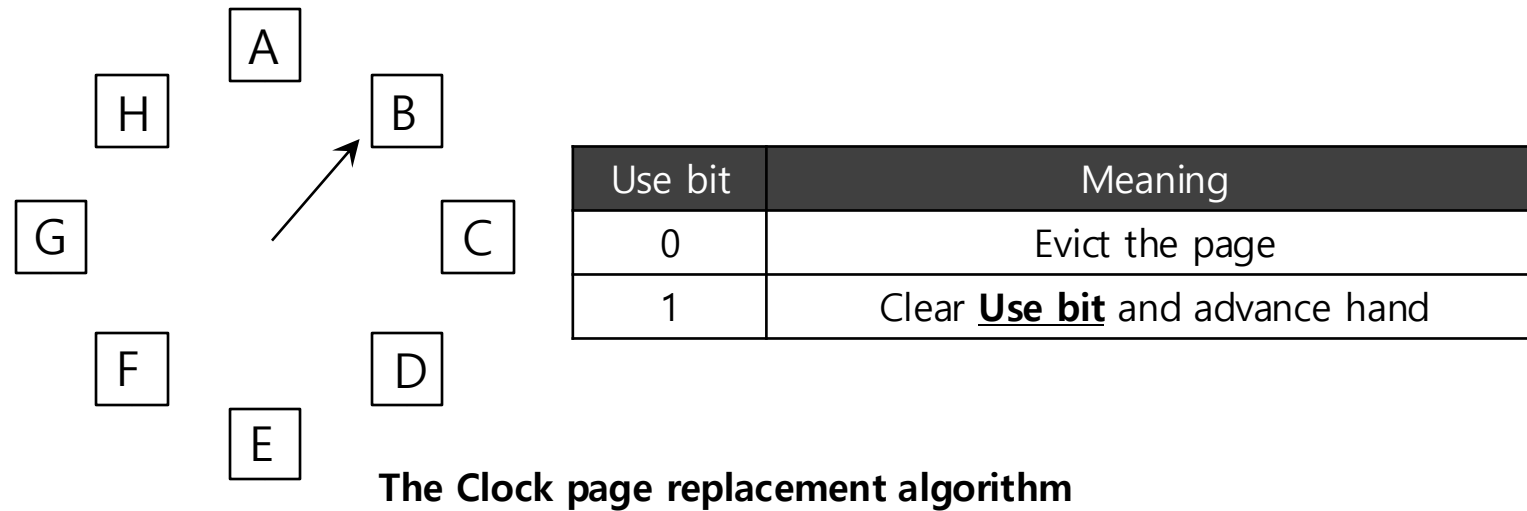
- To keep track of which pages have been least-and-recently used, the system has to do some accounting work on every memory reference
  - Add a little bit of hardware support

# Approximating LRU

- Require some hardware support, in the form of a use bit
  - Whenever a **page is referenced**, the use bit is set by hardware to 1
  - Hardware **never** clears the bit, though; that is the responsibility of the OS
- Clock Algorithm
  - All pages of the system arranges in a circular list
  - A clock hand points to some particular page to begin with

# Clock Algorithm


- The algorithm continues until it finds a use bit that is set to 0



When a page fault occurs, the page the hand is pointing to is inspected.  
The action taken depends on the Use bit

# Workload with Clock Algorithm

- Clock algorithm doesn't do as well as perfect LRU, it does better than approach that don't consider history at all



# Considering Dirty Pages

- The hardware include a modified bit (a.k.a dirty bit)
  - Page has been modified and is thus dirty, it must be written back to disk to evict it
  - Page has not been modified, the eviction is free



# Page Selection Policy

- The OS has to decide when to bring a page into memory
- Presents the OS with some different options



# Prefetching

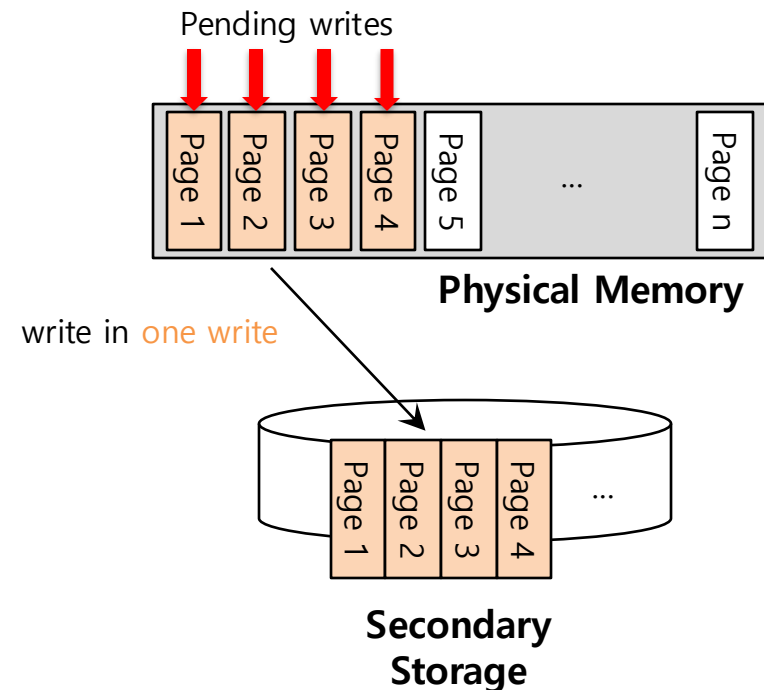
- The OS guess that a page is about to be used, and thus bring it in ahead of time



Page 2 likely soon be accessed and thus should be brought into memory too

# Clustering, Grouping

- Collect a number of **pending writes** together in memory and write them to disk in **one write**
  - Perform a **single large write** more efficiently than **many small ones**



# Thrashing

- Memory is **oversubscribed** and the memory demands of the set of running processes **exceeds** the available physical memory
  - Decide not to run a subset of processes
  - Reduced set of processes working sets fit in memory

