



Concurrency, Concepts and Case Studies

Thread



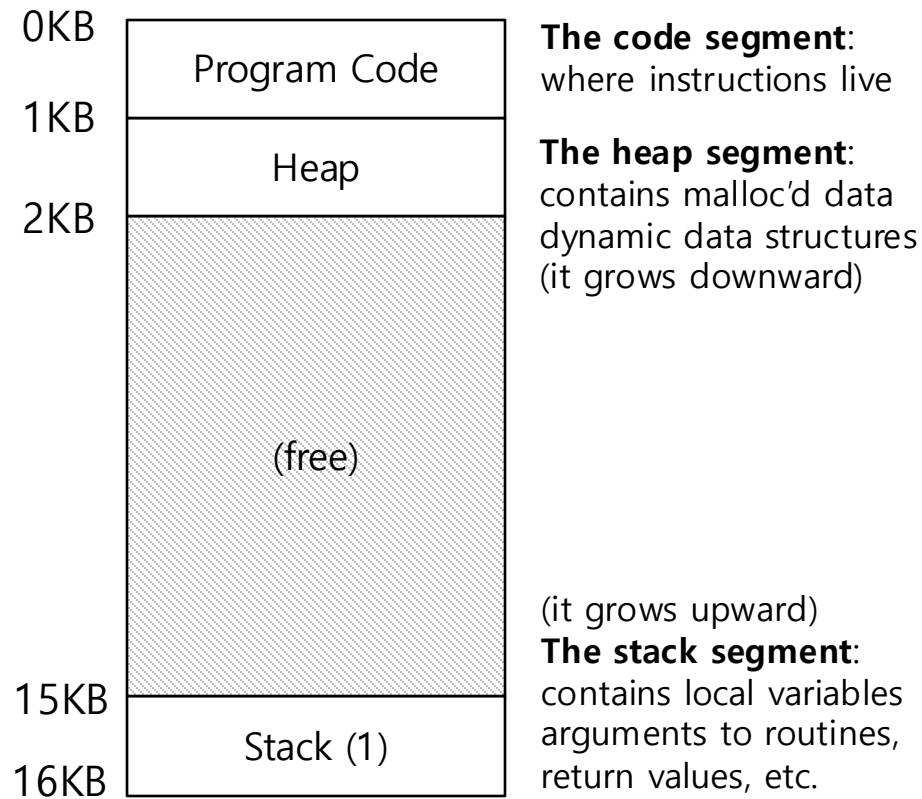
- A new abstraction for a single running process
- Multi-threaded program
 - A multi-threaded program has more than one point of execution
 - Multiple PCs (Program Counter)
 - They **share** the same **address space**

Context switch between threads


- Each thread has its own program counter and set of registers
 - One or more **thread control blocks(TCBs)** are needed to store the state of each thread
- When switching from running one (T1) to running the other (T2),
 - The register state of T1 be saved.
 - The register state of T2 restored.
 - The **address space remains** the same.

The stack of the relevant thread

- There will be **one stack per thread**



**A Single-Threaded
Address Space**



**Two threaded
Address Space**

Race condition

- Example with two threads
 - counter = counter + 1 (default is 50)
 - We expect the result is 52. However,

OS	Thread1	Thread2	(after instruction)		
			PC	%eax	counter
	before critical section		100	0	50
	mov 0x8049a1c, %eax		105	50	50
	add \$0x1, %eax		108	51	50
interrupt	save T1's state				
	restore T2's state		100	0	50
		mov 0x8049a1c, %eax	105	50	50
		add \$0x1, %eax	108	51	50
		mov %eax, 0x8049a1c	113	51	51
interrupt	save T2's state				
	restore T1's state		108	51	50
	mov %eax, 0x8049a1c		113	51	51

Critical section

- A piece of code that **accesses a shared variable** and must not be concurrently executed by more than one thread
 - Multiple threads executing critical section can result in a race condition
 - Need to support **atomicity** for critical sections (**mutual exclusion**)

Locks

- Ensure that any such critical section executes as if it were a single atomic instruction (**execute a series of instructions atomically**)

```
1  lock_t mutex;  
2  . . .  
3  lock(&mutex);  
4  balance = balance + 1;  
5  unlock(&mutex);
```

→ Critical section

Thread Creation

- How to create and control threads?

```
#include <pthread.h>

int
pthread_create(      pthread_t*      thread,
                    const pthread_attr_t* attr,
                    void*            (*start_routine)(void*),
                    void*            arg);
```

- `thread`: Used to interact with this thread
- `attr`: Used to specify any attributes this thread might have
 - Stack size, Scheduling priority, ...
- `start_routine`: the function this thread start running in
- `arg`: the argument to be passed to the function (`start_routine`)
 - *a void pointer* allows us to pass in *any type of* argument

Thread Creation (Cont.)

- If `start_routine` instead required another type argument, the declaration would look like this:
 - An integer argument:

```
int
pthread_create(..., // first two args are the same
                void*  (*start_routine)(int),
                int    arg);
```

- Return an integer:

```
int
pthread_create(..., // first two args are the same
                int   (*start_routine)(void*),
                void*  arg);
```

Example: Creating a Thread

```
#include <pthread.h>

typedef struct __myarg_t {
    int a;
    int b;
} myarg_t;

void *mythread(void *arg) {
    myarg_t *m = (myarg_t *) arg;
    printf("%d %d\n", m->a, m->b);
    return NULL;
}

int main(int argc, char *argv[]) {
    pthread_t p;
    int rc;

    myarg_t args;
    args.a = 10;
    args.b = 20;
    rc = pthread_create(&p, NULL, mythread, &args);
    ...
}
```

Wait for a thread to complete

```
int pthread_join(pthread_t thread, void **value_ptr);
```

- thread: Specify which thread *to wait for*
- value_ptr: A pointer to the return value
 - Because `pthread_join()` routine changes the value, you need to **pass in a pointer** to that value

Example: Waiting for Thread Completion

```
1  #include <stdio.h>
2  #include <pthread.h>
3  #include <assert.h>
4  #include <stdlib.h>
5
6  typedef struct __myarg_t {
7      int a;
8      int b;
9  } myarg_t;
10
11 typedef struct __myret_t {
12     int x;
13     int y;
14 } myret_t;
15
16 void *mythread(void *arg) {
17     myarg_t *m = (myarg_t *) arg;
18     printf("%d %d\n", m->a, m->b);
19     myret_t *r = malloc(sizeof(myret_t));
20     r->x = 1;
21     r->y = 2;
22     return (void *) r;
23 }
24
```

Example: Waiting for Thread Completion

```
25  int main(int argc, char *argv[]) {
26      int rc;
27      pthread_t p;
28      myret_t *m;
29
30      myarg_t args;
31      args.a = 10;
32      args.b = 20;
33      pthread_create(&p, NULL, mythread, &args);
34      pthread_join(p, (void **) &m); // this thread has been
                                     // waiting inside of the
                                     // pthread_join() routine.
35      printf("returned %d %d\n", m->x, m->y);
36      return 0;
37  }
```

Example: Dangerous code

- Be careful with how values are returned from a thread

```
1  void *mythread(void *arg) {
2      myarg_t *m = (myarg_t *) arg;
3      printf("%d %d\n", m->a, m->b);
4      myret_t r; // ALLOCATED ON STACK: BAD!
5      r.x = 1;
6      r.y = 2;
7      return (void *) &r;
8  }
```

- When the variable `r` returns, it is automatically **de-allocated**

Example: Simpler Argument Passing to a Thread

- Just passing in a single value

```
1  void *mythread(void *arg) {
2      int m = (int) arg;
3      printf("%d\n", m);
4      return (void *) (arg + 1);
5  }
6
7  int main(int argc, char *argv[]) {
8      pthread_t p;
9      int rc, m;
10     pthread_create(&p, NULL, mythread, (void *) 100);
11     pthread_join(p, (void **) &m);
12     printf("returned %d\n", m);
13     return 0;
14 }
```

Locks

- Provide **mutual exclusion** to a critical section

- Interface

```
int pthread_mutex_lock(pthread_mutex_t *mutex);  
int pthread_mutex_unlock(pthread_mutex_t *mutex);
```

- Usage (w/o *lock initialization* and *error check*)

```
pthread_mutex_t lock;  
pthread_mutex_lock(&lock);  
x = x + 1; // or whatever your critical section is  
pthread_mutex_unlock(&lock);
```

- No other thread holds the lock → the thread will acquire the lock and **enter the critical section**
 - If another thread hold the lock → the thread will **not return from the call** until it has acquired the lock

Locks (Cont.)

- All locks must be properly initialized
 - One way: using PTHREAD_MUTEX_INITIALIZER

```
pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER;
```

- The dynamic way: using pthread_mutex_init()

```
int rc = pthread_mutex_init(&lock, NULL);  
assert(rc == 0); // always check success!
```

Locks (Cont.)

- Check errors code when calling lock and unlock

- An example wrapper

```
// Use this to keep your code clean but check for failures
// Only use if exiting program is OK upon failure
void Pthread_mutex_lock(pthread_mutex_t *mutex) {
    int rc = pthread_mutex_lock(mutex);
    assert(rc == 0);
}
```

- These two calls are used in lock acquisition

```
int pthread_mutex_trylock(pthread_mutex_t *mutex);
int pthread_mutex_timelock(pthread_mutex_t *mutex,
                           struct timespec *abs_timeout);
```

- trylock: return failure if the lock is already held
- timelock: return after a timeout

Condition Variables

- **Condition variables** are useful when some kind of **signaling** must take place between threads

```
int pthread_cond_wait(pthread_cond_t *cond,  
                      pthread_mutex_t *mutex);  
int pthread_cond_signal(pthread_cond_t *cond);
```

- `pthread_cond_wait`:
 - Put the calling thread to sleep.
 - Wait for some other thread to signal it.
- `pthread_cond_signal`:
 - Unblock at least one of the threads that are blocked on the condition variable

Condition Variables (Cont.)

- A thread calling wait routine:

```
pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER;
pthread_cond_t init = PTHREAD_COND_INITIALIZER;

pthread_mutex_lock(&lock);
while (initialized == 0)
    pthread_cond_wait(&init, &lock);
pthread_mutex_unlock(&lock);
```

- The wait call **releases the lock** when putting said caller to sleep
- Before returning after being woken, the wait call **re-acquire the lock**

- A thread calling signal routine:

```
pthread_mutex_lock(&lock);
initialized = 1;
pthread_cond_signal(&init);
pthread_mutex_unlock(&lock);
```

Condition Variables (Cont.)

- The waiting thread **re-checks** the condition **in a while loop**, instead of a simple if statement

```
pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER;
pthread_cond_t init = PTHREAD_COND_INITIALIZER;

pthread_mutex_lock(&lock);
while (initialized == 0)
    pthread_cond_wait(&init, &lock);
pthread_mutex_unlock(&lock);
```

- Without rechecking, the waiting thread will continue thinking that the condition has changed *even though it has not*

Condition Variables (Cont.)

- Don't ever to this

- A thread calling wait routine:

```
while(initialized == 0)
    ; // spin
```

- A thread calling signal routine:

```
initialized = 1;
```

- It performs poorly in many cases → just wastes CPU cycles
- It is error prone

Compiling and Running

- To compile them, you must include the header `pthread.h`
 - Explicitly link with the **pthread library**, by adding the `-pthread` flag

```
prompt> gcc -o main main.c -Wall -pthread
```

- For more information,

```
man -k pthread
```