



# Locks

# Locks: The Basic Idea

- Ensure that any **critical section** executes as if it were **a single atomic instruction**
  - An example: the canonical update of a shared variable

```
balance = balance + 1;
```

- Add some code around the critical section

```
1 lock_t mutex; // some globally-allocated lock 'mutex'  
2 ...  
3 lock(&mutex);  
4 balance = balance + 1;  
5 unlock(&mutex);
```

# Locks: The Basic Idea

- Lock variable holds the state of the lock
  - **available** (or **unlocked** or **free**)
    - No thread holds the lock
  - **acquired** (or **locked** or **held**)
    - Exactly one thread holds the lock and presumably is in a critical section

# The semantics of the lock()

- lock()
  - Try to acquire the lock
  - If no other thread holds the lock, the thread will **acquire** the lock
  - **Enter** the *critical section*
    - This thread is said to be the owner of the lock
  - Other threads are *prevented from* entering the critical section while the first thread that holds the lock is in there

# Pthread Locks - mutex

- The name that the POSIX library uses for a lock
  - Used to provide **mutual exclusion** between threads

```
1  pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER;  
2  
3  Pthread_mutex_lock(&lock); // wrapper for pthread_mutex_lock()  
4  balance = balance + 1;  
5  Pthread_mutex_unlock(&lock);
```

- We may be using *different locks* to protect *different variables* → Increase **concurrency** (a more **fine-grained** approach)

# Building A Lock

- Efficient locks provided mutual exclusion at **low cost**
- Building a lock need some help from the **hardware** and the **OS**

# Evaluating locks – Basic criteria

- **Mutual exclusion**

- Does the lock work, preventing multiple threads from entering *a critical section*?

- **Fairness**

- Does each thread contending for the lock get a fair shot at acquiring it once it is free?  
(Starvation)

- **Performance**

- The time overheads added by using the lock

# Controlling Interrupts

- **Disable Interrupts** for critical sections
  - One of the earliest solutions used to provide mutual exclusion
  - Invented for single-processor systems
- Problem:
  - Require too much *trust* in applications
    - » Greedy (or malicious) program could monopolize the processor
  - Do not work on **multiprocessors**
  - Code that masks or unmasks interrupts be executed *slowly* by modern CPUs

```
1 void lock() {  
2     DisableInterrupts();  
3 }  
4 void unlock() {  
5     EnableInterrupts();  
6 }
```

# Why Hardware Support Needed?

- **First attempt:** Using a *flag* denoting whether the lock is held or not
  - The code below has problems

```
1  typedef struct __lock_t { int flag; } lock_t;
2
3  void init(lock_t *mutex) {
4      // 0 → lock is available, 1 → held
5      mutex->flag = 0;
6  }
7
8  void lock(lock_t *mutex) {
9      while (mutex->flag == 1) // TEST the flag
10          ; // spin-wait (do nothing)
11      mutex->flag = 1; // now SET it !
12  }
13
14 void unlock(lock_t *mutex) {
15     mutex->flag = 0;
16 }
```

# Why Hardware Support Needed?

- **Problem 1:** No Mutual Exclusion (assume `flag=0` to begin)

Thread1	Thread2
<pre>call lock() while (flag == 1) interrupt: switch to Thread 2</pre>	<pre>call lock() while (flag == 1) flag = 1; interrupt: switch to Thread 1</pre>

`flag = 1; // set flag to 1 (too!)`

- **Problem 2:** Spin-waiting wastes time waiting for another thread.
- So, we need an atomic instruction supported by **Hardware!**
  - *test-and-set* instruction, also known as *atomic exchange*

# Test And Set (Atomic Exchange)

- An instruction to support the creation of simple locks

```
1 int TestAndSet(int *ptr, int new) {  
2     int old = *ptr;    // fetch old value at ptr  
3     *ptr = new;       // store 'new' into ptr  
4     return old;        // return the old value  
5 }
```

- return**(testing) old value pointed to by the `ptr`
- Simultaneously update*(setting) said value to `new`
- This sequence of operations is **performed atomically**

# A Simple Spin Lock using *test-and-set*

```
1  typedef struct __lock_t {
2      int flag;
3  } lock_t;
4
5  void init(lock_t *lock) {
6      // 0 indicates that lock is available,
7      // 1 that it is held
8      lock->flag = 0;
9  }
10
11 void lock(lock_t *lock) {
12     while (TestAndSet(&lock->flag, 1) == 1)
13         ;           // spin-wait
14 }
15
16 void unlock(lock_t *lock) {
17     lock->flag = 0;
18 }
```

- **Note:** To work correctly on *a single processor*, it requires a preemptive scheduler

# Evaluating Spin Locks

- **Correctness:** yes
  - The spin lock only allows a single thread to entry the critical section
- **Fairness:** no
  - Spin locks don't provide any fairness guarantees
  - Indeed, a thread spinning may spin *forever*
- **Performance:**
  - In the single CPU, performance overheads can be quite *painful*
  - If the number of threads roughly equals the number of CPUs, spin locks work *reasonably well*

# Compare-And-Swap

- Test whether the value at the address(ptr) is equal to expected
  - If so, update the memory location pointed to by ptr with the new value
  - In either case, return the actual value at that memory location

```
1 int CompareAndSwap(int *ptr, int expected, int new) {  
2     int actual = *ptr;  
3     if (actual == expected)  
4         *ptr = new;  
5     return actual;  
6 }
```

Compare-and-Swap hardware atomic instruction (C-style)

```
1 void lock(lock_t *lock) {  
2     while (CompareAndSwap(&lock->flag, 0, 1) == 1)  
3         ; // spin  
4 }
```

Spin lock with compare-and-swap

# Compare-And-Swap

- C-callable x86-version of compare-and-swap

```
1  char CompareAndSwap(int *ptr, int old, int new) {
2      unsigned char ret;
3
4      // Note that sete sets a 'byte' not the word
5      __asm__ __volatile__ (
6          "lock\n"
7          "cmpxchgl %2,%1\n"
8          "sete %0\n"
9          : "=q" (ret), "=m" (*ptr)
10         : "r" (new), "m" (*ptr), "a" (old)
11         : "memory");
12
13 }
```

# Load-Linked and Store-Conditional

```
1 int LoadLinked(int *ptr) {
2     return *ptr;
3 }
4
5 int StoreConditional(int *ptr, int value) {
6     if (no one has updated *ptr since the LoadLinked to this address) {
7         *ptr = value;
8         return 1; // success!
9     } else {
10        return 0; // failed to update
11    }
12 }
```

## Load-linked And Store-conditional

- The store-conditional *only succeeds* if **no intermittent store** to the address has taken place
  - **success**: return 1 and update the value at `ptr` to `value`
  - **fail**: the value at `ptr` is not updates and 0 is returned

# Load-Linked and Store-Conditional

```
1 void lock(lock_t *lock) {
2     while (1) {
3         while (LoadLinked(&lock->flag) == 1)
4             ; // spin until it's zero
5         if (StoreConditional(&lock->flag, 1) == 1)
6             return; // if set-it-to-1 was a success: all done
7             // otherwise: try it all over again
8     }
9 }
10
11 void unlock(lock_t *lock) {
12     lock->flag = 0;
13 }
```

## Using LL/SC To Build A Lock

```
1 void lock(lock_t *lock) {
2     while (LoadLinked(&lock->flag) || !StoreConditional(&lock->flag, 1))
3         ; // spin
4 }
```

## A more concise form of the `lock()` using LL/SC

# Fetch-And-Add

- Atomically increment a value while returning the old value at a particular address

```
1 int FetchAndAdd(int *ptr) {  
2     int old = *ptr;  
3     *ptr = old + 1;  
4     return old;  
5 }
```

**Fetch-And-Add Hardware atomic instruction (C-style)**

# Ticket Lock

- Ticket lock can be built with fetch-and add
  - Ensure progress for all threads → fairness

```
1  typedef struct __lock_t {  
2      int ticket;  
3      int turn;  
4  } lock_t;  
5  
6  void lock_init(lock_t *lock) {  
7      lock->ticket = 0;  
8      lock->turn = 0;  
9  }  
10  
11 void lock(lock_t *lock) {  
12     int myturn = FetchAndAdd(&lock->ticket);  
13     while (lock->turn != myturn)  
14         ; // spin  
15 }  
16 void unlock(lock_t *lock) {  
17     FetchAndAdd(&lock->turn);  
18 }
```

# So Much Spinning

- Hardware-based spin locks are **simple** and they work
- In some cases, these solutions can be quite **inefficient**
  - Any time a thread gets caught *spinning*, it **wastes an entire time slice** doing nothing but checking a value

How To Avoid *Spinning*?  
We'll need **OS Support** too!

# A Simple Approach: Just Yield

- When you are going to spin, **give up the CPU** to another thread
  - OS system call moves the caller from the *running state* to the *ready state*
  - The cost of a **context switch** can be substantial and the **starvation** problem still exists

```
1 void init() {
2     flag = 0;
3 }
4
5 void lock() {
6     while (TestAndSet(&flag, 1) == 1)
7         yield(); // give up the CPU
8 }
9
10 void unlock() {
11     flag = 0;
12 }
```

Lock with Test-and-set and Yield

# Using Queues: Sleeping Instead of Spinning

- **Queue** to keep track of which threads are waiting to enter the lock
- `park()`
  - Put a calling thread to sleep
- `unpark(threadID)`
  - Wake up a particular thread as designated by `threadID`

# Using Queues: Sleeping Instead of Spinning

```
1  typedef struct __lock_t { int flag; int guard; queue_t *q; } lock_t;
2
3  void lock_init(lock_t *m) {
4      m->flag = 0;
5      m->guard = 0;
6      queue_init(m->q);
7  }
8
9  void lock(lock_t *m) {
10     while (TestAndSet(&m->guard, 1) == 1)
11         ; // acquire guard lock by spinning
12     if (m->flag == 0) {
13         m->flag = 1; // lock is acquired
14         m->guard = 0;
15     } else {
16         queue_add(m->q, gettid());
17         m->guard = 0;
18         park();
19     }
20 }
21 ...
```

**Lock With Queues, Test-and-set, Yield, And Wakeup**

# Using Queues: Sleeping Instead of Spinning

```
22 void unlock(lock_t *m) {  
23     while (TestAndSet(&m->guard, 1) == 1)  
24         ; // acquire guard lock by spinning  
25     if (queue_empty(m->q))  
26         m->flag = 0; // let go of lock; no one wants it  
27     else  
28         unpark(queue_remove(m->q)); // hold lock (for next thread!)  
29     m->guard = 0;  
30 }
```

**Lock With Queues, Test-and-set, Yield, And Wakeup (Cont.)**

# Wakeup/Waiting Race

- In case of releasing the lock (*thread A*) just before the call to park() (*thread B*) → Thread B would **sleep forever** (potentially)
- **Solaris** solves this problem by adding a third system call: setpark()
  - By calling this routine, a thread can indicate it *is about to* park
  - If it happens to be interrupted and another thread calls unpark before park is actually called, the subsequent park returns immediately instead of sleeping

```
1      queue_add(m->q, gettid());
2      setpark(); // new code
3      m->guard = 0;
4      park();
```

Code modification inside of lock()

# Futex

- Linux provides a **futex** (is similar to Solaris's park and unpark)
  - `futex_wait(address, expected)`
    - Put the calling thread to sleep
    - If the value at address is not equal to expected, the call returns immediately
  - `futex_wake(address)`
    - Wake one thread that is waiting on the queue

# Futex (Cont.)

- Snippet from `lowlevellock.h` in the **nptl** library
  - The high bit of the integer `v`: track whether the lock is held or not
  - All the other bits : the number of waiters

```
1 void mutex_lock(int *mutex) {
2     int v;
3     /* Bit 31 was clear, we got the mutex (this is the fastpath) */
4     if (atomic_bit_test_set(mutex, 31) == 0)
5         return;
6     atomic_increment(mutex);
7     while (1) {
8         if (atomic_bit_test_set(mutex, 31) == 0) {
9             atomic_decrement(mutex);
10            return;
11        }
12        /* We have to wait now. First make sure the futex value
13           we are monitoring is truly negative (i.e. locked). */
14        v = *mutex;
15        ...
```

## Linux-based Futex Locks

# Futex (Cont.)

```
16         if  (v >= 0)
17             continue;
18         futex_wait(mutex, v);
19     }
20 }
21
22 void mutex_unlock(int *mutex) {
23     /* Adding 0x80000000 to the counter results in 0 if and only if
24      there are not other interested threads */
25     if (atomic_add_zero(mutex, 0x80000000))
26         return;
27     /* There are other threads waiting for this mutex,
28      wake one of them up */
29     futex_wake(mutex);
30 }
```

## Linux-based Futex Locks (Cont.)

# Two-Phase Locks

- A two-phase lock realizes that **spinning can be useful** if the lock *is about to* be released
  - **First phase**
    - The lock spins for a while, *hoping that* it can acquire the lock
    - If the lock is not acquired during the first spin phase, a second phase is entered
  - **Second phase**
    - The caller is put to sleep
    - The caller is only woken up when the lock becomes free later