



# Conditional Variables

# Condition Variables

- There are many cases where a thread wishes to check whether a **condition** is true before continuing its execution
- Example
  - A parent thread might wish to check whether a child thread has *completed*
  - This is often called a `join()`

# Condition Variables (Cont.)



## A Parent Waiting For Its Child

```
1      void *child(void *arg) {
2          printf("child\n");
3          // XXX how to indicate we are done?
4          return NULL;
5      }
6
7      int main(int argc, char *argv[]) {
8          printf("parent: begin\n");
9          pthread_t c;
10         Pthread_create(&c, NULL, child, NULL); // create child
11         // XXX how to wait for child?
12         printf("parent: end\n");
13         return 0;
14     }
```

## What we would like to see here is:

```
parent: begin
child
parent: end
```

# Parent Waiting for Child: Spin-based Approach

```
1     volatile int done = 0;
2
3     void *child(void *arg) {
4         printf("child\n");
5         done = 1;
6         return NULL;
7     }
8
9     int main(int argc, char *argv[]) {
10        printf("parent: begin\n");
11        pthread_t c;
12        pthread_create(&c, NULL, child, NULL); // create child
13        while (done == 0)
14            ; // spin
15        printf("parent: end\n");
16        return 0;
17    }
```

- This is hugely inefficient as the parent spins and **wastes CPU time**

# How to Wait for a Condition

- Condition variable
  - **Waiting** on the condition
    - An explicit queue that threads can put themselves on when some state of execution is not as desired
  - **Signaling** on the condition
    - Some other thread, *when it changes the state*, can wake one of those waiting threads and allow them to continue

# Definition and Routines

- Declare condition variable

```
pthread_cond_t c;
```

- Proper initialization is required

- Operation (the POSIX calls)

```
pthread_cond_wait(pthread_cond_t *c, pthread_mutex_t *m);    // wait()
pthread_cond_signal(pthread_cond_t *c);                         // signal()
```

- The `wait()` call takes a mutex as a parameter
    - The `wait()` call release the lock and put the calling thread to sleep
    - When the thread wakes up, it must re-acquire the lock

# Parent Waiting for Child: Using a condition variable

```
1     int done = 0;
2     pthread_mutex_t m = PTHREAD_MUTEX_INITIALIZER;
3     pthread_cond_t c = PTHREAD_COND_INITIALIZER;
4
5     void thr_exit() {
6         pthread_mutex_lock(&m);
7         done = 1;
8         pthread_cond_signal(&c);
9         pthread_mutex_unlock(&m);
10    }
11
12    void *child(void *arg) {
13        printf("child\n");
14        thr_exit();
15        return NULL;
16    }
17
18    void thr_join() {
19        pthread_mutex_lock(&m);
20        while (done == 0)
21            pthread_cond_wait(&c, &m);
22        pthread_mutex_unlock(&m);
23    }
24
```

# Parent waiting for Child: Using a Condition Variable

```
(cont.)  
25     int main(int argc, char *argv[]) {  
26         printf("parent: begin\n");  
27         pthread_t p;  
28         pthread_create(&p, NULL, child, NULL);  
29         thr_join();  
30         printf("parent: end\n");  
31         return 0;  
32     }
```

# Parent waiting for Child: Using a Condition Variable

- **Parent:**

- Create the child thread and continues running itself
- Call into `thr_join()` to wait for the child thread to complete
  - Acquire the lock
  - Check if the child is done
  - Put itself to sleep by calling `wait()`
  - Release the lock

- **Child:**

- Print the message "child"
- Call `thr_exit()` to wake the parent thread
  - Grab the lock
  - Set the state variable `done`
  - Signal the parent thus waking it

# The Importance of the State Variable done

```
1     void thr_exit() {
2             Pthread_mutex_lock(&m);
3             Pthread_cond_signal(&c);
4             Pthread_mutex_unlock(&m);
5     }
6
7     void thr_join() {
8             Pthread_mutex_lock(&m);
9             Pthread_cond_wait(&c, &m);
10            Pthread_mutex_unlock(&m);
11    }
```

**thr\_exit() and thr\_join() without variable done**

- Imagine the case where the *child runs immediately*
  - The child will signal, but there is no thread asleep on the condition
  - When the parent runs, it will call wait and be stuck
  - No thread will ever wake it

# Another Poor Implementation

```
1     void thr_exit() {
2         done = 1;
3         Pthread_cond_signal(&c);
4     }
5
6     void thr_join() {
7         if (done == 0)
8             Pthread_cond_wait(&c);
9     }
```

- The issue here is a subtle **race condition**
  - The parent calls `thr_join()`
    - The parent checks the value of `done`
    - It will see that it is 0 and try to go to sleep
    - *Just before* it calls wait to go to sleep, the parent is interrupted and the child runs
  - The child changes the state variable `done` to 1 and signals
    - But no thread is waiting and thus no thread is woken
    - When the parent runs again, it sleeps forever

# The Producer / Consumer (Bound Buffer) Problem

- **Producer**
  - Produce data items
  - Wish to place data items in a buffer
- **Consumer**
  - Grab data items out of the buffer consume them in some way
- Example: Multi-threaded web server
  - *A producer* puts HTTP requests in to a work queue
  - *Consumer threads* take requests out of this queue and process them

# Bounded Buffer

- A bounded buffer is used when you pipe the output of one program into another
  - Example: `grep foo file.txt | wc -l`
    - The grep process is the producer
    - The wc process is the consumer
    - Between them is an in-kernel bounded buffer
  - Bounded buffer is Shared resource  
→ **Synchronized access** is required

# Put and Get Routines (Version 1)

```
1      int buffer;
2      int count = 0;    // initially, empty
3
4      void put(int value) {
5          assert(count == 0);
6          count = 1;
7          buffer = value;
8      }
9
10     int get() {
11         assert(count == 1);
12         count = 0;
13         return buffer;
14     }
```

- Only put data into the buffer when `count` is zero
  - i.e., when the buffer is *empty*
- Only get data from the buffer when `count` is one
  - i.e., when the buffer is *full*

# Producer/Consumer Threads (Version 1)

```
1      void *producer(void *arg) {
2          int i;
3          int loops = (int) arg;
4          for (i = 0; i < loops; i++) {
5              put(i);
6          }
7      }
8
9      void *consumer(void *arg) {
10         int i;
11         while (1) {
12             int tmp = get();
13             printf("%d\n", tmp);
14         }
15     }
```

- **Producer** puts an integer into the shared buffer loops number of times
- **Consumer** gets the data out of that shared buffer

# Producer/Consumer: Single CV and If Statement

- A single condition variable cond and associated lock mutex

```
1      cond_t cond;
2      mutex_t mutex;
3
4      void *producer(void *arg) {
5          int i;
6          for (i = 0; i < loops; i++) {
7              Pthread_mutex_lock(&mutex);           // p1
8              if (count == 1)                      // p2
9                  Pthread_cond_wait(&cond, &mutex); // p3
10             put(i);                         // p4
11             Pthread_cond_signal(&cond);       // p5
12             Pthread_mutex_unlock(&mutex);     // p6
13         }
14     }
15
16     void *consumer(void *arg) {
17         int i;
18         for (i = 0; i < loops; i++) {
19             Pthread_mutex_lock(&mutex);           // c1
```

# Producer/Consumer: Single CV and If Statement

```
20         if (count == 0)                                // c2
21             Pthread_cond_wait(&cond, &mutex);          // c3
22         int tmp = get();                             // c4
23         Pthread_cond_signal(&cond);                // c5
24         Pthread_mutex_unlock(&mutex);              // c6
25         printf("%d\n", tmp);
26     }
27 }
```

- p1-p3: A producer waits for the buffer to be empty
- c1-c3: A consumer waits for the buffer to be full
- With just *a single producer* and *a single consumer*, the code works

If we have **more than** one of producer and consumer?

# Thread Trace: Broken Solution (Version 1)

$T_{c1}$	State	$T_{c2}$	State	$T_p$	State	Count	Comment
c1	Running		Ready		Ready	0	
c2	Running		Ready		Ready	0	
c3	Sleep		Ready		Ready	0	Nothing to get
	Sleep		Ready	p1	Running	0	
	Sleep		Ready	p2	Running	0	
	Sleep		Ready	p4	Running	1	Buffer now full
	Ready		Ready	p5	Running	1	$T_{c1}$ awoken
	Ready		Ready	p6	Running	1	
	Ready		Ready	p1	Running	1	
	Ready		Ready	p2	Running	1	
	Ready		Ready	p3	Sleep	1	Buffer full; sleep
	Ready	c1	Running		Sleep	1	$T_{c2}$ sneaks in ...
c4	Ready	c2	Running		Sleep	1	
	Ready	c4	Running		Sleep	0	... and grabs data
	Ready	c5	Running		Ready	0	$T_p$ awoken
	Ready	c6	Running		Ready	0	
	Running		Ready		Ready	0	<b>Oh oh! No data</b>

# Thread Trace: Broken Solution (Version 1)

- The problem arises for a simple reason:
  - After the producer woke  $T_{c1}$ , but before  $T_{c1}$  ever ran, the state of the bounded buffer *changed by  $T_{c2}$*
  - There is no guarantee that when the woken thread runs, the state will still be as desired  
→ Mesa semantics
    - Virtually every system ever built employs *Mesa semantics*
  - Hoare semantics provides a stronger guarantee that the woken thread will run immediately upon being woken

# Producer/Consumer: Single CV and While

- Consumer  $T_{c1}$  wakes up and **re-checks** the state of the shared variable
  - If the buffer is empty, the consumer simply goes back to sleep

```
1      cond_t cond;
2      mutex_t mutex;
3
4      void *producer(void *arg) {
5          int i;
6          for (i = 0; i < loops; i++) {
7              Pthread_mutex_lock(&mutex);           // p1
8              while (count == 1)                  // p2
9                  Pthread_cond_wait(&cond, &mutex); // p3
10             put(i);
11             Pthread_cond_signal(&cond);        // p4
12             Pthread_mutex_unlock(&mutex);     // p5
13         }
14     }
15 }
```

# Producer/Consumer: Single CV and While

(Cont.)

```
16     void *consumer(void *arg) {
17         int i;
18         for (i = 0; i < loops; i++) {
19             Pthread_mutex_lock(&mutex); // c1
20             while (count == 0) // c2
21                 Pthread_cond_wait(&cond, &mutex); // c3
22             int tmp = get(); // c4
23             Pthread_cond_signal(&cond); // c5
24             Pthread_mutex_unlock(&mutex); // c6
25             printf("%d\n", tmp);
26         }
27     }
```

- A simple rule to remember with condition variables is to **always use while loops**
- However, this code still has a bug (*next page*)

# Thread Trace: Broken Solution (Version 2)

$T_{c1}$	State	$T_{c2}$	State	$T_p$	State	Count	Comment
c1	Running		Ready		Ready	0	
c2	Running		Ready		Ready	0	
c3	Sleep		Ready		Ready	0	Nothing to get
	Sleep	c1	Running		Ready	0	
	Sleep	c2	Running		Ready	0	
	Sleep	c3	Sleep		Ready	0	Nothing to get
	Sleep		Sleep	p1	Running	0	
	Sleep		Sleep	p2	Running	0	
	Sleep		Sleep	p4	Running	1	Buffer now full
	Ready		Sleep	p5	Running	1	$T_{c1}$ awoken
	Ready		Sleep	p6	Running	1	
	Ready		Sleep	p1	Running	1	
	Ready		Sleep	p2	Running	1	
	Ready		Sleep	p3	Sleep	1	Must sleep (full)
c2	Running		Sleep		Sleep	1	Recheck condition
c4	Running		Sleep		Sleep	0	$T_{c1}$ grabs data
c5	Running		Ready		Sleep	0	<b>Oops! Woke <math>T_{c2}</math></b>

# Thread Trace: Broken Solution (Version 2)

$T_{c1}$	State	$T_{c2}$	State	$T_p$	State	Count	Comment
...	...	...	...	...	...	...	(cont.)
c6	Running		Ready		Sleep	0	
c1	Running		Ready		Sleep	0	
c2	Running		Ready		Sleep	0	
c3	Sleep		Ready		Sleep	0	Nothing to get
	Sleep	c2	Running		Sleep	0	
	Sleep	c3	Sleep		Sleep	0	<b>Everyone asleep ...</b>

- A consumer should not wake other consumers, only producers, and vice-versa

# Single Buffer Producer/Consumer Solution

- Use **two** condition variables and while
  - **Producer** threads wait on the condition `empty`, and signals `fill`
  - **Consumer** threads wait on `fill` and signal `empty`

```
1      cond_t empty, fill;
2      mutex_t mutex;
3
4      void *producer(void *arg) {
5          int i;
6          for (i = 0; i < loops; i++) {
7              Pthread_mutex_lock(&mutex);
8              while (count == 1)
9                  Pthread_cond_wait(&empty, &mutex);
10             put(i);
11             Pthread_cond_signal(&fill);
12             Pthread_mutex_unlock(&mutex);
13         }
14     }
```

# Single Buffer Producer/Consumer Solution

```
(Cont.)  
16     void *consumer(void *arg) {  
17         int i;  
18         for (i = 0; i < loops; i++) {  
19             Pthread_mutex_lock(&mutex);  
20             while (count == 0)  
21                 Pthread_cond_wait(&fill, &mutex);  
22             int tmp = get();  
23             Pthread_cond_signal(&empty);  
24             Pthread_mutex_unlock(&mutex);  
25             printf("%d\n", tmp);  
26         }  
27     }
```

# Final Producer/Consumer Solution

- More concurrency and efficiency → Add more buffer slots
  - Allow concurrent production or consuming to take place
  - Reduce context switches

```
1      int buffer[MAX];
2      int fill = 0;
3      int use = 0;
4      int count = 0;
5
6      void put(int value) {
7          buffer[fill] = value;
8          fill = (fill + 1) % MAX;
9          count++;
10     }
11
12     int get() {
13         int tmp = buffer[use];
14         use = (use + 1) % MAX;
15         count--;
16         return tmp;
17     }
```

The Final Put and Get Routines

# Final Producer/Consumer Solution

```
1      cond_t empty, fill;
2      mutex_t mutex;
3
4      void *producer(void *arg) {
5          int i;
6          for (i = 0; i < loops; i++) {
7              Pthread_mutex_lock(&mutex); // p1
8              while (count == MAX) // p2
9                  Pthread_cond_wait(&empty, &mutex); // p3
10             put(i);
11             Pthread_cond_signal(&fill); // p4
12             Pthread_mutex_unlock(&mutex); // p5
13         } // p6
14     }
15
16     void *consumer(void *arg) {
17         int i;
18         for (i = 0; i < loops; i++) {
19             Pthread_mutex_lock(&mutex); // c1
20             while (count == 0) // c2
21                 Pthread_cond_wait(&fill, &mutex); // c3
22             int tmp = get(); // c4
```

# Final Producer/Consumer Solution

(Cont.)

```
23             Pthread_cond_signal (&empty) ;           // c5  
24             Pthread_mutex_unlock (&mutex) ;         // c6  
25             printf ("%d\n", tmp) ;  
26         }  
27     }
```

The Final Working Solution (Cont.)

- p2: A **producer** only sleeps if all buffers are currently filled
- c2: A **consumer** only sleeps if all buffers are currently empty

# Covering Conditions

- Assume there are zero bytes free
  - Thread  $T_a$  calls `allocate(100)`
  - Thread  $T_b$  calls `allocate(10)`
  - Both  $T_a$  and  $T_b$  wait on the condition and go to sleep
  - Thread  $T_c$  calls `free(50)`

**Which waiting thread should be woken up?**

# Covering Conditions

```
1      // how many bytes of the heap are free?
2      int bytesLeft = MAX_HEAP_SIZE;
3
4      // need lock and condition too
5      cond_t c;
6      mutex_t m;
7
8      void *
9      allocate(int size) {
10         Pthread_mutex_lock(&m);
11         while (bytesLeft < size)
12             Pthread_cond_wait(&c, &m);
13         void *ptr = ...;           // get mem from heap
14         bytesLeft -= size;
15         Pthread_mutex_unlock(&m);
16         return ptr;
17     }
18
19     void free(void *ptr, int size) {
20         Pthread_mutex_lock(&m);
21         bytesLeft += size;
22         Pthread_cond_signal(&c);    // whom to signal??
23         Pthread_mutex_unlock(&m);
24     }
```

# Covering Conditions

- Solution (Suggested by Lampson and Redell)
  - Replace `pthread_cond_signal()` with `pthread_cond_broadcast()`
  - `pthread_cond_broadcast()`
    - Wake up **all waiting threads**
    - Cost: too many threads might be woken
    - Threads that shouldn't be awake will simply wake up, re-check the condition, and then go back to sleep