



Semaphores, Advanced Locks, and Synchronization Problems

Semaphore Definition

- An object **with an integer value**
 - We can manipulate with two routines; `sem_wait()` and `sem_post()`
 - Initialization

```
1 #include <semaphore.h>
2 sem_t s;
3 sem_init(&s, 0, 1); // initialize s to the value 1
```

- Declare a semaphore `s` and initialize it to the value 1
- The second argument, 0, indicates that the semaphore is shared between *threads in the same process*

Semaphore Interfaces

■ sem_wait()

```
1 int sem_wait(sem_t *s) {  
2     decrement the value of semaphore s by one  
3     wait if value of semaphore s is negative  
4 }
```

- If the value of the semaphore was *one* or *higher* when called `sem_wait()`, **return right away**
- It will cause the caller to suspend execution waiting for a subsequent post
- When negative, the value of the semaphore is equal to the number of waiting threads

Semaphore Interfaces



■ sem_post()

```
1 int sem_post(sem_t *s) {  
2     increment the value of semaphore s by one  
3     if there are one or more threads waiting, wake one  
4 }
```

- Simply **increments** the value of the semaphore
- If there is a thread waiting to be woken, **wakes** one of them up

Binary Semaphores (Locks)

- What should **x** be?
 - The initial value should be **1**

```
1 sem_t m;
2 sem_init(&m, 0, X); // initialize semaphore to X; what should X be?
3
4 sem_wait(&m);
5 //critical section here
6 sem_post(&m);
```

Single Thread Using a Semaphore

Value of Semaphore	Thread 0	Thread 1
1		
1	call sema_wait()	
0	sem_wait() returns	
0	(crit sect)	
0	call sem_post()	
1	sem_post() returns	

Two Threads Using A Semaphore

Value	Thread 0	State	Thread 1	State
1		Running		Ready
1	call sem_wait()	Running		Ready
0	sem_wait() returns	Running		Ready
0	(crit set: begin)	Running		Ready
0	<i>Interrupt; Switch → T1</i>	Ready		Running
0		Ready	call sem_wait()	Running
-1		Ready	decrement sem	Running
-1		Ready	(sem < 0) → sleep	sleeping
-1		Running	<i>Switch → T0</i>	sleeping
-1	(crit sect: end)	Running		sleeping
-1	call sem_post()	Running		sleeping
0	increment sem	Running		sleeping
0	wake(T1)	Running		Ready
0	sem_post() returns	Running		Ready
0	<i>Interrupt; Switch → T1</i>	Ready		Running
0		Ready	sem_wait() returns	Running
0		Ready	(crit sect)	Running
0		Ready	call sem_post()	Running
1		Ready	sem_post() returns	Running

Semaphores As Condition Variables

```
1  sem_t s;
2
3  void *
4  child(void *arg) {
5      printf("child\n");
6      sem_post(&s); // signal here: child is done
7      return NULL;
8 }
9
10 int
11 main(int argc, char *argv[]) {
12     sem_init(&s, 0, X); // what should X be?
13     printf("parent: begin\n");
14     pthread_t c;
15     pthread_create(c, NULL, child, NULL);
16     sem_wait(&s); // wait here for child
17     printf("parent: end\n");
18     return 0;
19 }
```

parent: begin
child
parent: end

The execution result

A Parent Waiting For Its Child

- What should **x** be?
 - The value of semaphore should be set to is **0**

Parent Waiting For Child (Case 1)

- The parent call `sem_wait()` before the child has called `sem_post()`

Value	Parent	State	Child	State
0	Create(Child)	Running	(Child exists; is runnable)	Ready
0	call <code>sem_wait()</code>	Running		Ready
-1	decrement sem	Running		Ready
-1	$(\text{sem} < 0) \rightarrow \text{sleep}$	sleeping		Ready
-1	<i>Switch→Child</i>	sleeping	child runs	Running
-1		sleeping	call <code>sem_post()</code>	Running
0		sleeping	increment sem	Running
0		Ready	wake (Parent)	Running
0		Ready	<code>sem_post()</code> returns	Running
0		Ready	<i>Interrupt; Switch→Parent</i>	Ready
0	<code>sem_wait()</code> retruns	Running		Ready

Parent Waiting For Child (Case 2)

- The child runs to completion before the parent call `sem_wait()`

Value	Parent	State	Child	State
0	Create(Child)	Running	(Child exists; is runnable)	Ready
0	<i>Interrupt; switch→Child</i>	Ready	child runs	Running
0		Ready	call <code>sem_post()</code>	Running
1		Ready	increment sem	Running
1		Ready	wake (nobody)	Running
1		Ready	<code>sem_post()</code> returns	Running
1	parent runs	Running	<i>Interrupt; Switch→Parent</i>	Ready
1	call <code>sem_wait()</code>	Running		Ready
0	decrement sem	Running		Ready
0	(sem<0)→awake	Running		Ready
0	<code>sem_wait()</code> retruns	Running		Ready

The Producer/Consumer (Bounded-Buffer) Problem

- **Producer:** `put ()` interface
 - Wait for a buffer to become *empty* in order to put data into it
- **Consumer:** `get ()` interface
 - Wait for a buffer to become *filled* before using it

```
1  int buffer[MAX];
2  int fill = 0;
3  int use = 0;
4
5  void put(int value) {
6      buffer[fill] = value;      // line f1
7      fill = (fill + 1) % MAX; // line f2
8  }
9
10 int get() {
11     int tmp = buffer[use];    // line g1
12     use = (use + 1) % MAX;   // line g2
13     return tmp;
14 }
```

The Producer/Consumer (Bounded-Buffer) Problem

```
1  sem_t empty;
2  sem_t full;
3
4  void *producer(void *arg) {
5      int i;
6      for (i = 0; i < loops; i++) {
7          sem_wait(&empty);           // line P1
8          put(i);                   // line P2
9          sem_post(&full);         // line P3
10     }
11 }
12
13 void *consumer(void *arg) {
14     int i, tmp = 0;
15     while (tmp != -1) {
16         sem_wait(&full);        // line C1
17         tmp = get();            // line C2
18         sem_post(&empty);       // line C3
19         printf("%d\n", tmp);
20     }
21 }
22 ...
```

First Attempt: Adding the Full and Empty Conditions

The Producer/Consumer (Bounded-Buffer) Problem

```
21 int main(int argc, char *argv[]) {  
22     // ...  
23     sem_init(&empty, 0, MAX);           // MAX buffers are empty to begin with...  
24     sem_init(&full, 0, 0);             // ... and 0 are full  
25     // ...  
26 }
```

First Attempt: Adding the Full and Empty Conditions (Cont.)

- Imagine that `MAX` is greater than 1
 - If there are multiple producers, **race condition** can happen at line p1
 - It means that the old data may be overwritten
- We've forgotten **mutual exclusion**
 - Filling of a buffer and incrementing of the index into the buffer is in a **critical section**

Adding Mutual Exclusion

```
1 sem_t empty;
2 sem_t full;
3 sem_t mutex;
4
5 void *producer(void *arg) {
6     int i;
7     for (i = 0; i < loops; i++) {
8         sem_wait(&mutex);           // line p0 (NEW LINE)
9         sem_wait(&empty);          // line p1
10        put(i);                  // line p2
11        sem_post(&full);          // line p3
12        sem_post(&mutex);          // line p4 (NEW LINE)
13    }
14 }
15
(Cont.)
```

Adding Mutual Exclusion (Incorrectly)

Adding Mutual Exclusion

```
(Cont.)  
16 void *consumer(void *arg) {  
17     int i;  
18     for (i = 0; i < loops; i++) {  
19         sem_wait(&mutex);           // line c0 (NEW LINE)  
20         sem_wait(&full);          // line c1  
21         int tmp = get();          // line c2  
22         sem_post(&empty);          // line c3  
23         sem_post(&mutex);          // line c4 (NEW LINE)  
24         printf("%d\n", tmp);  
25     }  
26 }
```

Adding Mutual Exclusion (Incorrectly)

Adding Mutual Exclusion

- Imagine two threads: one producer and one consumer
 - The consumer **acquire** the mutex (line c0)
 - The consumer **calls** `sem_wait()` on the full semaphore (line c1)
 - The consumer is **blocked** and **yield** the CPU
 - The consumer still holds the mutex!
 - The producer **calls** `sem_wait()` on the binary mutex semaphore (line p0)
 - The producer is now **stuck** waiting too
 - **This is a classic deadlock situation**

Finally, A Working Solution

```
1 sem_t empty;
2 sem_t full;
3 sem_t mutex;
4
5 void *producer(void *arg) {
6     int i;
7     for (i = 0; i < loops; i++) {
8         sem_wait(&empty);           // line p1
9         sem_wait(&mutex);          // line p1.5 (MOVED MUTEX HERE...)
10        put(i);                  // line p2
11        sem_post(&mutex);         // line p2.5 (... AND HERE)
12        sem_post(&full);          // line p3
13    }
14 }
15
(Cont.)
```

Adding Mutual Exclusion (Correctly)

Finally, A Working Solution

```
(Cont.)  
16 void *consumer(void *arg) {  
17     int i;  
18     for (i = 0; i < loops; i++) {  
19         sem_wait(&full);           // line c1  
20         sem_wait(&mutex);        // line c1.5 (MOVED MUTEX HERE...)  
21         int tmp = get();         // line c2  
22         sem_post(&mutex);        // line c2.5 (... AND HERE)  
23         sem_post(&empty);        // line c3  
24         printf("%d\n", tmp);  
25     }  
26 }  
27  
28 int main(int argc, char *argv[]) {  
29     // ...  
30     sem_init(&empty, 0, MAX); // MAX buffers are empty to begin with ...  
31     sem_init(&full, 0, 0);    // ... and 0 are full  
32     sem_init(&mutex, 0, 1);  // mutex=1 because it is a lock  
33     // ...  
34 }
```

Adding Mutual Exclusion (Correctly)

Reader-Writer Locks

- Imagine a number of concurrent list operations, including **inserts** and simple **lookups**
 - **insert:**
 - Change the state of the list
 - A traditional critical section makes sense
 - **lookup:**
 - Simply *read* the data structure
 - As long as we can guarantee that no insert is on-going, we can allow many lookups to proceed **concurrently**

This special type of lock is known as a **reader-write lock**

A Reader-Writer Locks

- Only a **single writer** can acquire the lock
- Once a reader has acquired a **read lock**,
 - **More readers** will be allowed to acquire the read lock too
 - A writer will have to wait until all readers are finished

```
1  typedef struct _rwlock_t {
2      sem_t lock;          // binary semaphore (basic lock)
3      sem_t writelock;    // used to allow ONE writer or MANY readers
4      int readers;        // count of readers reading in critical section
5  } rwlock_t;
6
7  void rwlock_init(rwlock_t *rw) {
8      rw->readers = 0;
9      sem_init(&rw->lock, 0, 1);
10     sem_init(&rw->writelock, 0, 1);
11 }
12
13 void rwlock_acquire_readlock(rwlock_t *rw) {
14     sem_wait(&rw->lock);
15     ...
16 }
```

A Reader-Writer Locks

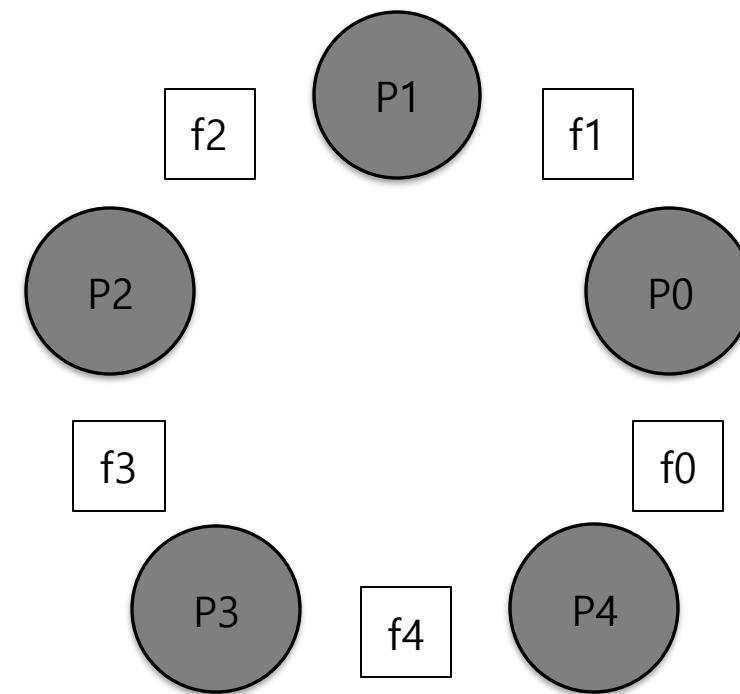
```
15     rw->readers++;
16     if (rw->readers == 1)
17         sem_wait(&rw->writelock); // first reader acquires writelock
18     sem_post(&rw->lock);
19 }
20
21 void rwlock_release_readlock(rwlock_t *rw) {
22     sem_wait(&rw->lock);
23     rw->readers--;
24     if (rw->readers == 0)
25         sem_post(&rw->writelock); // last reader releases writelock
26     sem_post(&rw->lock);
27 }
28
29 void rwlock_acquire_writelock(rwlock_t *rw) {
30     sem_wait(&rw->writelock);
31 }
32
33 void rwlock_release_writelock(rwlock_t *rw) {
34     sem_post(&rw->writelock);
35 }
```

A Reader-Writer Locks

- The reader-writer locks have **fairness problem**
 - It would be relatively easy for reader to **starve writer**
 - How to prevent more readers from entering the lock once a writer is waiting?

The Dining Philosophers

- Assume there are five “**philosophers**” sitting around a table
 - Between each pair of philosophers is a single fork (five total)
 - The philosophers each have times where they **think**, and don’t need any forks, and times where they **eat**
 - In order to *eat*, a philosopher needs **two forks**, both the one on their *left* and the one on their *right*
 - The contention for these forks**



The Dining Philosophers

- Key challenge

- There is **no deadlock**
- **No** philosopher **starves** and never gets to eat
- **Concurrency** is high

```
while (1) {  
    think();  
    getforks();  
    eat();  
    putforks();  
}
```

Basic loop of each philosopher

```
// helper functions  
int left(int p) { return p; }  
  
int right(int p) {  
    return (p + 1) % 5;  
}
```

Helper functions (Downey's solutions)

- Philosopher p wishes to refer to the fork on their left \rightarrow call $\text{left}(p)$
- Philosopher p wishes to refer to the fork on their right \rightarrow call $\text{right}(p)$

The Dining Philosophers

- We need some **semaphore**, one for each fork: `sem_t forks[5]`

```
1 void getforks() {  
2     sem_wait(forks[left(p)]);  
3     sem_wait(forks[right(p)]);  
4 }  
5  
6 void putforks() {  
7     sem_post(forks[left(p)]);  
8     sem_post(forks[right(p)]);  
9 }
```

The `getforks()` and `putforks()` Routines (Broken Solution)

- **Deadlock** occur!
 - If each philosopher happens to **grab the fork on their left** before any philosopher can grab the fork on their right
 - Each will be stuck *holding one fork* and waiting for another, *forever*

A Solution: Breaking The Dependency

- Change how forks are acquired
 - Let's assume that philosopher 4 acquire the forks in a *different order*

```
1 void getforks () {
2     if (p == 4) {
3         sem_wait(forks[right(p)]);
4         sem_wait(forks[left(p)]);
5     } else {
6         sem_wait(forks[left(p)]);
7         sem_wait(forks[right(p)]);
8     }
9 }
```

- There is no situation where each philosopher grabs one fork and is stuck waiting for another
- **The cycle of waiting is broken**

How To Implement Semaphores

- Build our own version of semaphores called **Zemaphores**

```
1  typedef struct __Zem_t {  
2      int value;  
3      pthread_cond_t cond;  
4      pthread_mutex_t lock;  
5  } Zem_t;  
6  
7  // only one thread can call this  
8  void Zem_init(Zem_t *s, int value) {  
9      s->value = value;  
10     Cond_init(&s->cond);  
11     Mutex_init(&s->lock);  
12 }  
13  
14 void Zem_wait(Zem_t *s) {  
15     Mutex_lock(&s->lock);  
16     while (s->value <= 0)  
17         Cond_wait(&s->cond, &s->lock);  
18     s->value--;  
19     Mutex_unlock(&s->lock);  
20 }  
21 ...
```

How To Implement Semaphores

```
22 void Zem_post(Zem_t *s) {  
23     Mutex_lock(&s->lock);  
24     s->value++;  
25     Cond_signal(&s->cond);  
26     Mutex_unlock(&s->lock);  
27 }
```

- Zemaphore don't maintain the invariant that the value of the semaphore reflects the number of waiting threads
 - The value never be lower than zero
 - This behavior is **easier** to implement and **matches** the current Linux implementation