



# Concurrent Data Structures

# Lock-based Concurrent Data Structures

- Adding locks to a data structure makes the structure thread safe
  - How locks are added determine both the **correctness** and **performance** of the data structure

# Example: Concurrent Counters without Locks

- Incorrect version

```
1     typedef struct __counter_t {
2         int value;
3     } counter_t;
4
5     void init(counter_t *c) {
6         c->value = 0;
7     }
8
9     void increment(counter_t *c) {
10        c->value++;
11    }
12
13    void decrement(counter_t *c) {
14        c->value--;
15    }
16
17    int get(counter_t *c) {
18        return c->value;
19    }
```

# Example: Concurrent Counters with Locks

- Add a **single lock**
  - The lock is acquired when calling a routine that manipulates the data structure

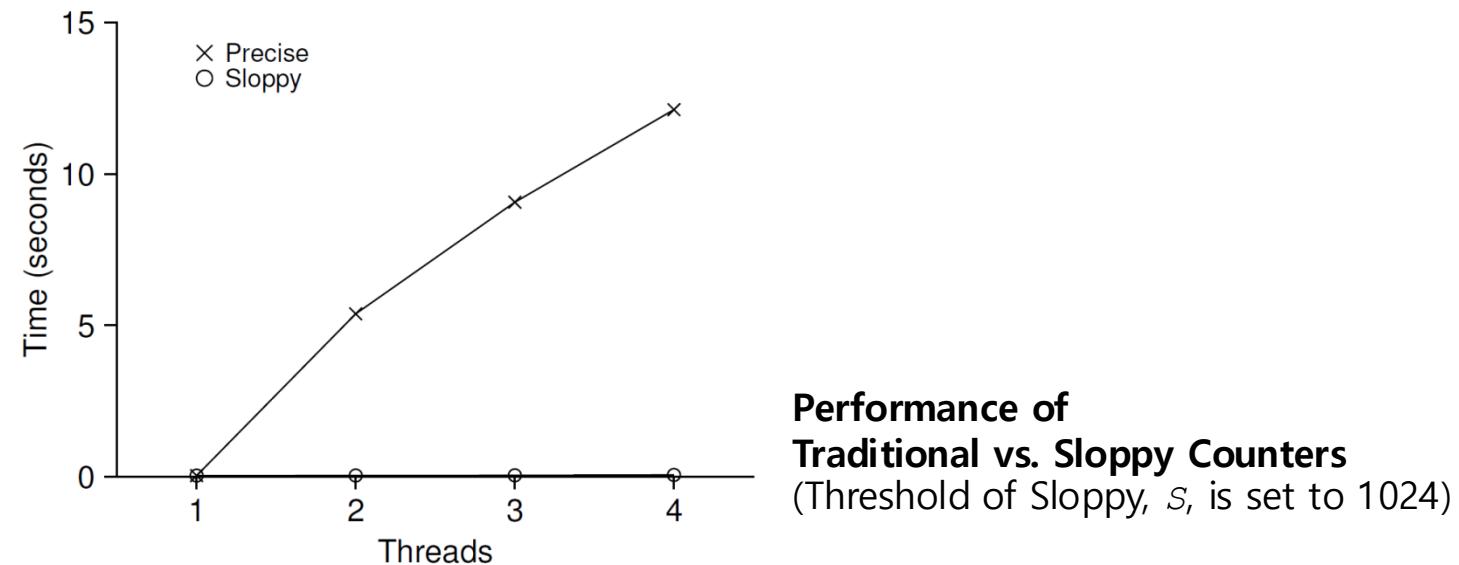
```
1     typedef struct __counter_t {  
2         int value;  
3         pthread_lock_t lock;  
4     } counter_t;  
5  
6     void init(counter_t *c) {  
7         c->value = 0;  
8         Pthread_mutex_init(&c->lock, NULL);  
9     }  
10  
11    void increment(counter_t *c) {  
12        Pthread_mutex_lock(&c->lock);  
13        c->value++;  
14        Pthread_mutex_unlock(&c->lock);  
15    }  
16
```

# Example: Concurrent Counters with Locks

```
(Cont.)  
17     void decrement(counter_t *c) {  
18         Pthread_mutex_lock(&c->lock);  
19         c->value--;  
20         Pthread_mutex_unlock(&c->lock);  
21     }  
22  
23     int get(counter_t *c) {  
24         Pthread_mutex_lock(&c->lock);  
25         int rc = c->value;  
26         Pthread_mutex_unlock(&c->lock);  
27         return rc;  
28     }
```

# Overhead of Simple Approach

- Each thread updates a single shared counter
  - Each thread updates the counter one million times
  - iMac with four Intel 2.7GHz i5 CPUs



Synchronized counter scales poorly

# Perfect Scaling

- Even though more work is done, it is **done in parallel**
- The time taken to complete the task is *not increased*

# Sloppy counter

- The sloppy counter works by representing ...
  - A single **logical counter** via numerous local physical counters, one per CPU core
  - A single **global counter**
  - There is only one lock for the global counter
- Example: on a machine with four CPUs
  - Four local counters
  - One global counter

# Basic Ideas of Sloppy Counting

- When a thread running on a core wishes to increment the counter
  - It increments its local counter
  - Each CPU has its own local counter
    - Threads across CPUs can update local counters *without contention*
    - Thus counter updates are **scalable**
  - The local values are periodically transferred to the global counter
    - Acquire the global lock
    - Increment it by the local counter's value
    - Local counter is then reset to zero

# The basic idea of sloppy counting

- How often the local-to-global transfer occurs is determined by a threshold,  $S$  (sloppiness)
  - The smaller  $S$ :
    - The more the counter behaves like the *non-scalable counter*
  - The bigger  $S$ :
    - The more scalable the counter
    - The further off the global value might be from the *actual count*

# Sloppy Counter Example

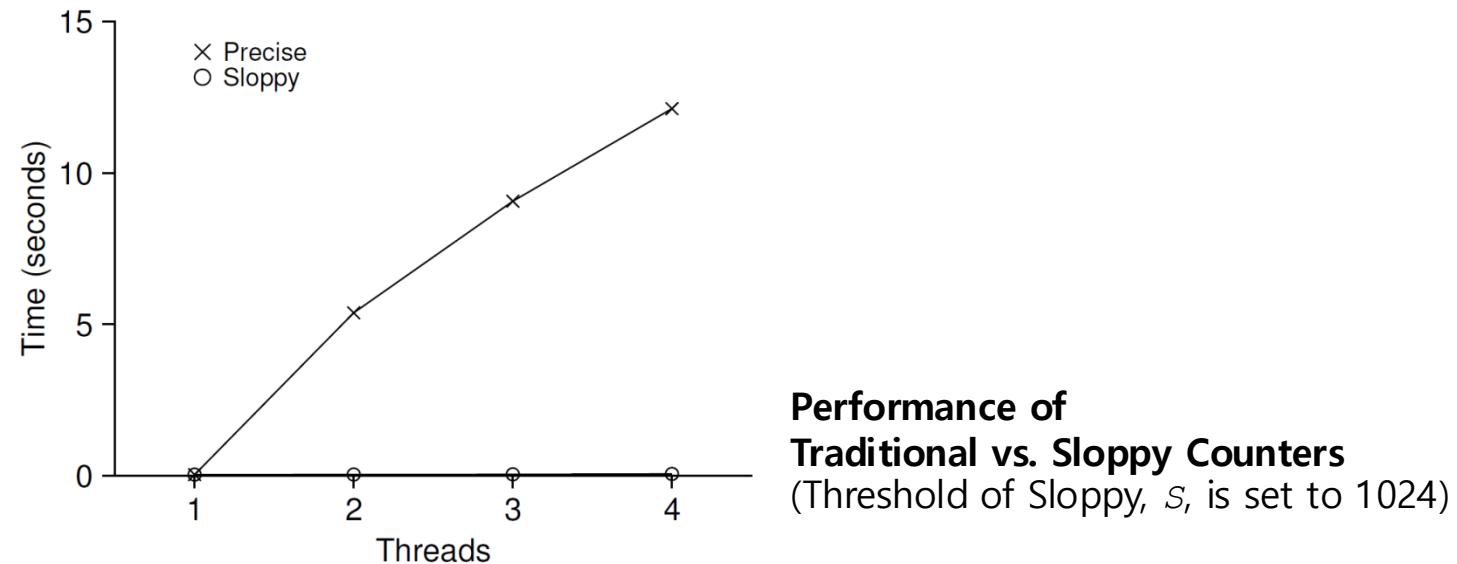
- Tracing the Sloppy Counters

- The threshold S is set to 5
- There are threads on each of four CPUs
- Each thread updates their local counters  $L_1 \dots L_4$

Time	$L_1$	$L_2$	$L_3$	$L_4$	$G$
0	0	0	0	0	0
1	0	0	1	1	0
2	1	0	2	1	0
3	2	0	3	1	0
4	3	0	3	2	0
5	4	1	3	3	0
6	5 → 0	1	3	4	5 (from $L_1$ )
7	0	2	4	5 → 0	10 (from $L_4$ )

# Sloppy Counter Performance

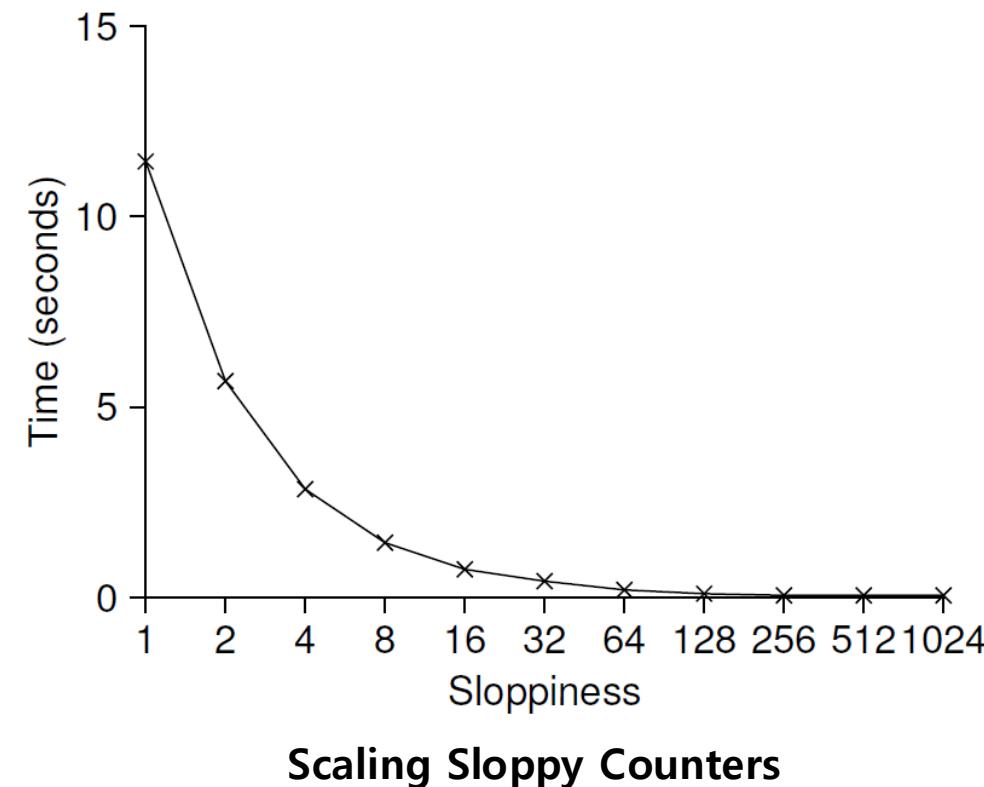
- Each thread updates a single shared counter
  - Each thread updates the counter one million times
  - iMac with four Intel 2.7GHz i5 CPUs



Sloppy counter scales well

# Importance of Threshold Value $S$

- Each four threads increments counter 1 million times, respectively
  - Low  $S \rightarrow$  **poor** performance, accurate global counter
  - High  $S \rightarrow$  Performance is **excellent**, The global count **lags**.



# Sloppy Counter Implementation

```
1  typedef struct __counter_t {
2      int global;           // global count
3      pthread_mutex_t glock; // global lock
4      int local[NUMCPUS];   // local count (per cpu)
5      pthread_mutex_t llock[NUMCPUS]; // ... and locks
6      int threshold;        // update frequency
7  } counter_t;
8
9  // init: record threshold, init locks, init values
10 //       of all local counts and global count
11 void init(counter_t *c, int threshold) {
12     c->threshold = threshold;
13
14     c->global = 0;
15     pthread_mutex_init(&c->glock, NULL);
16
17     int i;
18     for (i = 0; i < NUMCPUS; i++) {
19         c->local[i] = 0;
20         pthread_mutex_init(&c->llock[i], NULL);
21     }
22 }
23 }
```

# Sloppy Counter Implementation

```
(Cont.)  
24    // update: usually, just grab local lock and update local amount  
25    //          once local count has risen by 'threshold', grab global  
26    //          lock and transfer local values to it  
27    void update(counter_t *c, int threadID, int amt) {  
28        pthread_mutex_lock(&c->llock[threadID]);  
29        c->local[threadID] += amt;           // assumes amt > 0  
30        if (c->local[threadID] >= c->threshold) { // transfer to global  
31            pthread_mutex_lock(&c->glock);  
32            c->global += c->local[threadID];  
33            pthread_mutex_unlock(&c->glock);  
34            c->local[threadID] = 0;  
35        }  
36        pthread_mutex_unlock(&c->llock[threadID]);  
37    }  
38  
39    // get: just return global amount (which may not be perfect)  
40    int get(counter_t *c) {  
41        pthread_mutex_lock(&c->glock);  
42        int val = c->global;  
43        pthread_mutex_unlock(&c->glock);  
44        return val;      // only approximate!  
45    }
```

# Concurrent Linked Lists

```
1 // basic node structure
2 typedef struct __node_t {
3     int key;
4     struct __node_t *next;
5 } node_t;
6
7 // basic list structure (one used per list)
8 typedef struct __list_t {
9     node_t *head;
10    pthread_mutex_t lock;
11 } list_t;
12
13 void List_Init(list_t *L) {
14     L->head = NULL;
15     pthread_mutex_init(&L->lock, NULL);
16 }
17
(Cont.)
```

# Concurrent Linked Lists

(Cont.)

```
18     int List_Insert(list_t *L, int key) {
19         pthread_mutex_lock(&L->lock);
20         node_t *new = malloc(sizeof(node_t));
21         if (new == NULL) {
22             perror("malloc");
23             pthread_mutex_unlock(&L->lock);
24             return -1; } // fail
25         new->key = key;
26         new->next = L->head;
27         L->head = new;
28         pthread_mutex_unlock(&L->lock);
29         return 0; // success
30     }
31 }
```

(Cont.)

# Concurrent Linked Lists

(Cont.)

```
32
32     int List_Lookup(list_t *L, int key) {
33         pthread_mutex_lock(&L->lock);
34         node_t *curr = L->head;
35         while (curr) {
36             if (curr->key == key) {
37                 pthread_mutex_unlock(&L->lock);
38                 return 0; // success
39             }
40             curr = curr->next;
41         }
42         pthread_mutex_unlock(&L->lock);
43         return -1; // failure
44     }
```

# Concurrent Linked Lists

- The code **acquires** a lock in the insert routine upon entry
- The code **releases** the lock upon exit
  - If `malloc()` happens to *fail*, the code must also release the lock before failing the insert
  - This kind of exceptional control flow has been shown to be **quite error prone**
  - **Solution:** The lock and release *only surround* the actual critical section in the insert code

# Concurrent Linked List: Rewritten

```
1      void List_Init(list_t *L) {
2          L->head = NULL;
3          pthread_mutex_init(&L->lock, NULL);
4      }
5
6      void List_Insert(list_t *L, int key) {
7          // synchronization not needed
8          node_t *new = malloc(sizeof(node_t));
9          if (new == NULL) {
10              perror("malloc");
11              return;
12          }
13          new->key = key;
14
15          // just lock critical section
16          pthread_mutex_lock(&L->lock);
17          new->next = L->head;
18          L->head = new;
19          pthread_mutex_unlock(&L->lock);
20      }
21
```

# Concurrent Linked List: Rewritten

```
(Cont.)  
22     int List_Lookup(list_t *L, int key) {  
23         int rv = -1;  
24         pthread_mutex_lock(&L->lock);  
25         node_t *curr = L->head;  
26         while (curr) {  
27             if (curr->key == key) {  
28                 rv = 0;  
29                 break;  
30             }  
31             curr = curr->next;  
32         }  
33         pthread_mutex_unlock(&L->lock);  
34         return rv; // now both success and failure  
35     }
```

# Scaling Linked List

- Hand-over-hand locking (lock coupling)
  - Add **a lock per node** of the list instead of having a single lock for the entire list
  - When traversing the list
    - First grabs the next node's lock
    - And then releases the current node's lock
  - Enable a high degree of concurrency in list operations
    - However, in practice, the overheads of acquiring and releasing locks for each node of a list traversal is *prohibitive*

# Michael and Scott Concurrent Queues

- There are two locks
  - One for the **head** of the queue
  - One for the **tail**
  - The goal of these two locks is to enable concurrency of *enqueue* and *dequeue* operations
- Add a dummy node
  - Allocated in the queue initialization code
  - Enable the separation of head and tail operations



Me with Prof. Scott in 2024

# Concurrent Queues

```
1  typedef struct __node_t {
2      int value;
3      struct __node_t *next;
4  } node_t;
5
6  typedef struct __queue_t {
7      node_t *head;
8      node_t *tail;
9      pthread_mutex_t headLock;
10     pthread_mutex_t tailLock;
11 } queue_t;
12
13 void Queue_Init(queue_t *q) {
14     node_t *tmp = malloc(sizeof(node_t));
15     tmp->next = NULL;
16     q->head = q->tail = tmp;
17     pthread_mutex_init(&q->headLock, NULL);
18     pthread_mutex_init(&q->tailLock, NULL);
19 }
20
(Cont.)
```

# Concurrent Queues



```
(Cont.)  
21     void Queue_Enqueue(queue_t *q, int value) {  
22         node_t *tmp = malloc(sizeof(node_t));  
23         assert(tmp != NULL);  
24  
25         tmp->value = value;  
26         tmp->next = NULL;  
27  
28         pthread_mutex_lock(&q->tailLock);  
29         q->tail->next = tmp;  
30         q->tail = tmp;  
31         pthread_mutex_unlock(&q->tailLock);  
32     }  
(Cont.)
```

# Concurrent Queues

(Cont.)

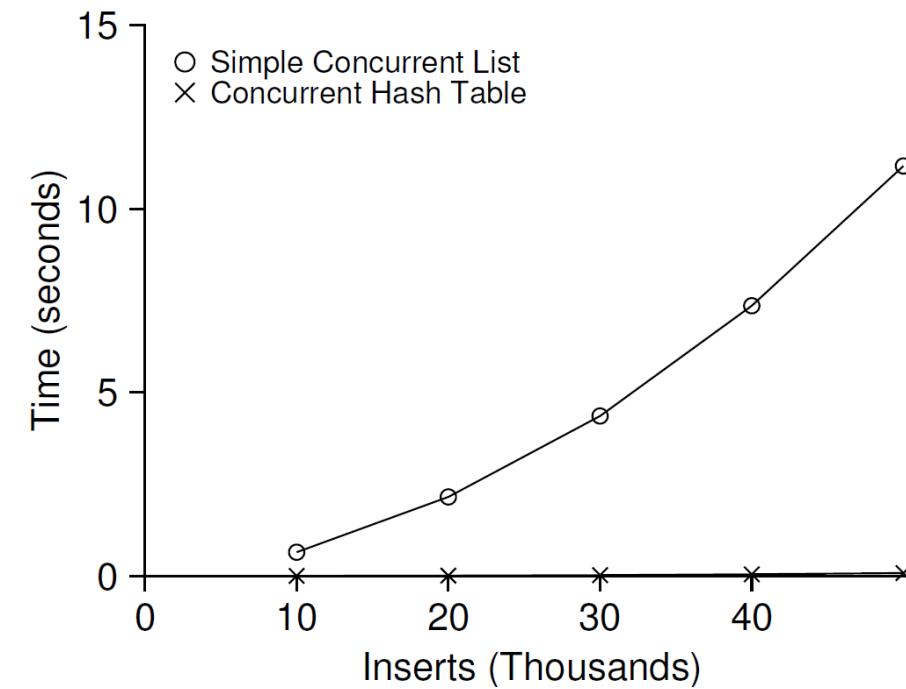
```
33     int Queue_Dequeue(queue_t *q, int *value) {
34         pthread_mutex_lock(&q->headLock);
35         node_t *tmp = q->head;
36         node_t *newHead = tmp->next;
37         if (newHead == NULL) {
38             pthread_mutex_unlock(&q->headLock);
39             return -1; // queue was empty
40         }
41         *value = newHead->value;
42         q->head = newHead;
43         pthread_mutex_unlock(&q->headLock);
44         free(tmp);
45         return 0;
46     }
```

# Concurrent Hash Table

- Focus on a simple hash table
  - The hash table does not resize
  - Built using the concurrent lists
  - It uses a **lock per hash bucket** each of which is represented by *a list*

# Performance of Concurrent Hash Table

- From 10,000 to 50,000 concurrent updates from each of four threads
  - iMac with four Intel 2.7GHz i5 CPUs



The simple concurrent hash table  
scales magnificently

# Concurrent Hash Table

```
1      #define BUCKETS (101)
2
3      typedef struct __hash_t {
4          list_t lists[BUCKETS];
5      } hash_t;
6
7      void Hash_Init(hash_t *H) {
8          int i;
9          for (i = 0; i < BUCKETS; i++) {
10              List_Init(&H->lists[i]);
11          }
12      }
13
14      int Hash_Insert(hash_t *H, int key) {
15          int bucket = key % BUCKETS;
16          return List_Insert(&H->lists[bucket], key);
17      }
18
19      int Hash_Lookup(hash_t *H, int key) {
20          int bucket = key % BUCKETS;
21          return List_Lookup(&H->lists[bucket], key);
22      }
```