



Common Concurrency Bugs and Deadlocks

What Types Of Bugs Exist?

- Focus on four major open-source applications
 - MySQL, Apache, Mozilla, OpenOffice

| Application | What it does | Non-Deadlock | Deadlock |
|--------------|-----------------|--------------|-----------|
| MySQL | Database Server | 14 | 9 |
| Apache | Web Server | 13 | 4 |
| Mozilla | Web Browser | 41 | 16 |
| Open Office | Office Suite | 6 | 2 |
| Total | | 74 | 31 |

Bugs In Modern Applications

Non-Deadlock Bugs

- Make up a majority of concurrency bugs
- Two major types of non deadlock bugs
 - Atomicity violation
 - Order violation

Atomicity-Violation Bugs

- The desired **serializability** among multiple memory accesses is *violated*
 - Simple Example found in MySQL
 - Two different threads access the field `proc_info` in the struct `thd`

```
1  Thread1::  
2  if(thd->proc_info){  
3      ...  
4      fputs(thd->proc_info , ...);  
5      ...  
6  }  
7  
8  Thread2::  
9  thd->proc_info = NULL;
```

Atomicity-Violation Bugs

- **Solution:** Simply add locks around the shared-variable references

```
1  pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER;
2
3  Thread1::
4  pthread_mutex_lock(&lock);
5  if(thd->proc_info) {
6      ...
7      fputs(thd->proc_info , ...);
8      ...
9  }
10 pthread_mutex_unlock(&lock);
11
12 Thread2::
13 pthread_mutex_lock(&lock);
14 thd->proc_info = NULL;
15 pthread_mutex_unlock(&lock);
```

Order-Violation Bugs

- The **desired order** between two memory accesses is flipped
 - i.e., **A** should always be executed before **B**, but the order is not enforced during execution
 - **Example:**
 - The code in Thread2 seems to assume that the variable `mThread` has already been *initialized* (and is not `NULL`)

```
1  Thread1::  
2  void init() {  
3      mThread = PR_CreateThread(mMain, ...);  
4  }  
5  
6  Thread2::  
7  void mMain(...) {  
8      mState = mThread->State  
9  }
```

Order-Violation Bugs

- **Solution:** Enforce ordering using **condition variables**

```
1  pthread_mutex_t mtLock = PTHREAD_MUTEX_INITIALIZER;
2  pthread_cond_t mtCond = PTHREAD_COND_INITIALIZER;
3  int mtInit = 0;
4
5  Thread 1::
6  void init() {
7      ...
8      mThread = PR_CreateThread(mMain, ...);
9
10     // signal that the thread has been created.
11     pthread_mutex_lock(&mtLock);
12     mtInit = 1;
13     pthread_cond_signal(&mtCond);
14     pthread_mutex_unlock(&mtLock);
15     ...
16 }
17
18 Thread2::
19 void mMain(...) {
20     ...
}
```

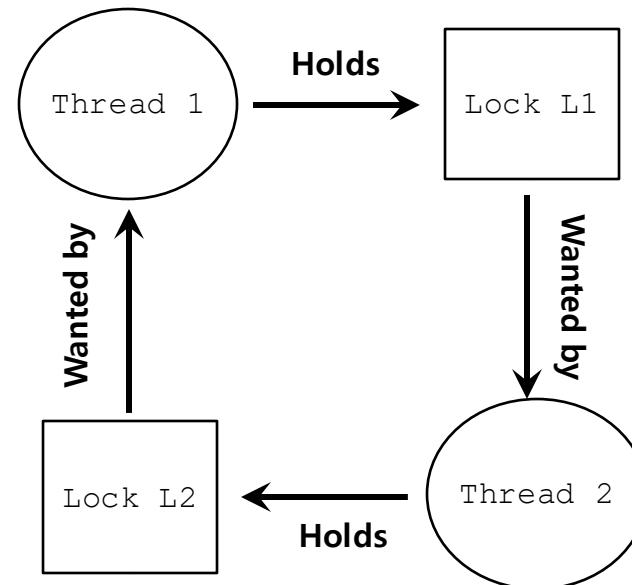
Order-Violation Bugs

```
21     // wait for the thread to be initialized ...
22     pthread_mutex_lock(&mtLock);
23     while(mtInit == 0)
24         pthread_cond_wait(&mtCond, &mtLock);
25     pthread_mutex_unlock(&mtLock);
26
27     mState = mThread->State;
28     ...
29 }
```

Deadlock Bugs

```
Thread 1:          Thread 2:  
lock(L1);          lock(L2);  
lock(L2);          lock(L1);
```

- The presence of a cycle
 - Thread1 is holding a lock L1 and waiting for another one, L2.
 - Thread2 that holds lock L2 is waiting for L1 to be release.



Why Do Deadlocks Occur?



- Reason 1
 - In large code bases, **complex dependencies** arise between components
- Reason 2
 - Due to the nature of **encapsulation**
 - Hide details of implementations and make software easier to build in a modular way
 - Such **modularity** *does not mesh well with locking*

Why Do Deadlocks Occur?



- **Example:** Java Vector class and the method AddAll()

```
1  Vector v1,v2;  
2  v1.AddAll(v2);
```

- **Locks** for both the vector being added to (`v1`) and the parameter (`v2`) *need to be acquired*
 - The routine acquires said locks in some arbitrary order (`v1` then `v2`)
 - If some other thread calls `v2.AddAll(v1)` at nearly the same time → We have the potential for **deadlock**

Conditional for Deadlock

- Four conditions need to hold for a deadlock to occur

| Condition | Description |
|------------------|--|
| Mutual Exclusion | Threads claim exclusive control of resources that they require. |
| Hold-and-wait | Threads hold resources allocated to them while waiting for additional resources |
| No preemption | Resources cannot be forcibly removed from threads that are holding them. |
| Circular wait | There exists a circular chain of threads such that each thread holds one more resources that are being requested by the next thread in the chain |

- If any of these four conditions are not met, **deadlock cannot occur**

Prevention – Circular Wait

- Provide a total ordering on lock acquisition
 - This approach requires *careful design* of global locking strategies
- **Example**
 - There are two locks in the system (L1 and L2)
 - We can prevent deadlock by always acquiring L1 before L2

Prevention – Hold-and-wait

- Acquire all locks at once, atomically

```
1  lock (prevention) ;  
2  lock (L1) ;  
3  lock (L2) ;  
4  ...  
5  unlock (prevention) ;
```

- This code guarantees that **no untimely thread switch can occur** *in the midst of lock acquisition*
- **Problem**
 - Require us to know when calling a routine exactly which locks must be held and to acquire them ahead of time
 - Decrease *concurrency*

Prevention – No Preemption

- **Multiple lock acquisition** often gets us into trouble because when waiting for one lock we are holding another
- `trylock()`
 - Used to build a *deadlock-free, ordering-robust* lock acquisition protocol
 - Grab the lock (if it is available)
 - Or, return -1: you should try again later

```
1  top:  
2      lock(L1);  
3      if( tryLock(L2) == -1 ) {  
4          unlock(L1);  
5          goto top;  
6      }
```

Prevention – No Preemption



- livelock

- Both systems are running through the code sequence *over and over again*
- Progress is not being made
- Solution
 - Add a **random delay** before looping back and trying the entire thing over again

Prevention – Mutual Exclusion

- wait-free
 - Using powerful **hardware instruction**
 - You can build data structures in a manner that *does not require explicit locking*

```
1 int CompareAndSwap(int *address, int expected, int new) {  
2     if (*address == expected) {  
3         *address = new;  
4         return 1; // success  
5     }  
6     return 0;  
7 }
```

Prevention – Mutual Exclusion

- We now wanted to **atomically increment** a value by a certain amount

```
1 void AtomicIncrement(int *value, int amount){  
2     do{  
3         int old = *value;  
4         }while( CompareAndSwap(value, old, old+amount)==0);  
5 }
```

- Repeatedly tries to update the value to *the new amount* and uses the compare-and-swap to do so
- **No lock** is acquired
- **No deadlock** can arise
- **livelock** is still a possibility

Prevention – Mutual Exclusion

- More complex example: list insertion

```
1 void insert(int value){  
2     node_t * n = malloc(sizeof(node_t));  
3     assert( n != NULL );  
4     n->value = value ;  
5     n->next = head;  
6     head = n;  
7 }
```

- If called by multiple threads at the "*same time*", this code has a **race condition**

Prevention – Mutual Exclusion

■ Solution

- Surrounding this code with a **lock acquire** and **release**

```
1 void insert(int value){  
2     node_t * n = malloc(sizeof(node_t));  
3     assert( n != NULL );  
4     n->value = value ;  
5     lock(listlock); // begin critical section  
6     n->next = head;  
7     head = n;  
8     unlock(listlock) ; //end critical section  
9 }
```

- wait-free manner** using the compare-and-swap instruction

```
1 void insert(int value) {  
2     node_t *n = malloc(sizeof(node_t));  
3     assert(n != NULL);  
4     n->value = value;  
5     do {  
6         n->next = head;  
7     } while (CompareAndSwap(&head, n->next, n) == 0);  
8 }
```

Deadlock Avoidance via Scheduling

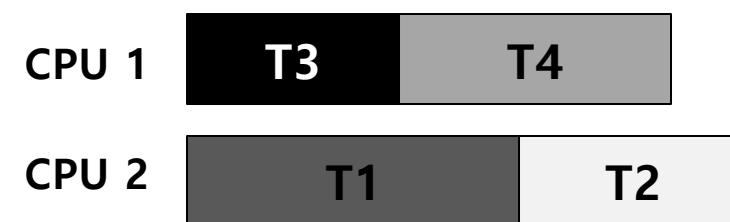
- In some scenarios **deadlock avoidance** is preferable
 - **Global knowledge** is required
 - Which locks various threads might grab during their execution
 - Subsequently schedules said threads in a way as to guarantee no deadlock can occur

Example of Deadlock Avoidance via Scheduling (1)

- We have two processors and four threads
 - Lock acquisition demands of the threads

| | T1 | T2 | T3 | T4 |
|----|-----|-----|-----|----|
| L1 | yes | yes | no | no |
| L2 | yes | yes | yes | no |

- A smart scheduler could compute that as long as T1 and T2 are not run at the same time, **no deadlock** could ever arise



Example of Deadlock Avoidance via Scheduling (2)

- More contention for the same resources

| | T1 | T2 | T3 | T4 |
|----|-----|-----|-----|----|
| L1 | yes | yes | yes | no |
| L2 | yes | yes | yes | no |

- A possible schedule that guarantees that *no deadlock* could ever occur



- The total time to complete the jobs is lengthened considerably

Detect and Recover

- Allow deadlock to occasionally occur and then *take some action*
 - Example: if an OS froze, you would reboot it
- Many database systems employ *deadlock detection* and *recovery technique*
 - A deadlock detector **runs periodically**
 - Building a **resource graph** and checking it for cycles
 - In deadlock, the system **need to be restarted**