

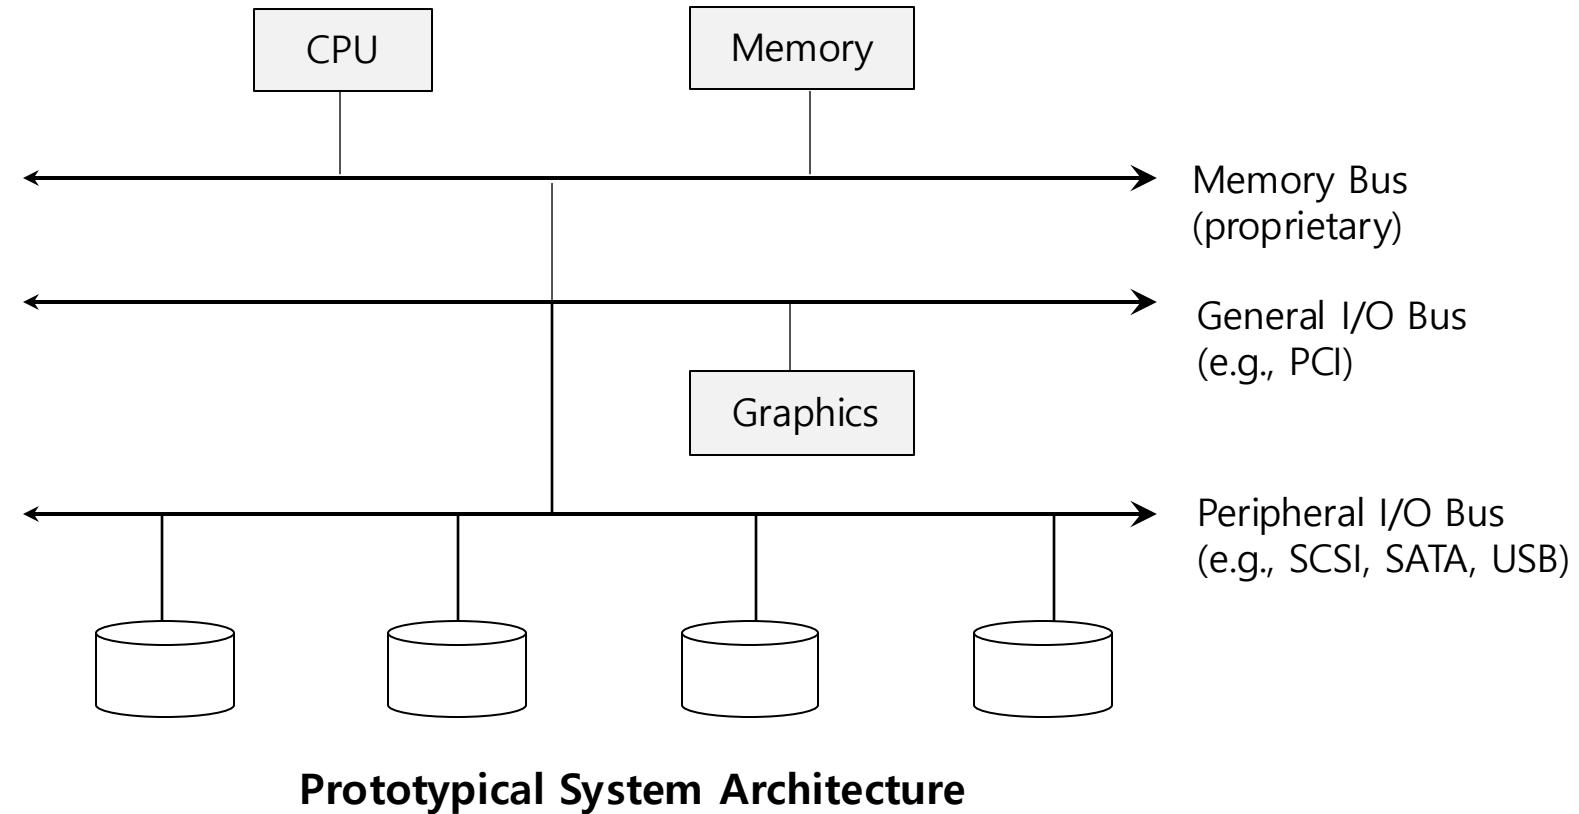


# I/O Subsystems

# I/O Devices

- I/O is critical to computer systems to interact with users or other devices
- Issues
  - How should I/O be integrated into OS?
  - What are the general mechanisms?
  - How can we make it efficiently?

# Structure of Input/Output (I/O) Devices



CPU is attached to the main memory of the system via some kind of **memory bus**

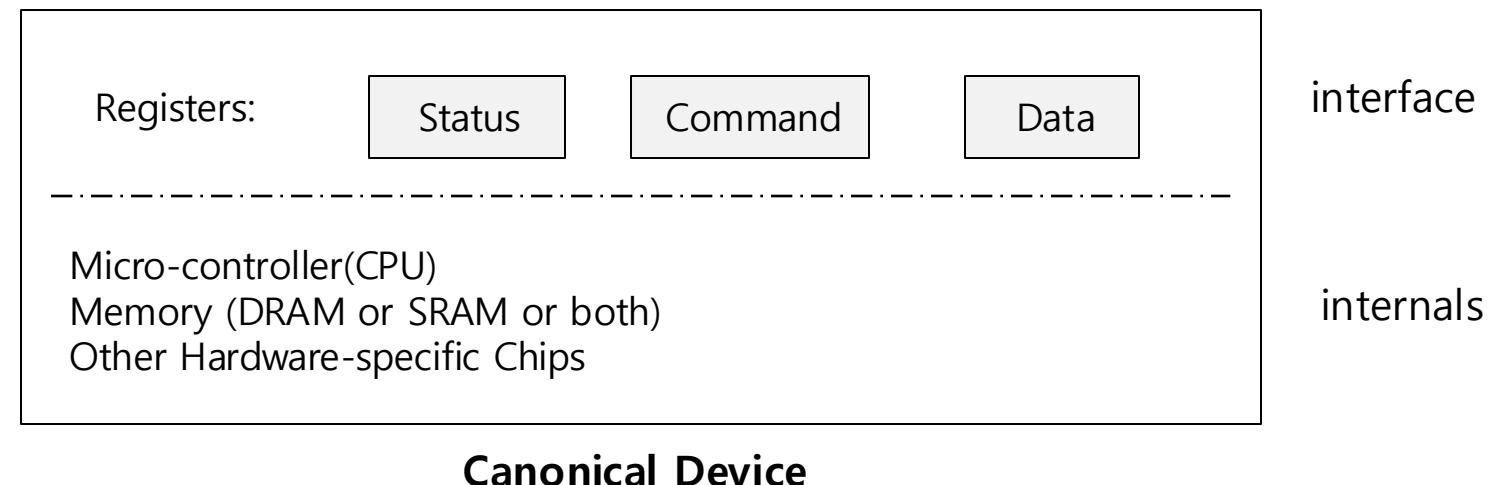
Some devices are connected to the system via a **general I/O bus**

# I/O Architecture

- Buses
  - Data paths that transfer information between CPU(s), RAM, and I/O devices
- I/O Bus
  - Data path that connects a CPU to an I/O device
  - I/O bus is connected to I/O device by three hardware components: I/O ports, interfaces and device controllers

# Canonical Device

- Canonical devices has two important components
  - Hardware interface allows the system software to control its operation
  - Internals which is implementation specific



# Hardware Interface of Canonical Device

- Status registers
  - Show the current status of the device
- Command registers
  - Tell the device to perform a certain task
- Data registers
  - Pass data to the device, or get data from the device

By reading and writing above **three registers**,  
the operating system can **control device behavior**

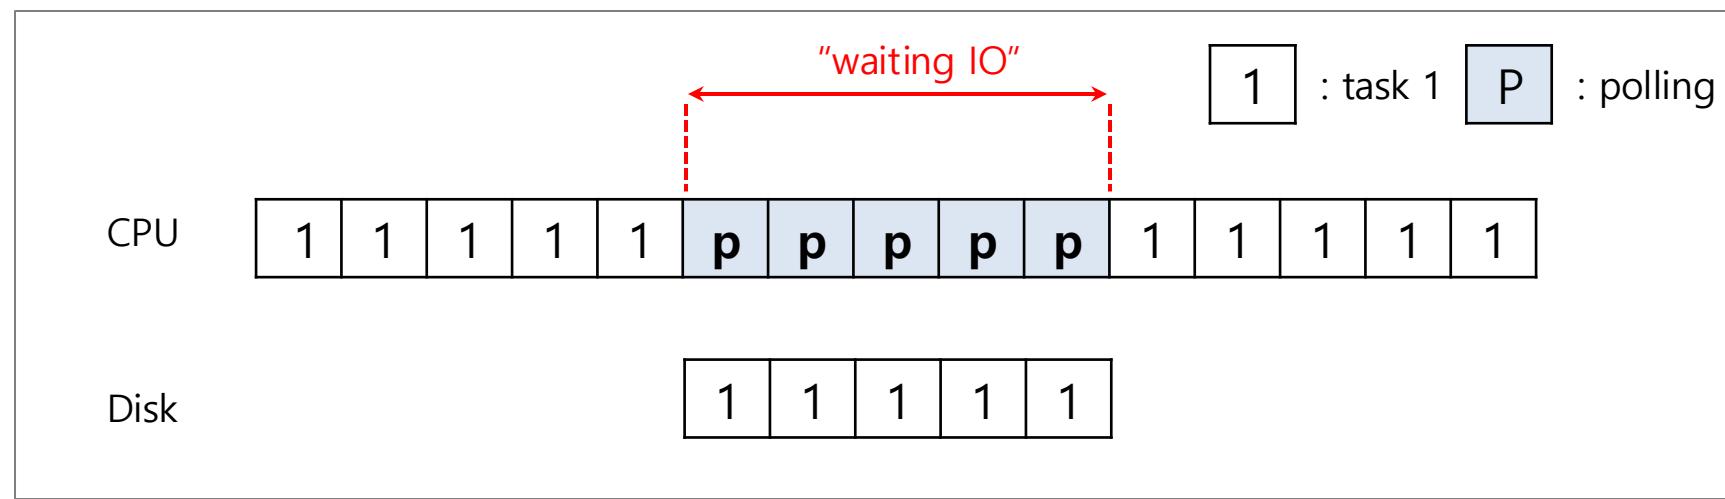
# Hardware Interface of Canonical Device

- Typical interaction example

```
while ( STATUS == BUSY)
    ; //wait until device is not busy
write data to data register
write command to command register
    Doing so starts the device and executes the command
while ( STATUS == BUSY)
    ; //wait until device is done with your request
```

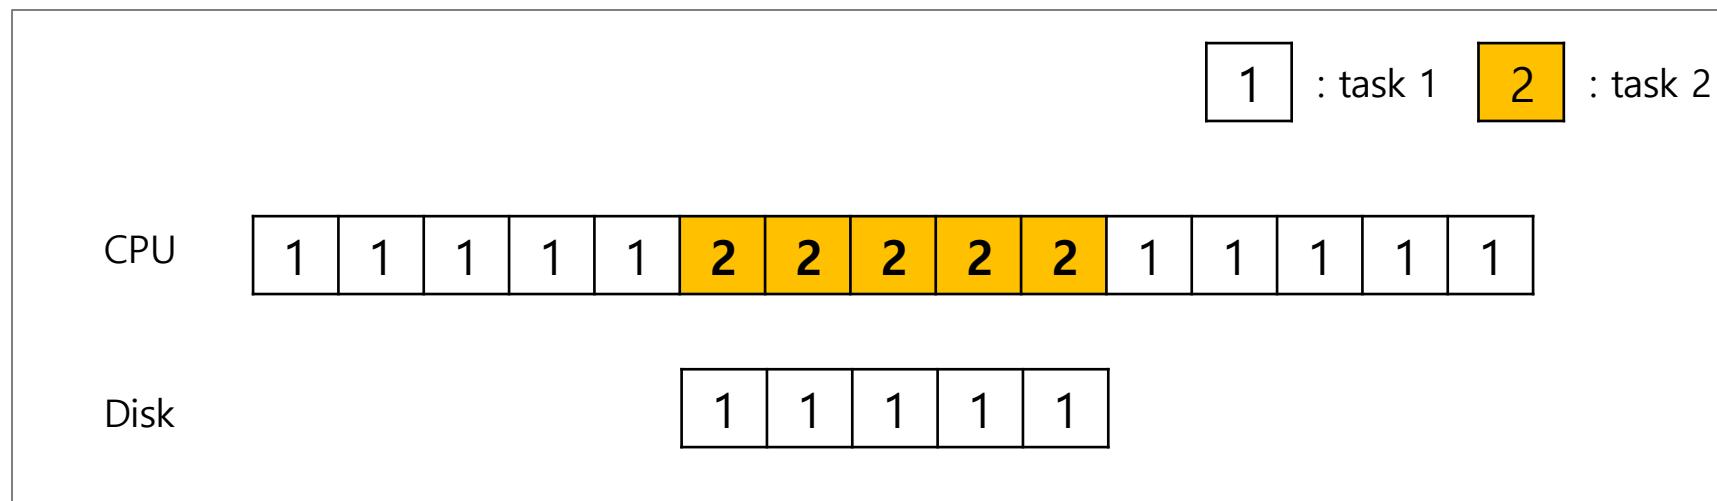
# Polling

- Operating system waits until the device is ready by repeatedly reading the status register
  - Positive aspect: simple and working
  - However, it wastes CPU time just waiting for the device
    - Switching to another ready process is better utilizing the CPU



# Interrupts

- Put the I/O request process to sleep and context switch to another
- When the device is finished, wake the process waiting for the I/O by interrupt
  - Positive aspect: allowing CPU and disks to be properly utilized



**Diagram of CPU utilization by interrupt**

# Polling vs interrupts

- However, “interrupts is not always the best solution”
  - If, device performs very quickly, interrupt will “slow down” the system
  - Because context switch is expensive (switching to another process)

If a device is fast → **poll** is best  
If it is slow → **interrupts** is better

# CPU is once again over-burdened

- CPU wastes a lot of time to copy large chunks of data from memory to the device

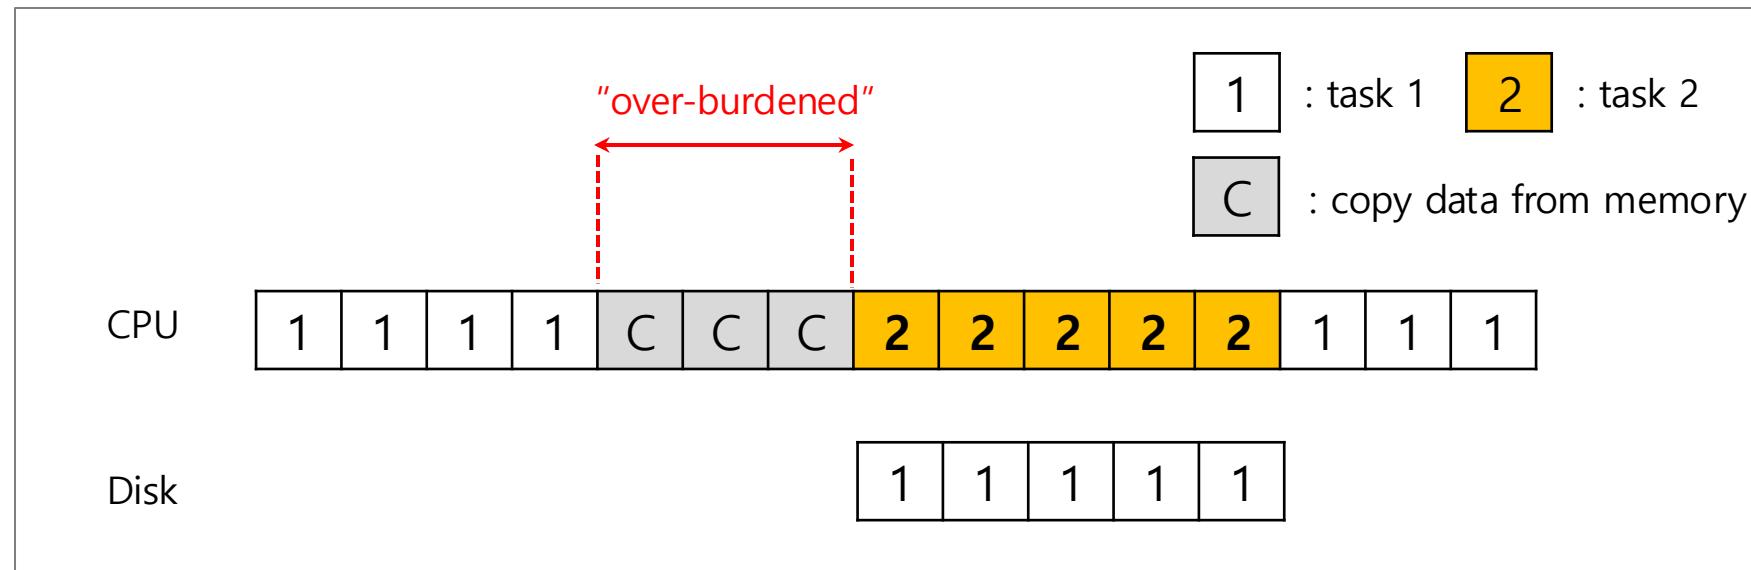


Diagram of CPU utilization

# DMA (Direct Memory Access)

- Copy data in memory by knowing “where the data lives in memory, how much data to copy”
- When completed, DMA raises an interrupt, I/O begins on disk

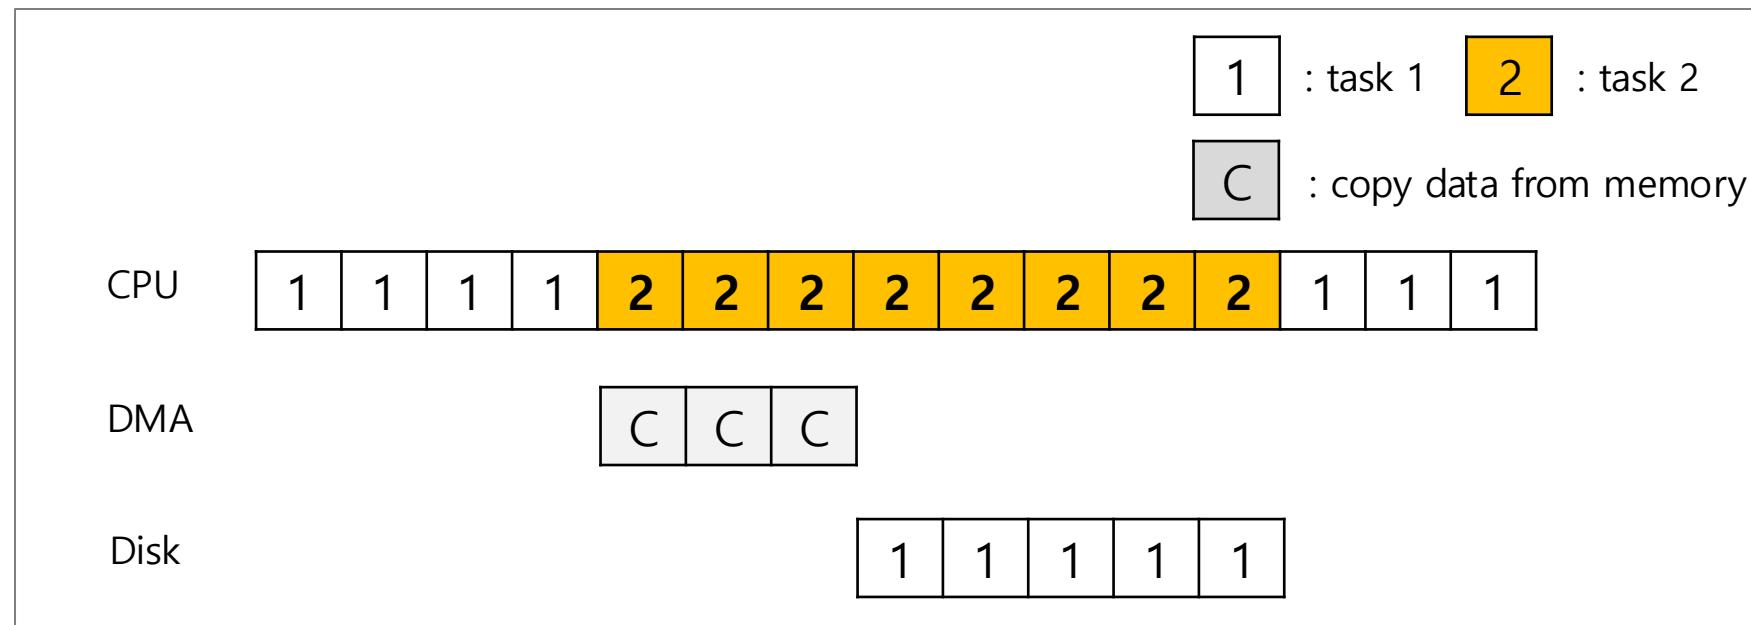


Diagram of CPU utilization by DMA

# Device Interaction

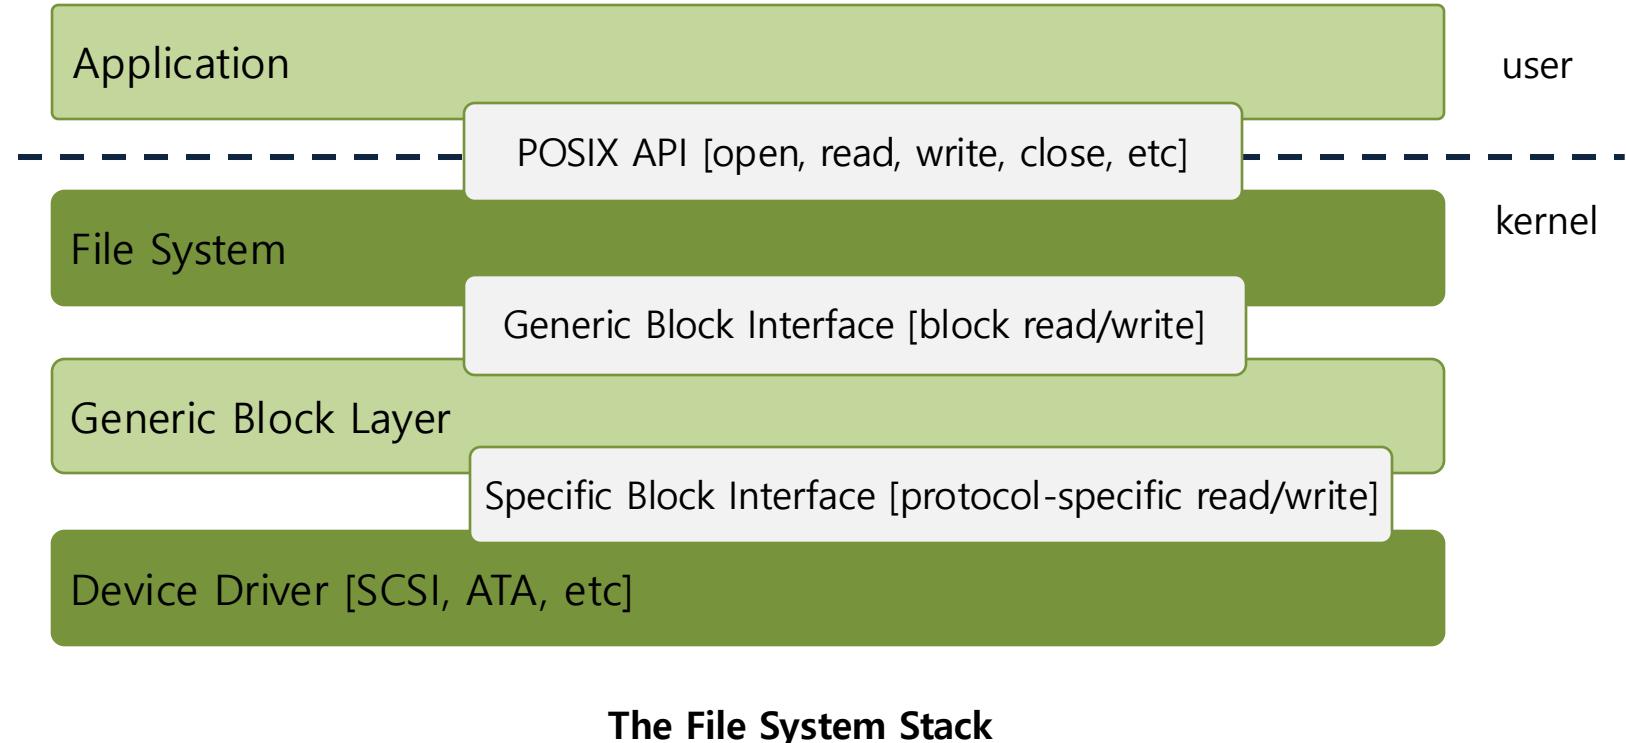
- How the OS communicates with the device?
- Solutions
  - I/O instructions: a way for the OS to send data to specific device registers
    - Ex) in and out instructions on x86
  - memory-mapped I/O
    - Device registers available as if they were memory locations
    - The OS load (to read) or store (to write) to the device instead of main memory

# Device Interaction

- How the OS interact with different specific interfaces?
  - Ex) We'd like to build a file system that worked on top of SCSI disks, IDE disks, USB keychain drivers, and so on
- Solutions: Abstraction
  - Abstraction encapsulate any specifics of device interaction

# File system Abstraction

- File system specifics of which disk class it is using
  - Ex) It issues block read and write request to the generic block layer



# Problem of File System Abstraction

- If there is a device having many special capabilities, these capabilities will go unused in the generic interface layer
- Over 70% of OS code is found in device drivers
  - Any device drivers are needed because you might plug it to your system
  - They are primary contributor to kernel crashes, making more bugs

# A Simple IDE Disk Driver

- Four types of registers
  - Control, command block, status and error
  - Memory mapped IO
  - in and out I/O instruction

- Control Registers:  
Address 0x3F6 = 0x80 (0000 1RE0): R=reset, E=0 means "enable interrupt"
- Command Block Registers:  
Address 0x1F0 = Data Port  
Address 0x1F1 = Error  
Address 0x1F2 = Sector Count  
Address 0x1F3 = LBA low byte  
Address 0x1F4 = LBA mid byte  
Address 0x1F5 = LBA hi byte  
Address 0x1F6 = 1B1D TOP4LBA: B=LBA, D=drive  
Address 0x1F7 = Command/status

- Status Register (Address 0x1F7):

7	6	5	4	3	2	1	0
BUSY	READY	FAULT	SEEK	DRQ	CORR	IDDEX	ERROR

- Error Register (Address 0x1F1): (check when Status ERROR==1)

7	6	5	4	3	2	1	0
BBK	UNC	MC	IDNF	MCR	ABRT	T0NF	AMNF

- BBK = Bad Block
- UNC = Uncorrectable data error
- MC = Media Changed
- IDNF = ID mark Not Found
- MCR = Media Change Requested
- ABRT = Command aborted
- T0NF = Track 0 Not Found
- AMNF = Address Mark Not Found

- **Wait for drive to be ready.** Read Status Register (0x1F7) until drive is not busy and READY
- **Write parameters to command registers.** Write the sector count, logical block address (LBA) of the sectors to be accessed, and drive number (master=0x00 or slave=0x10, as IDE permits just two drives) to command registers (0x1F2-0x1F6)
- **Start the I/O.** by issuing read/write to command register. Write READ—WRITE command to command register (0x1F7)
- **Data transfer (for writes):** Wait until drive status is READY and DRQ (drive request for data); write data to data port
- **Handle interrupts.** In the simplest case, handle an interrupt for each sector transferred; more complex approaches allow batching and thus one final interrupt when the entire transfer is complete
- Error handling. After each operation, read the status register. If the ERROR bit is on, read the error register for details

```
static int ide_wait_ready() {  
    while (((int r = inb(0x1f7)) & IDE_BSY) ||  
          !(r & IDE_DRDY))  
    ; // loop until drive isn't busy  
}
```

```
static void ide_start_request(struct buf *b) {  
    ide_wait_ready();  
    outb(0x3f6, 0); // generate interrupt  
    outb(0x1f2, 1); // how many sectors?  
    outb(0x1f3, b->sector & 0xff); // LBA goes here ...  
    outb(0x1f4, (b->sector >> 8) & 0xff); // ... and here  
    outb(0x1f5, (b->sector >> 16) & 0xff); // ... and here!  
    outb(0x1f6, 0xe0 | ((b->dev&1)<<4) | ((b->sector>>24)&0x0f));  
    if(b->flags & B_DIRTY){  
        outb(0x1f7, IDE_CMD_WRITE); // this is a WRITE  
        outsl(0x1f0, b->data, 512/4); // transfer data too!  
    } else {  
        outb(0x1f7, IDE_CMD_READ); // this is a READ (no data)  
    }  
}
```

```
void ide_rw(struct buf *b) {
    acquire(&ide_lock);
    for (struct buf **pp = &ide_queue; *pp; pp=&(*pp)->qnext)
        ; // walk queue
    *pp = b; // add request to end
    if (ide_queue == b) // if q is empty
        ide_start_request(b); // send req to disk
    while ((b->flags & (B_VALID|B_DIRTY)) != B_VALID)
        sleep(b, &ide_lock); // wait for completion
    release(&ide_lock);
}
```

```
void ide_intr() {
    struct buf *b;
    acquire(&ide_lock);
    if (!(b->flags & B_DIRTY) && ide_wait_ready(1) >= 0)
        insl(0x1f0, b->data, 512/4); // if READ: get data
    b->flags |= B_VALID;
    b->flags &= B_DIRTY;
    wakeup(b); // wake waiting process
    if ((ide_queue = b->qnext) != 0) // start next request
        ide_start_request(ide_queue); // (if one exists)
    release(&ide_lock);
}
```