



Hard Disks and Their Management

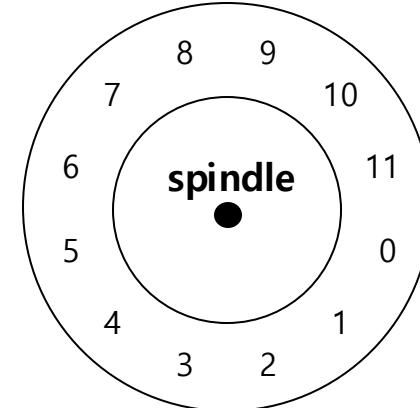
Hard Disk Drive

- Hard disk drive has been **the main form of persistent data storage** in computer systems for decades
 - The drive consists of a large number of **sectors** (512-byte blocks)
 - **Address Space :**
 - We can view the disk with n sectors as an array of sectors; 0 to $n-1$

Interface

- The only guarantee is that a single 512-byte write is **atomic**
- Multi-sector operations are possible
 - Many file systems will read or write 4KB at a time
 - **Torn write:**
 - If an untimely power loss occurs, only a portion of a larger write may complete
- Accessing blocks in a **contiguous chunk** is the fastest access mode
 - A sequential read or write
 - Much faster than any more random access pattern

Basic Geometry



A Disk with Just A Single Track (12 sectors)

- **Platter** (Aluminum coated with a thin magnetic layer)
 - A circular hard surface
 - Data is stored persistently by inducing magnetic changes to it
 - Each platter has 2 sides, each of which is called a **surface**

Basic Geometry



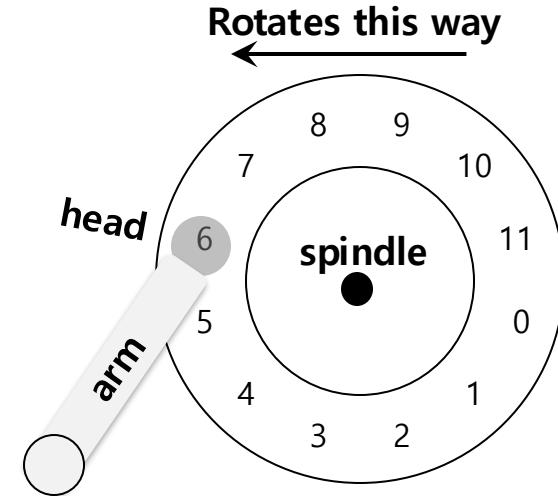
■ Spindle

- Spindle is connected to a motor that spins the platters around
- The rate of rotations is measured in **RPM** (Rotations Per Minute)
 - Typical modern values : 7,200 RPM to 15,000 RPM
 - E.g., 10000 RPM : A single rotation takes about 6 ms

■ Track

- Concentric circles of sectors
- Data is encoded on each surface in a track
- A single surface contains many thousands and thousands of tracks

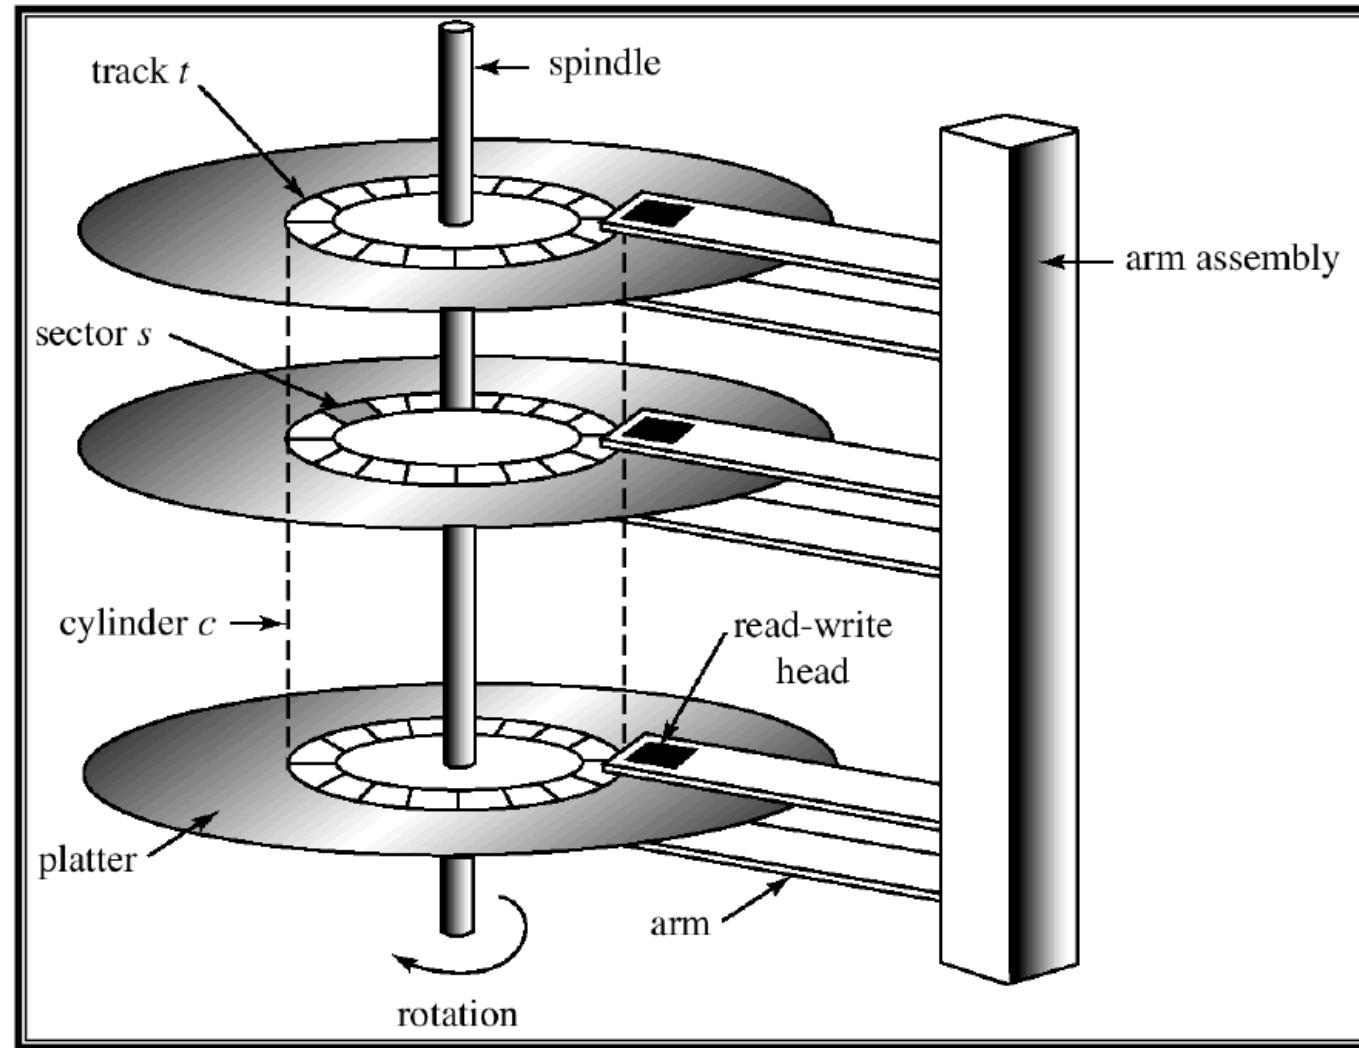
A Simple Disk Drive



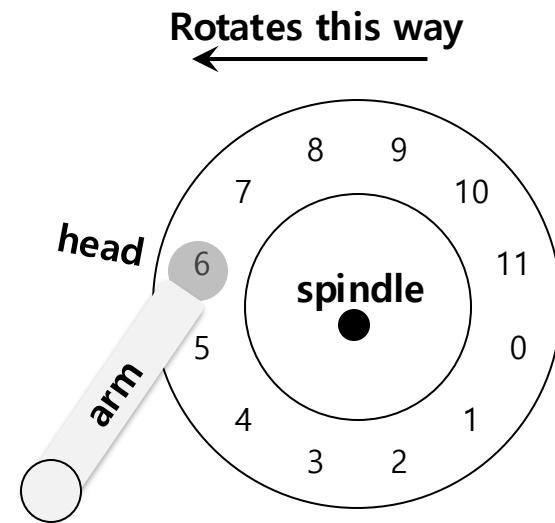
A Single Track Plus A Head

- **Disk head** (One head per surface of the drive)
 - The process of *reading* and *writing* is accomplished by the **disk head**
 - Attached to a single disk arm, which moves across the surface

Example of a Disk



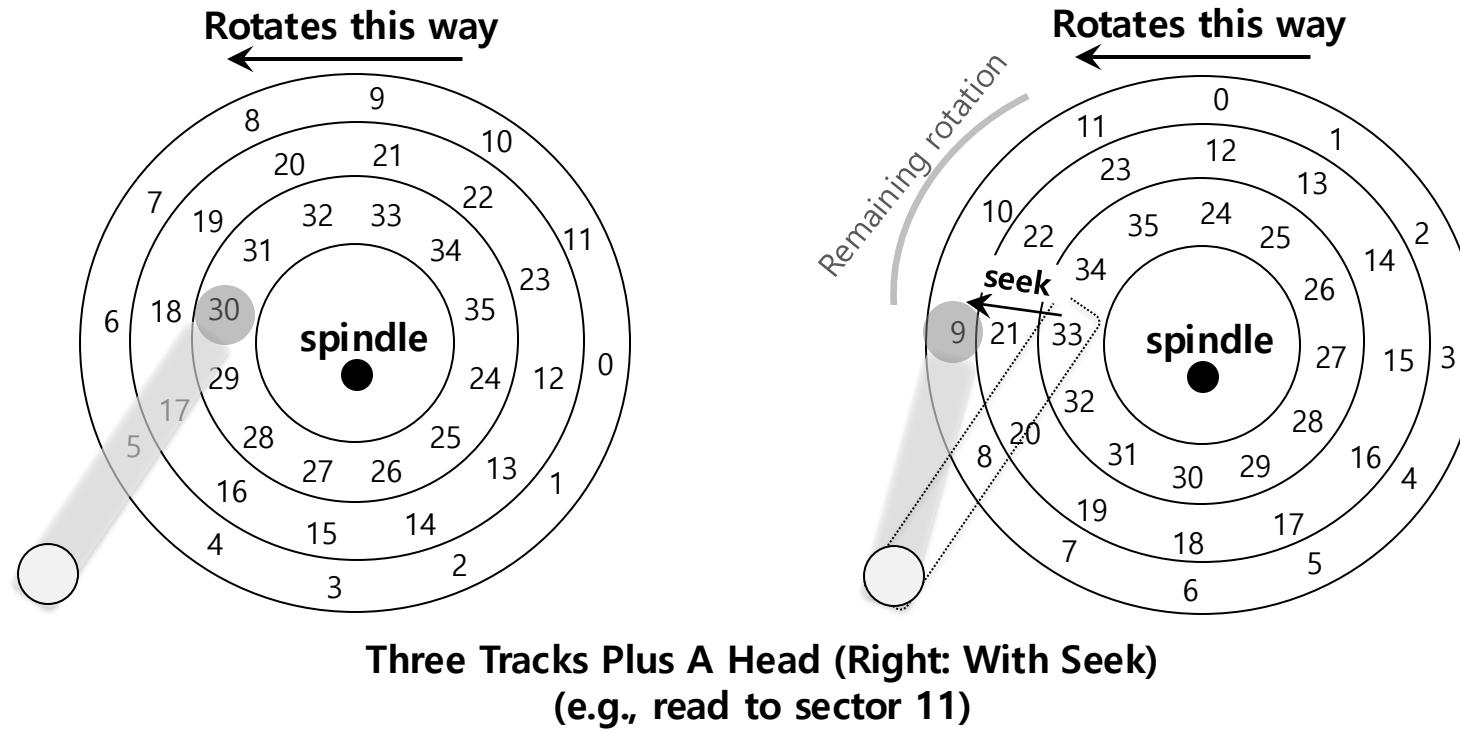
Single-track Latency: The Rotational Delay



A Single Track Plus A Head

- **Rotational delay:** Time for the desired sector to rotate
 - Ex) Full rotational delay is R and we start at sector 6
 - Read sector 0: Rotational delay = $\frac{R}{2}$
 - Read sector 5: Rotational delay = $R-1$ (worst case)

Multiple Tracks: Seek Time



- **Seek:** Move the disk arm to the correct track
 - **Seek time:** Time to move head to the track contain the desired sector
 - One of the most costly disk operations

Phases of Seek

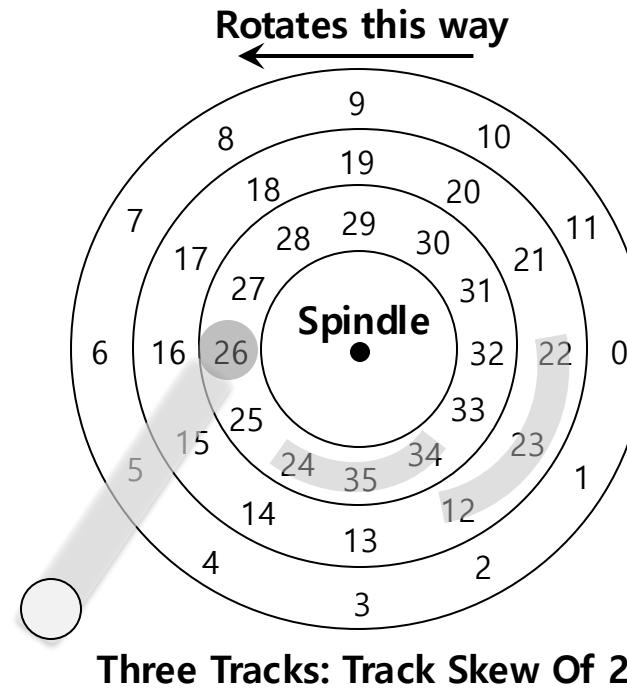
- Acceleration → Coasting → Deceleration → Settling
 - **Acceleration:** The disk arm gets moving
 - **Coasting:** The arm is moving at full speed
 - **Deceleration:** The arm slows down
 - **Settling:** The head is *carefully positioned* over the correct track
 - The settling time is often quite significant, e.g., 0.5 to 2ms

Transfer

- The final phase of I/O
 - Data is either *read from* or *written* to the surface
- Complete I/O time:
 - **Seek**
 - Waiting for the **rotational delay**
 - **Transfer**

Track Skew

- Make sure that sequential reads can be properly serviced **even when crossing track boundaries**



- Without track skew, the head would be moved to the next track but the desired next block would have already rotated under the head.*

Cache (Track Buffer)

- **Hold data** read from or written to the disk
 - Allow the drive to quickly respond to requests
 - Small amount of memory (usually around 8 or 16 MB)

Write on Cache

- **Writeback** (Immediate reporting)
 - Acknowledge a write has completed when it has **put the data in its memory**
 - faster but dangerous
- **Write through**
 - Acknowledge a write has completed after the write has **actually been written to disk**

I/O Time: Doing The Math

- I/O time ($T_{I/O}$):
$$T_{I/O} = T_{seek} + T_{rotation} + T_{transfer}$$

- The rate of I/O ($R_{I/O}$):
$$R_{I/O} = \frac{Size_{Transfer}}{T_{I/O}}$$

	Cheetah 15K.5	Barracuda
Capacity	300 GB	1 TB
RPM	15,000	7,200
Average Seek	4 ms	9 ms
Max Transfer	125 MB/s	105 MB/s
Platters	4	4
Cache	16 MB	16/32 MB
Connects Via	SCSI	SATA

Disk Drive Specs: SCSI Versus SATA

I/O Time Example

- **Random workload:** Issue 4KB read to random locations on the disk
- **Sequential workload:** Read 100MB consecutively from the disk

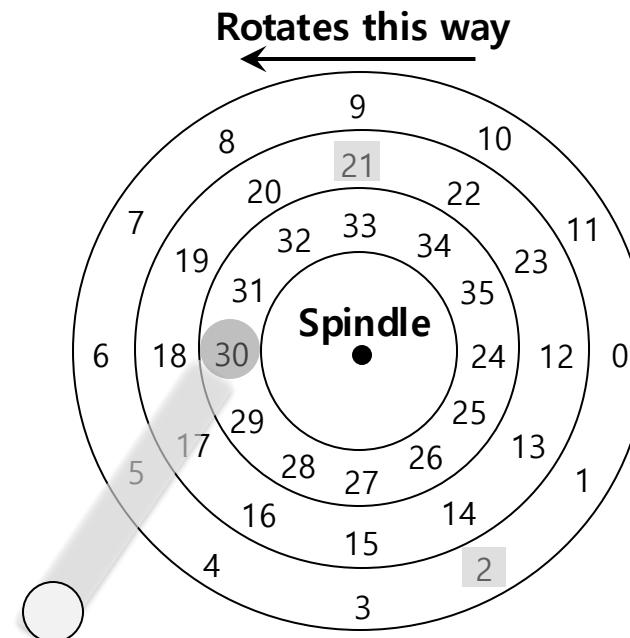
		Cheetah 15K.5	Barracuda
Random	T_{seek}	4 ms	9 ms
	$T_{rotation}$	2 ms	4.2 ms
	$T_{transfer}$	30 microsecs	38 microsecs
	$T_{I/O}$	6 ms	13.2 ms
Sequential	$R_{I/O}$	0.66 MB/s	0.31 MB/s
	$T_{transfer}$	800 ms	950 ms
	$T_{I/O}$	806 ms	963.2 ms
	$R_{I/O}$	125 MB/s	105 MB/s

Disk Drive Performance: SCSI Versus SATA

There is a huge gap in drive performance between **random** and **sequential** workloads

Disk Scheduling

- **Disk Scheduler** decides which I/O request to schedule next
- **SSTF (Shortest Seek Time First)**
 - Order the queue of I/O request by track
 - Pick requests on the nearest track to complete first



SSTF: Scheduling Request 21 and 2
Issue the request to 21 → issue the request to 2

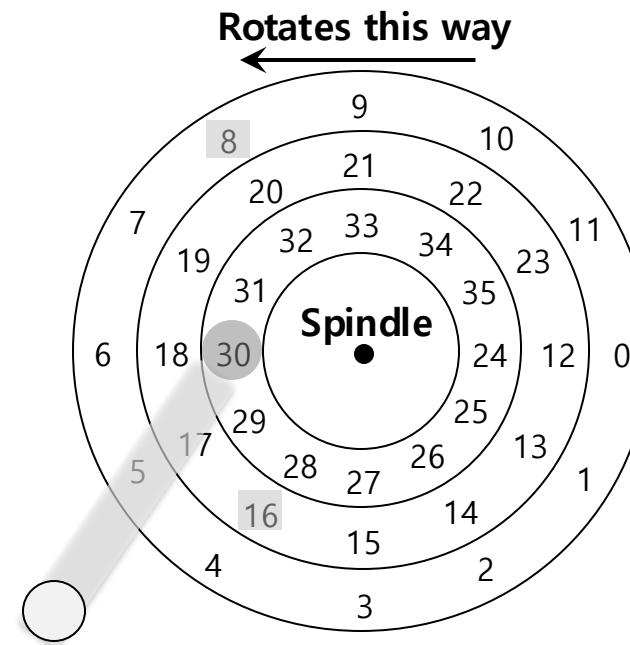
SSTF is not a Panacea

- **Problem 1:** The drive geometry is not available to the host OS
 - Solution: OS can simply implement Nearest-block-first (NBF)
- **Problem 2:** Starvation
 - If there were a steady stream of request to the inner track, request to other tracks would then be ignored completely.

Elevator (a.k.a. SCAN or C-SCAN)

- Move across the disk servicing requests in order across the tracks
 - **Sweep (Elevator)**
 - A single pass across the disk. If a request comes for a block on a track that has already been serviced during this sweep of the disk, it is queued until the next sweep.
 - **F-SCAN:**
 - Freezes the queue to be serviced when it starts a sweep.
 - This helps avoid starvation of far-away requests by ensuring that all requests in the current queue are processed before any new requests are added to the queue.
 - **C-SCAN (Circular SCAN)**
 - Sweeps from the outermost track to the innermost track, and then immediately returns to the outermost track without servicing any requests on the return trip, before starting the next sweep from outer to inner again.
 - This ensures uniform wait times and prevents starvation of requests.

How to Account for Disk Rotation Costs?



SSTF: Sometimes Not Good Enough

- If rotation is faster than seek : request 16 → request 8
 - If seek is faster than rotation : request 8 → request 16

On modern drives, both seek and rotation are roughly equivalent:
Thus, SPTF (Shortest Positioning Time First) is useful

I/O Merging

- **Reduce the number of request** sent to the disk and lowers overhead
 - E.g., read blocks 33, then 8, then 34:
 - The scheduler merge the request for blocks 33 and 34 *into a single two-block request*