

File System Implementation

Unix operating system



Data structures

- The Good Thing
 - Simple and supports the basic abstractions
 - Easy to use file system
- The Problem
 - Terrible performance

Problem of Unix File System

- Unix file system treated the disk as a random-access memory
 - Example of random-access blocks with four files
 - Data blocks for each file can accessed by going back and forth the disk, because they are contiguous
 - File b and d is deleted
 - File E is created with free blocks (spread across the block)
- Other Problem is that the original block size was too small (512 bytes)

FFS: Disk Awareness is Solution

- FFS is Fast File system designed by a group at UC Berkeley
- The design of FFS is that file system structures and allocation policies to be "disk aware" and improve performance
 - Keeps the same APIs with UNIX file system (open(), read(), write(), etc)
 - Changed the internal implementation

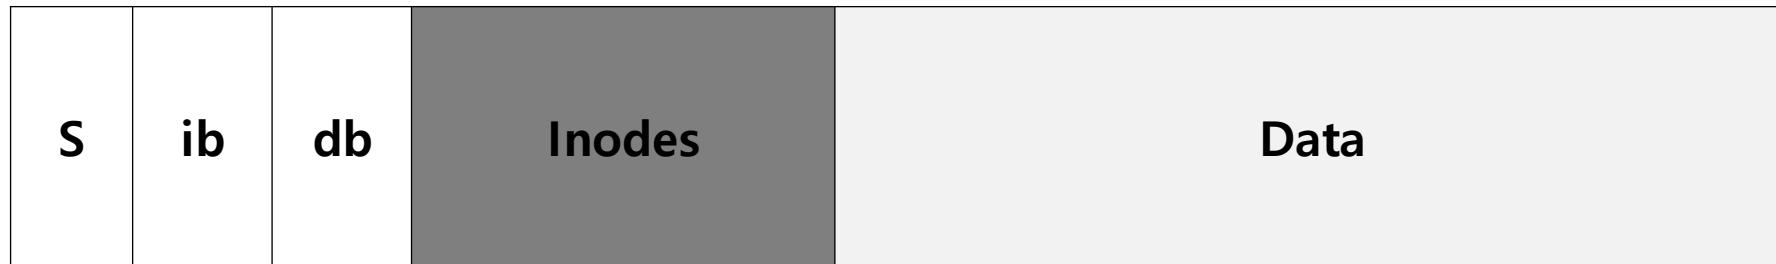
Organizing Structure: Cylinder Group

- FFS divides the disk into a bunch of groups (Cylinder Group)
 - Modern file system call cylinder group as block group

G0	G1	G2	G3	G4	G5	G6	G7	G8	G9
----	----	----	----	----	----	----	----	----	----

- These groups are uses to improve seek performance
 - By placing two files within the same group
 - Accessing one after the other will not be long seeks across the disk
 - FFS needs to allocate files and directories within each of these groups

Organizing Structure: Cylinder Group



- Data structure for each cylinder group
 - A copy of the super block(S) for reliability reason
 - inode bitmap(ib) and data bitmap(db) to track free inode and data block
 - inodes and data block are same to the previous very-simple file system(VSFS)

How To Allocate Files and Directories?

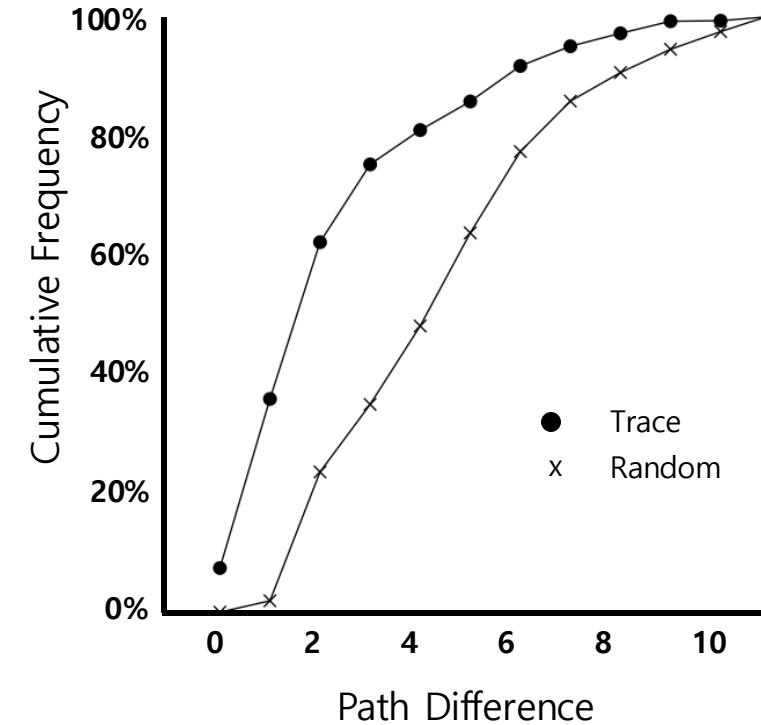
- Policy is “keep related stuff together”
- The placement of directories
 - Find the cylinder group with a low number of allocated directories and a high number of free inodes
 - Put the directory data and inode in that group
- The placement of files
 - Allocate data blocks of a file in the same group as its inode
 - It places all files in the same group as their directory

FFS Locality for SEER Traces

- How “**far away**” file accesses were from one another in the directory tree.

proc/**src**/foo.c
proc/**src**/bar.c
the distance of two file access is 1

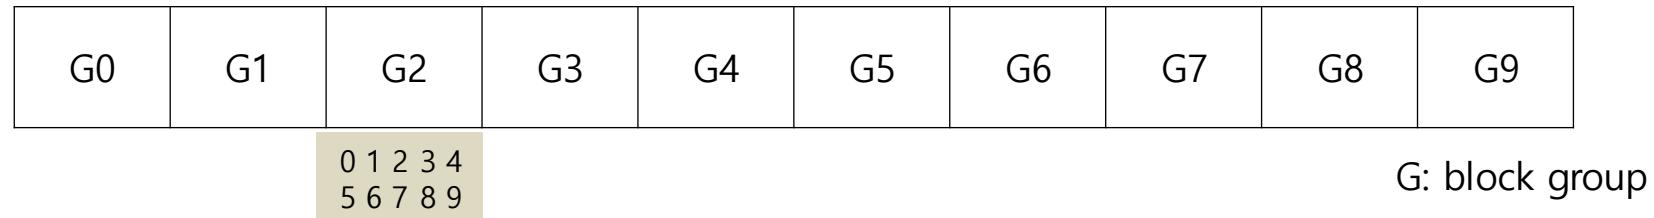
proc/**src**/foo.c
proc/**obj**/foo.o
the distance of two file access is 2



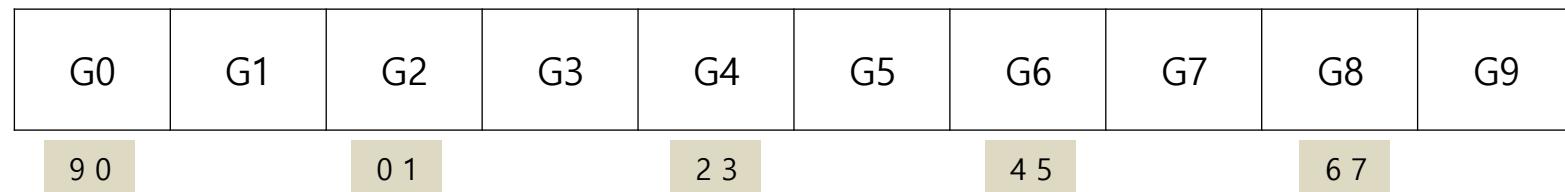
- 7% of file accesses to the same file
- Nearly 40% of file accesses in the same directory
- 25% of file accesses were two distances

Large-File Exception

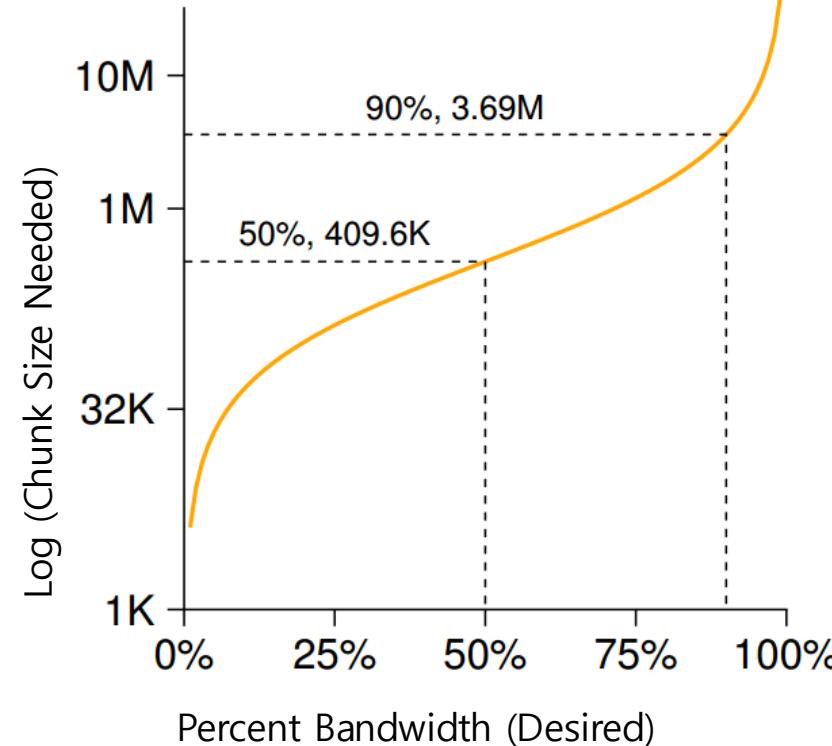
- General policy of file placement
 - Entirely fill the block group it is first place within
 - Hurt file-access locality from “related” file being placed



- For large files, chunks are spread across the disk
 - Hurt performance, but it can be addressed by choosing chunk size
 - Amortization: reducing overhead by doing more work



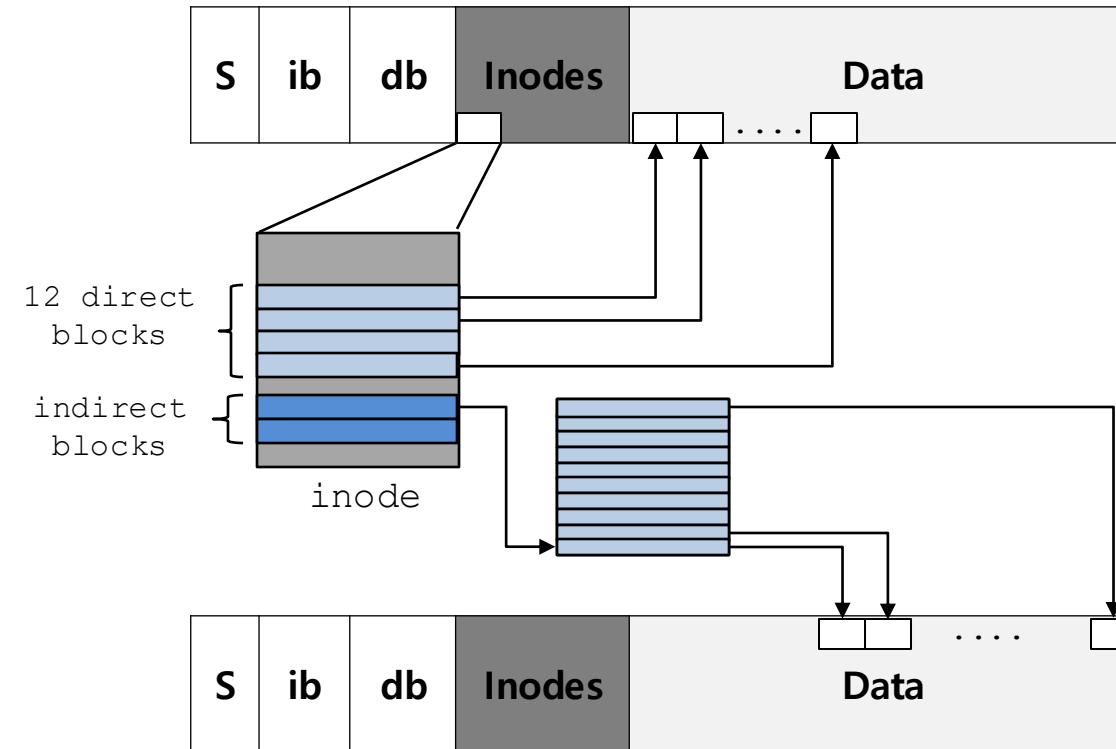
Amortization: How Big Do Chunks Have To Be?



- Computation of the size of chunk
 - Desire 50% of peak disk performance
 - half of time seeking and half of time transferring
 - Disk bandwidth: 40 MB/s
 - Positioning time: 10ms
 - $$\frac{40 \text{ MB}}{\text{sec}} \cdot \frac{1024 \text{ KB}}{1 \text{ MB}} \cdot \frac{1 \text{ sec}}{1000 \text{ ms}} \cdot 10 \text{ ms} = 409.6 \text{ KB}$$
 - Transfer only 409.6 KB every time seeking
 - 90% of peak performance on 3.69MB chunk size

Large-File Exception in FSS

- A simple approach based on the structure of inode
 - Each subsequent **indirect blocks**, and all the **blocks it pointed to**, placed in a **different block group**
 - Every **1024 blocks (4MB)** of the **file in a separate group**



A Few Other Things about FFS

- Internal fragmentation
- Sub-blocks
 - Ex) Create a file with 1 KB : use two sub-blocks, not an entire 4-KB blocks
- Parameterization
- Track buffer
- Long file names
 - Enabling more expressive names in the file system
- Symbolic link