



Crash Consistency


Crash Consistency

- Unlike most data structure, file system data structures must **persist**
 - They must survive over the long haul, stored on devices that retain data despite power loss
- One major challenge faced by a file system is how to update persistent data structure despite the presence of a **power loss or system crash**
- We'll begin by examining the approach taken by older file systems
 - **fsck**(file system checker)
 - **journaling**(write-ahead logging)

A Detailed Example

- Workload

- Append of a single data block(4KB) to an existing file
- open() → lseek() → write() → close()




- Before append a single data block

- single inode is allocated (inode number 2)
- single allocated data block (data block 4)
- The inode is denoted I[v1]

A Detailed Example


- Inside of I[v1] (inode, before update)



- Size of the file is 1 (one block allocated)
- First direct pointer points to block4 (Da)
- All 3 other direct pointers are set to null(unused)

A Detailed Example

- After update



- Data bitmap is updated
- Inode is updated (**I[v2]**)
- New data block is allocated (**Db**)

A Detailed Example

- To achieve the transition, the system perform three separate writes to the disk
 - One each of inode I[v2]
 - Data bitmap B[v2]
 - Data block (Db)
- These writes usually don't happen immediately
 - dirty inode, bitmap, and new data will sit in main memory
 - **page cache or buffer cache**
- If a crash happens after one or two of these write have taken place, but not all three, the file system could be left in a funny state

Crash Scenario

- Imagine only a single write succeeds; there are thus three possible outcomes
 1. Just the data block(Db) is written to disk
 - The data is on disk, but there is no inode
 - Thus, it is as if the write never occurred
 - This case is not a problem at all
 2. Just the updated inode(I[v2]) is written to disk
 - The inode points to the disk address (5, Db)
 - But, the Db has not yet been written there
 - We will read **garbage** data(old contents of address 5) from the disk
 - **Problem : file-system inconsistency**

Crash Scenario

- Imagine only a single write succeeds; there are thus three possible outcomes (Cont.)
 3. Just the updated bitmap ($B[v2]$) is written to disk
 - The bitmap indicates that block 5 is allocated
 - But there is no inode that points to it
 - Thus, the file system is inconsistent again
 - **Problem : space leak**, as block 5 would never be used by the file system

Crash Scenario

- There are also three more crash scenarios. In these cases, two writes succeed and the last one fails
 1. The inode(I[v2]) and bitmap(B[v2]) are written to disk, but not data(Db)
 - The file system metadata is completely consistent
 - **Problem : Block 5 has garbage in it**
 2. The inode(I[v2]) and the data block(Db) are written, but not the bitmap(B[v2])
 - We have the inode pointing to the correct data on disk
 - **Problem : inconsistency between the inode and the old version of the bitmap(B1)**

Crash Scenario

- There are also three more crash scenarios. In these cases, two writes succeed and the last one fails (Cont.)
 3. The bitmap($B[v2]$) and data block(D_b) are written, but not the inode($I[v2]$)
 - **Problem : inconsistency between the inode and the data bitmap**
 - We have no idea which file it belongs to

Crash Consistency Problem

- What we'd like to do ideally is move the file system from one consistent state to another **atomically**
- Unfortunately, we can't do this easily
 - The disk only commits one write at a time
 - Crashes or power loss may occur between these updates
- We call this general problem the **crash-consistency problem**

Solution #1: File System Checker

File System Checker

- The File System Checker (**fsck**)

- `fsck` is a Unix tool for finding inconsistencies and repairing them
- `fsck` check super block, Free block, Inode state, Inode links, etc
- Such an approach can't fix all problems
 - example : The file system looks consistent but the inode points to garbage data
- The only real goal is to make sure the file system metadata is internally consistent

File System Checker

- Basic summary of what fsck does:
 - **Superblock**
 - fsck first checks if the superblock looks reasonable
 - » Sanity checks : file system size > number of blocks allocated
 - Goal : to find suspect superblock
 - In this case, the system may decide to use an alternate copy of the superblock
 - **Free blocks**
 - fsck scans the inodes, indirect blocks, double indirect blocks, and so on
 - The only real goal is to make sure the file system metadata is internally consistent

File System Checker

- Basic summary of what `fsck` does: (Cont.)

- **Inode state**

- Each inode is checked for corruption or other problem
 - » Example : type checking(regular file, directory, symbolic link, etc)
 - If there are problems with the inode fields that are not easily fixed
 - » The inode is considered suspect and cleared by `fsck`

- **Inode Links**

- `fsck` also verifies the link count of each allocated inode
 - » To verify the link count, `fsck` scans through the entire directory tree
 - If there is a mismatch between the newly-calculated count and that found within an inode, corrective action must be taken
 - » Usually by fixing the count with in the inode

File System Checker

- Basic summary of what fsck does: (Cont.)
 - **Inode Links** (Cont.)
 - If an allocated inode is discovered but no directory refers to it, it is moved to the lost+found directory
 - **Duplicates**
 - fsck also checks for duplicated pointers
 - Example : Two different inodes refer to the same block
 - » If one inode is obviously bad, it may be cleared
 - » Alternately, the pointed-to block could be copied

File System Checker

- Basic summary of what fsck does: (Cont.)

- **Bad blocks**

- A check for bad block pointers is also performed while scanning through the list of all pointers
 - A pointer is considered “bad” if it obviously points to something outside its valid range
 - Example : It has an address that refers to a block greater than the partition size
 - » In this case, fsck can't do anything too intelligent; it just removes the pointer

File System Checker

- Basic summary of what `fsck` does: (Cont.)
 - **Directory checks**
 - `fsck` does not understand the contents of user files
 - » However, directories hold specifically formatted information created by the file system itself
 - » Thus, `fsck` performs additional integrity checks on the contents of each directory
 - Example
 - » making sure that “.” and “..” are the first entries
 - » each inode referred to in a directory entry is allocated?
 - » ensuring that no directory is linked to more than once in the entire hierarchy

File System Checker

- Building a working `fsck` requires intricate knowledge of the filesystem
- `fsck` have a bigger and fundamental problem: **too slow**
 - scanning the entire disk may take many minutes or hours
 - Performance of `fsck` became prohibitive
 - as disk grew in capacity and RAIDs grew in popularity
- At a higher level, the basic premise of `fsck` seems just a tad irrational
 - It is incredibly expensive to scan the entire disk
 - It works but is wasteful
 - Thus, as disk(and RAIDs) grew, researchers started to look for other solutions