



File System Journaling

Journaling

▪ Journaling (Write-Ahead Logging)

- When updating the disk, before over writing the structures in place, first write down a little note describing what you are about to do
- Writing this note is the “write ahead” part, and we write it to a structure that we organize as a “log”
- By writing the note to disk, you are guaranteeing that if a crash takes places during the update of the structures you are updating, you can go back and look at the note you made and try again
- Thus, you will know exactly what to fix after a crash, instead of having to scan the entire disk

Journaling

- We'll describe how Linux ext3 incorporates journaling into the file system
 - Most of the on-disk structures are identical to Linux ext2
 - The new key structure is the journal itself
 - It occupies some small amount of space within the partition or on another device



Fig.1 Ext2 File system structure




Fig.2 Ext3 File system structure

Data Journaling

- Data journaling is available as a mode with the ext3 file system
- Example : our canonical update again
 - We wish to update inode ($I[v2]$), bitmap ($B[v2]$), and data block (D_b) to disk
 - Before writing them to their final disk locations, we are now first going to write them to the log(a.k.a. journal)

Data Journaling

- Example : our canonical update again (Cont.)



- TxB: Transaction begin block
 - It contains some kind of transaction identifier(TID)
- The middle three blocks just contain the exact content of the blocks themselves
 - This is known as physical logging
- TxE: Transaction end block
 - Marker of the end of this transaction
 - It also contain the TID

Data Journaling



■ Checkpoint

- Once this transaction is safely on disk, we are ready to overwrite the old structures in the file system
- This process is called **checkpointing**
- Thus, to checkpoint the file system, we issue the writes I[v2], B[v2], and Db to their disk locations

Data Journaling

- Our initial sequence of operations:

1. Journal write

- Write the transaction to the log and wait for these writes to complete
- TxB, all pending data, metadata updates, TxE

2. Checkpoint

- Write the pending metadata and data updates to their final locations

Data Journaling

- When a crash occurs during the writes to the journal

1. Transaction each one at a time

- 5 transactions (TxB, I[v2], B[v2], Dnb, TxE)
 - This is slow because of waiting for each to complete

2. Transaction all block writes at once


- Five writes -> a single sequential write : Faster way
 - However, this is unsafe
 - » Given such a big write, the disk internally may perform scheduling and complete small pieces of the big write in any order

Data Journaling

- When a crash occurs during the writes to the journal

2. Transaction all block writes at once (Cont.)

- Thus, the disk internally may (1) write TxB, I[v2], B[v2], and TxE and only later (2) write Db
- Unfortunately, if the disk loses power between (1) and (2)



Data Journaling

- When a crash occurs during the writes to the journal
 - Transaction all block writes at once (Cont.)
 - To avoid this problem, the file system issues the transactional write in two steps
 - First, writes all blokes except the TxE block to journal
- Second, 
 - An important aspect of this process is the atomicity guarantee provided by the disk
 - The disk guarantees that any 512-byte write either happen or not
 - Thus, TxE should be a single 512-byte block

Data Journaling

- When a crash occurs during the writes to the journal
 - Transaction all block writes at once (Cont.)
 - Thus, our current protocol to update the file system, with each of its three phases labeled:
 - » Journal write : write the contents of the transaction to the log
 - » Journal commit (added) : write the transaction commit block
 - » Checkpoint : write the contents of the update to their locations

Data Journaling




- Recovery
 - If the crash happens before the transactions is written to the log
 - The pending update is skipped
 - If the crash happens after the transactions is written to the log, but before the checkpoint
 - Recover the update as follow:
 - » Scan the log and lock for transactions that have committed to the disk
 - » Transactions are replayed

Batching Log Updates

- If we create two files in the same directory, the same inode, directory entry block is written to the log and committed twice
- To reduce excessive write traffic to disk, journaling manages the **global transaction**
 - Write the content of the global transaction forced by synchronous request
 - Write the content of the global transaction after timeout of 5 seconds

Making The log Finite

- The log is of a finite size
 - Two problems arise when the log becomes full




1. The larger the log, the longer recovery will take
 - Simpler but less critical
2. No further transactions can be committed to the disk
 - Thus making the file system “less than useful”

Making The log Finite

- To address these problems, journaling file systems treat the log as a circular data structure, re-using it over and over
 - This is why the journal is referred to as a circular log.
- To do so, the file system must take action some time after a checkpoint
 - Specifically, once a transaction has been checkpointed, the file system should free the space

Making The log Finite

- journal super block
 - Mark the oldest and newest transactions in the log.
 - The journaling system records which transactions have not been check pointed



Making The log Finite


- journal super block (Cont.)
 - Thus, we add another step to our basic protocol
 1. Journal write
 2. Journal commit
 3. checkpoint
 - 4. Free**
 - » Some time later, mark the transaction free in the journal by updating the journal Superblock

Metadata Journaling

- There is a still problem : writing every data block to disk **twice**
 - Commit to log (journal file)
 - Checkpoint to on-disk location
- People have tried a few different things in order to speed up performance
 - Example : A simpler form of journaling is called **ordered journaling (metadata journaling)**
 - User data is not written to the journal

Metadata Journaling

- Thus, the following information would be written to the journal



- The data block Db, previously written to the log, would instead be written to the file system proper

Metadata Journaling


- The modification does raise an interesting question: **when should we write data blocks to disk?**
- Let's consider an example
 1. Write Data to disk after the transaction
 - Unfortunately, this approach has a problem
 - The file system is consistent but $I[v2]$ may end up pointing to garbage data
 2. Write Data to disk before the transaction
 - It ensures the problems

Metadata Journaling


- Specifically, the protocol is as follows:
 1. **Data Write(added)**: Write data to final location
 2. **Journal metadata write(added)**: Write the begin and metadata to the log
 3. Journal commit
 4. Checkpoint metadata
 5. Free

Tricky case: Block Reuse

- Some metadatas should not be replayed
- Example
 - 1. Directory "foo" is updated.



- 2. Directory "foo" is deleted. block 1000 is freed up for reuse
- 3. User Create a file "foobar", reusing block 1000 for data



Tricky case: Block Reuse

- Now assume a crash occurs and all of this information is still in the log

Journal	TxB id=1	I[foo] ptr:1000	D[foo] [final addr:1000]	TxE id=1	TxB id=2	I[foobar] ptr:1000	TxE id=2	→
---------	-------------	--------------------	-----------------------------	-------------	-------------	-----------------------	-------------	---

- During replay, the recovery process replays everything in the log
 - Including the write of directory data in block 1000
- The replay thus overwrites the user data of current file `foobar` with old directory contents

Tricky case: Block Reuse

- Solution

- What Linux ext3 does instead is to add a **new type** of record to the journal, Known as a **revoke** record
- When replaying the journal, the system first scans for such revoke records
- Any such revoked data is never replayed

Data Journaling Timeline

Journal contents			TxE	File System	
TxB	(metadata)	(data)		Metadata	Data
issue	issue	issue			
complete					
	complete				
		complete			
			issue		
			complete		
				issue	
				complete	
					issue
					complete

Data Journaling Timeline

Metadata Journaling Timeline

