Flash Memory and SSDs

Memory Types

FLASH

- High-density
- Low-cost
- High-speed
- Low-power
- High reliability

DRAM

- High-density
- Low-cost
- High-speed
- High-power

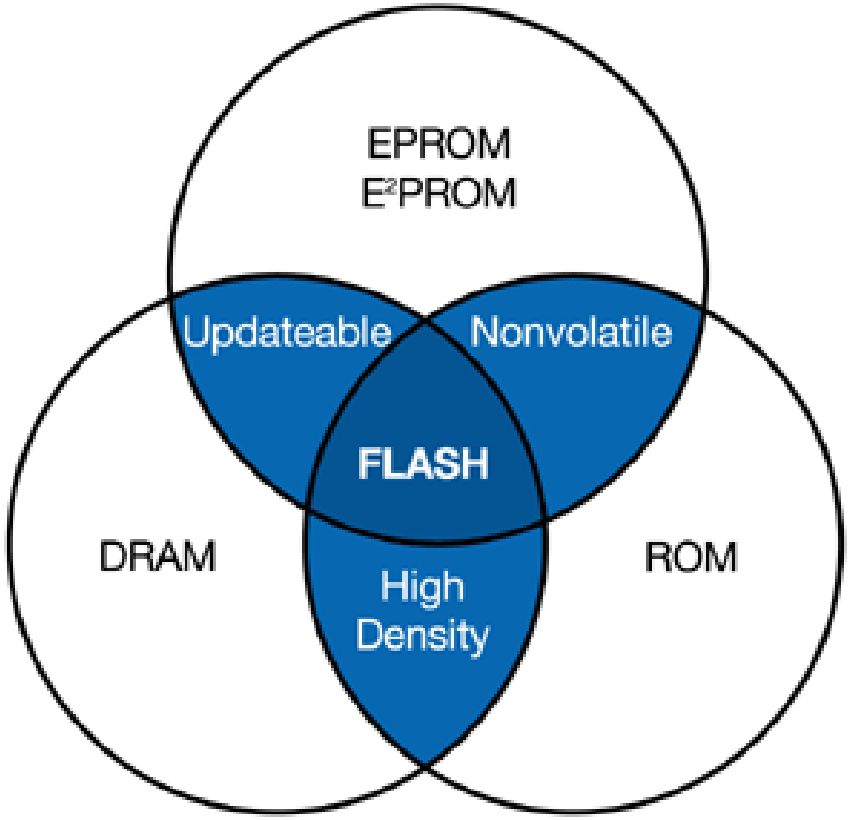

EPROM

- Non-volatile
- High-density
- Ultraviolet light for erasure

EEPROM

- Non-volatile
- Lower reliability
- Higher cost
- Lowest density
- Electrically byte-erasable

ROM

- High-density
- Reliable
- Low-cost
- Suitable for high production with stable code

Source: Intel Corporation.

Flash Memory Characteristics

- Erase-before-write
  - Read
  - Write or Program: 1 → 0
  - Erase: 0 → 1
- Bulk erase
  - Program unit:
    - NOR: byte or word
    - NAND: page
  - Erase unit: block

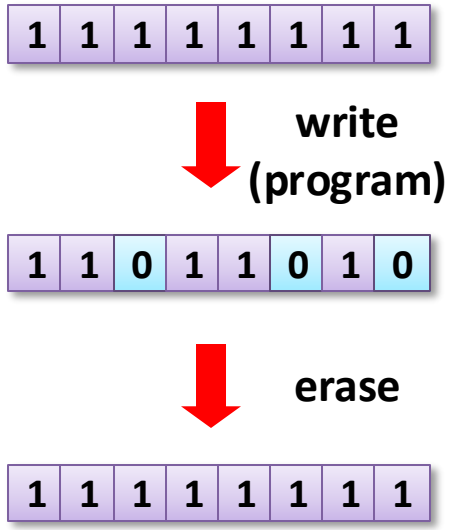

Logical View of NAND Flash

- A collection of blocks
- Each block has a number of pages
- The size of a block or a page depends on the technology (but, it's getting larger)

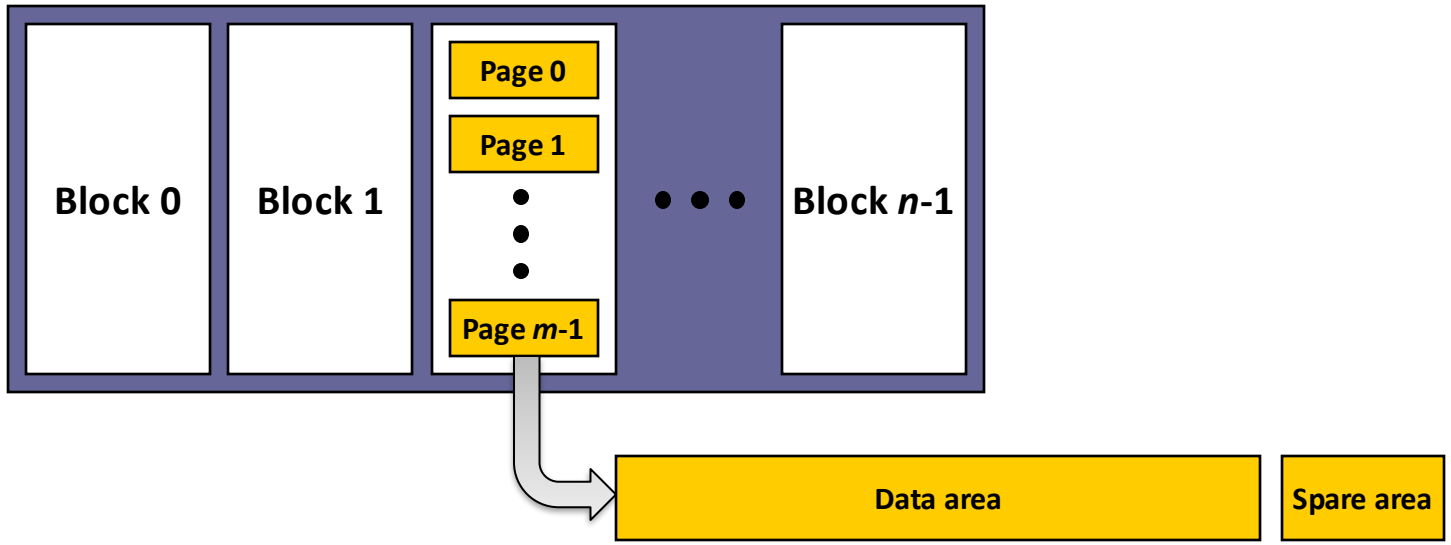

NAND Flash Types

- SLC NAND
  - Single Level Cell
  - 1 bit/cell
- MLC NAND
  - Multi-level Cell (misnomer)
  - 2 bits/cell
- TLC NAND
  - Triple-level Cell
  - 3 bits/cell
- QLC NAND
  - Quad-level Cell
  - 4 bits/cell

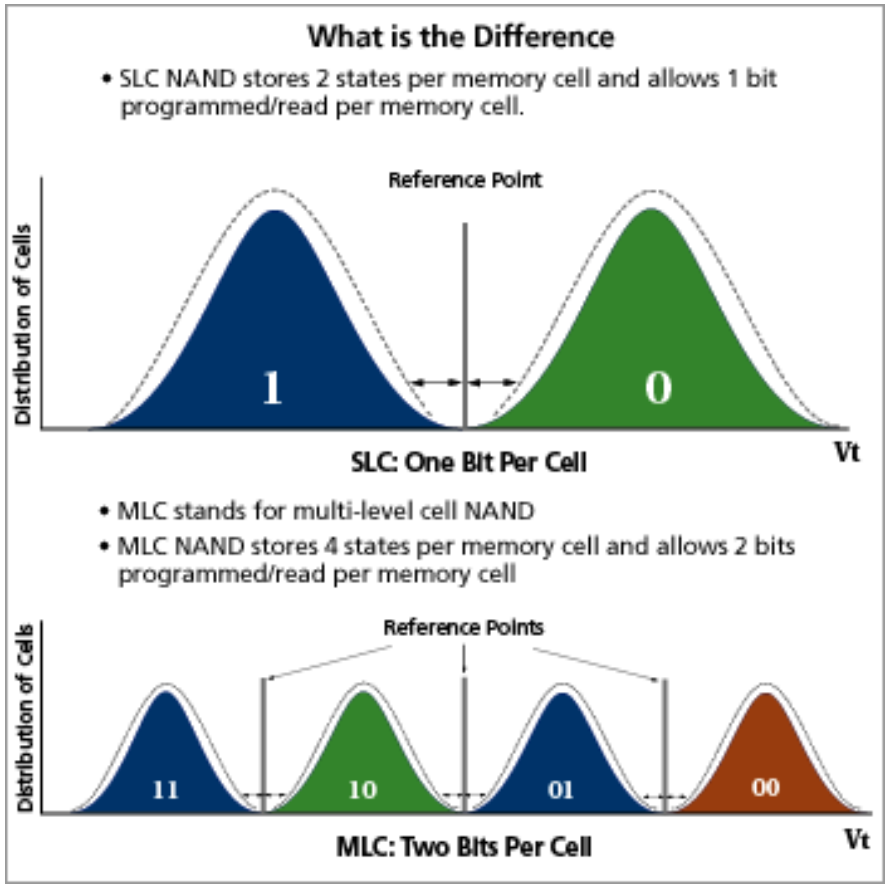

Source: Micron Technology, Inc.

NAND Applications

- USB Flash Drives
- Flash cards
  - CompactFlash, MMC, SD, Memory stick, …
- Smartphones
  - eMMC (Embedded MMC)
  - UFS (Universal Flash Storage)
- SSDs (Solid State Drives)
- Other embedded devices
  - MP3 players, Digital TVs, Set-top boxes, Car navigators, …

Anatomy of an SSD

- Samsung 850 Evo

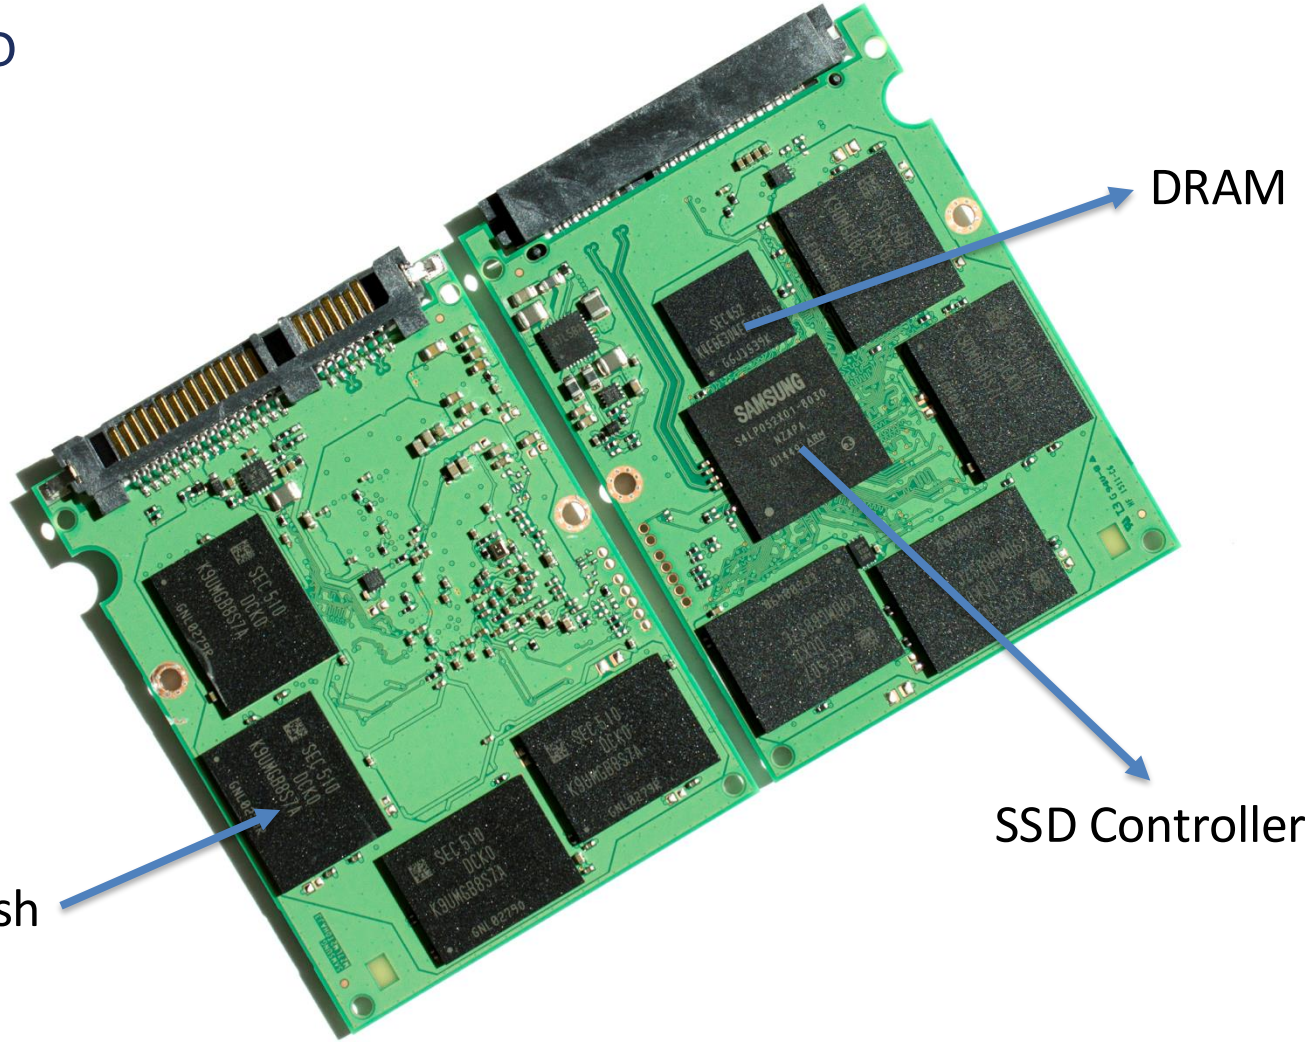

http://www.anandtech.com/show/9451/the-2tb-samsung-850-pro-evo-ssd-review

SSD Internals

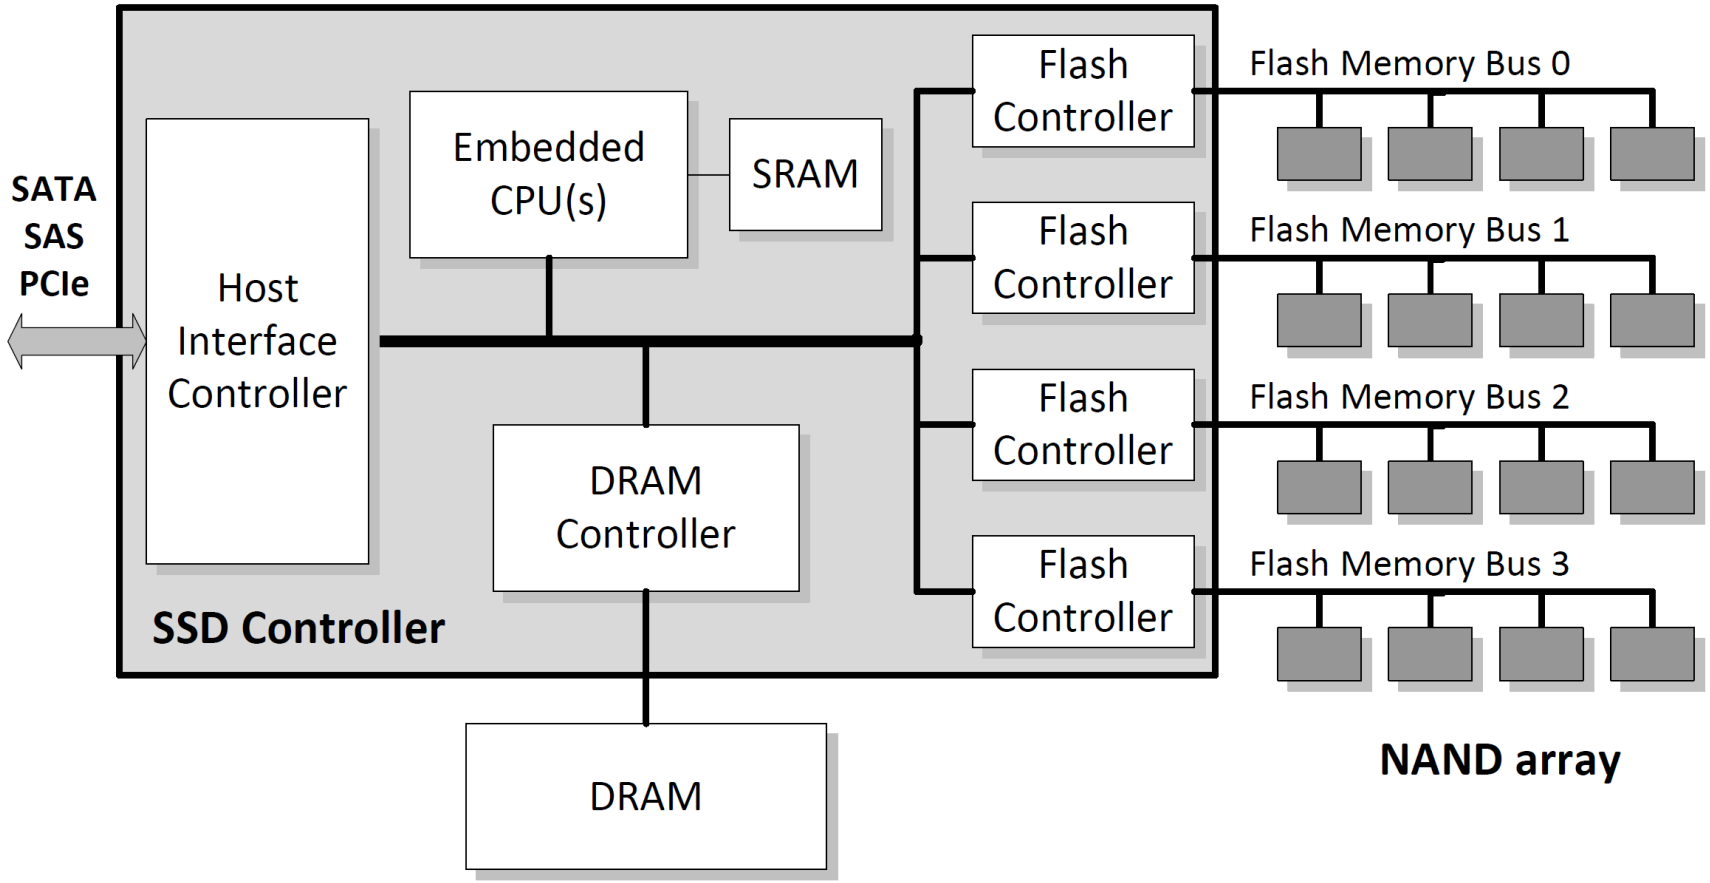

HDDs vs. SSDs

Feature	SSD (Samsung)	HDD (Seagate)
Model	MZ-75E2T0B (850 Evo)	ST2000LM003 (SpinPoint M9T)
Capacity	2TB (128Gb 32-Layer 3D V-NAND TLC x 16 die/channel x 8 channels)	2TB (3 Discs, 6 Heads, 5400 RPM)
Form factor	2.5”, 66g	2.5”, 130g
DRAM	2 GB	32 MB
Host interface	SATA-3 (6.0 Gbps)	SATA-3 (6.0 Gbps)
Power consumption (Active / Idle / Sleep)	3.7, 4.7 W / 0.5 W / 0.05 W	2.3 W / 0.7 W / 0.18 W
Performance **850 Evo**[1]: Sequential: 128KB/QD2 Random: 4KB/QD32 **M9T**[2]: Sequential: 2MB Random: 4KB	Sequential read: 544 MB/s Sequential write: 520 MB/s Random read: 97,687 IOPS Random write: 89,049 IOPS Random read: 11,335 IOPS (QD1) Random write: 38,433 IOPS (QD1)	Sequential read: 124 MB/s Sequential write: 124 MB/s Random read: 56 IOPS Random write: 98 IOPS Power-on to ready: 3.5 sec Average seek: 12/14 ms Average latency: 5.6 ms
Price[3]	1,009,380 won (505won/GB)	117,060 won (59won/GB)

[1] *http://www.tomshardware.com/reviews/samsung-850-evo-850-pro-2tb-ssd,4205.html*
[1] *http://www.storagereview.com/samsung_spinpoint_m9t_hard_drive_review* [3] *http://www.enuri.com (As of Sep. 27, 2015)*

NAND Constraints

- No in-place update
  - Require sector remapping (or address translation)
- Bit errors
  - Require the use of error correction codes (ECCs)
- Bad blocks
  - Factory-marked and run-time bad blocks
  - Require bad block remapping
- Limited program/erase cycles
  - < 100K for SLCs, < 3K for MLCs, < 1K for TLCs
  - Require wear-leveling

Flash Translation Layer (FTL)

- A software layer to make NAND flash fully emulate traditional block devices (e.g. disks)

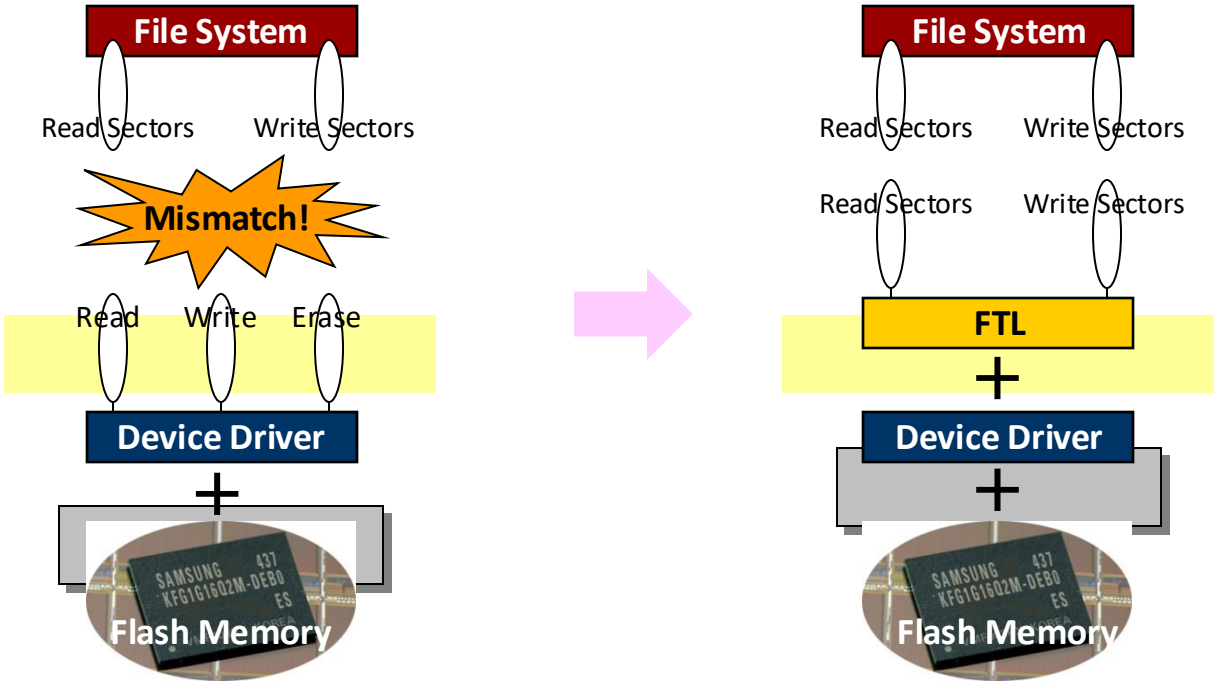

Source: Zeen Info. Tech.

Address Mapping

- Required since flash pages cannot be overwritten

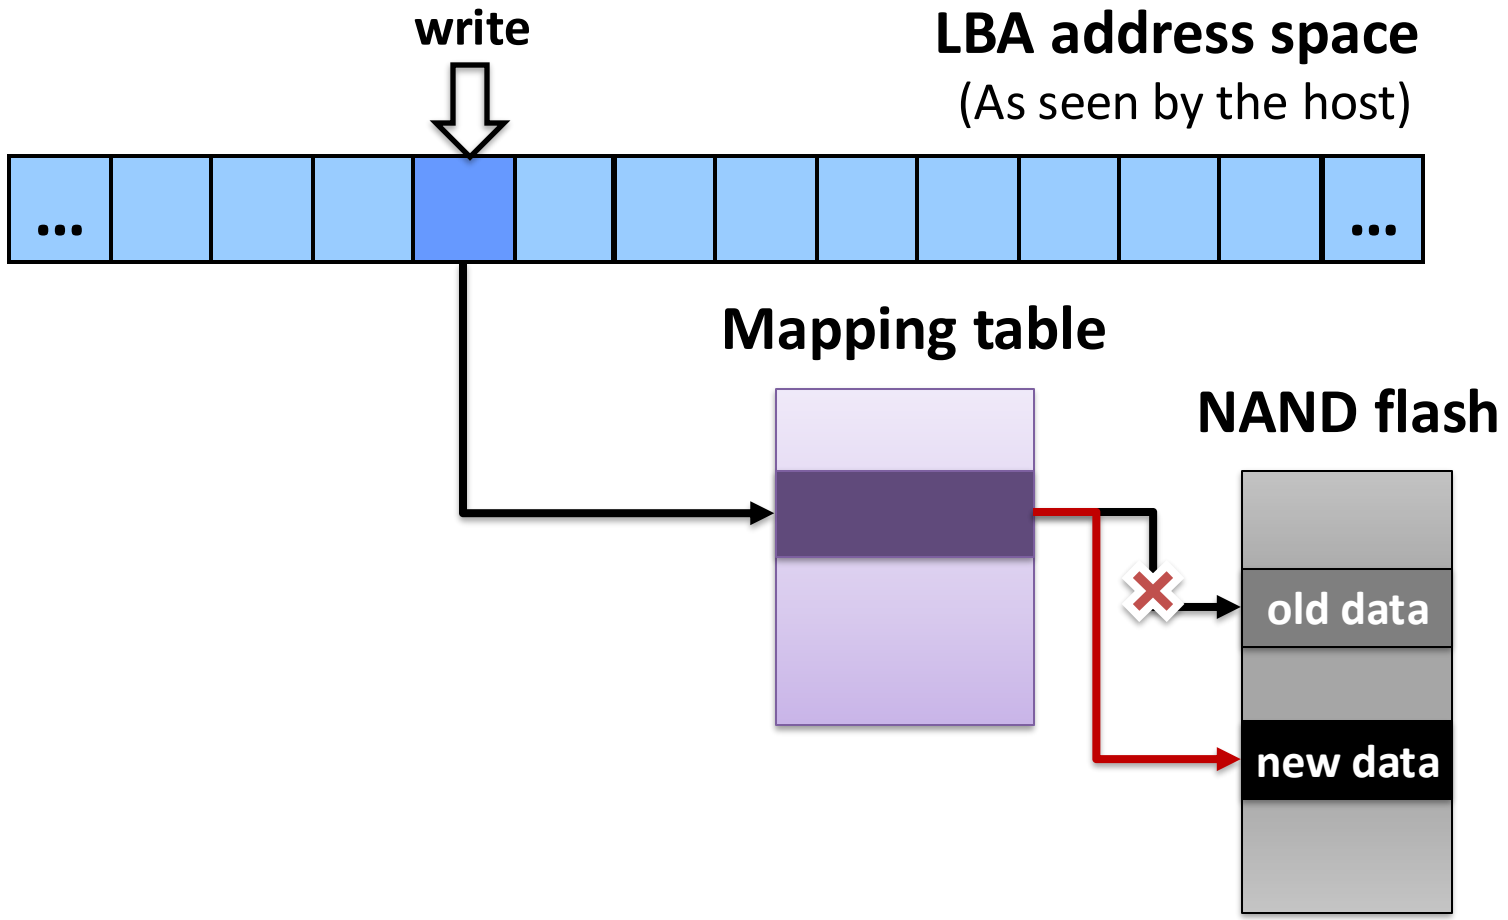

Example: Page Mapping

- Flash configuration
  - Page size: 4KB
  - # of pages / block = 4
- Current state
  - Written to page 0, 1, 2, 8, 4, 5
- Reading page 5

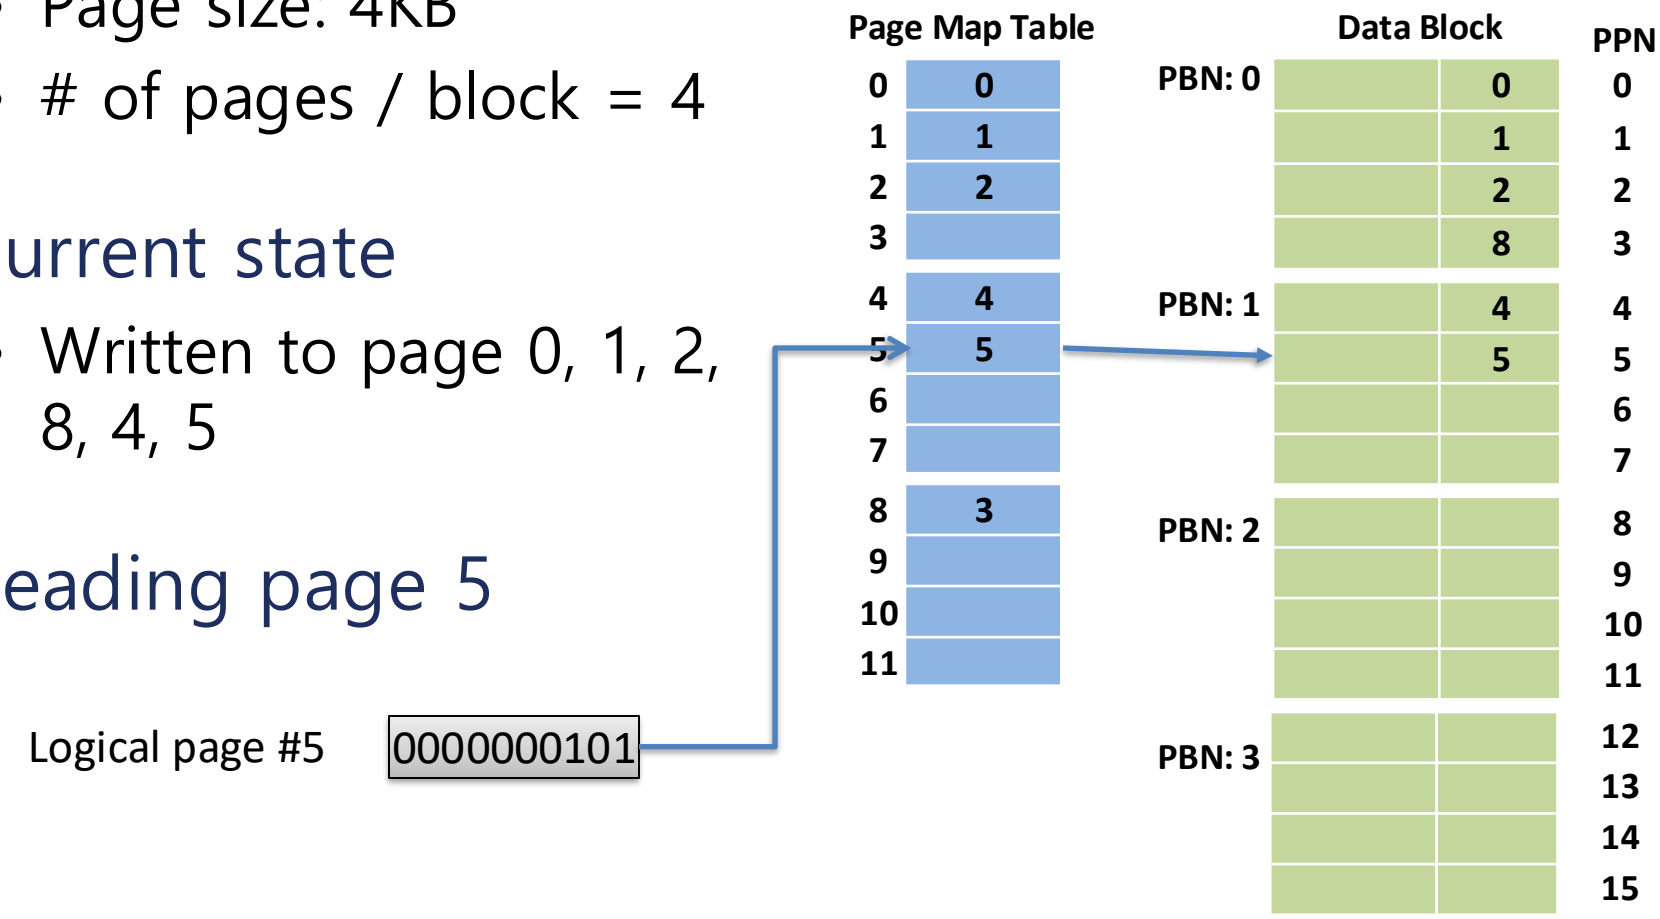

Example: Page Mapping

- Flash configuration
  - Page size: 4KB
  - # of pages / block = 4
- Current state
  - Written to page 0, 1, 2, 8, 4, 5
- New requests (in order)
  - Write to page 9
  - Write to page 3
  - Write to page 5

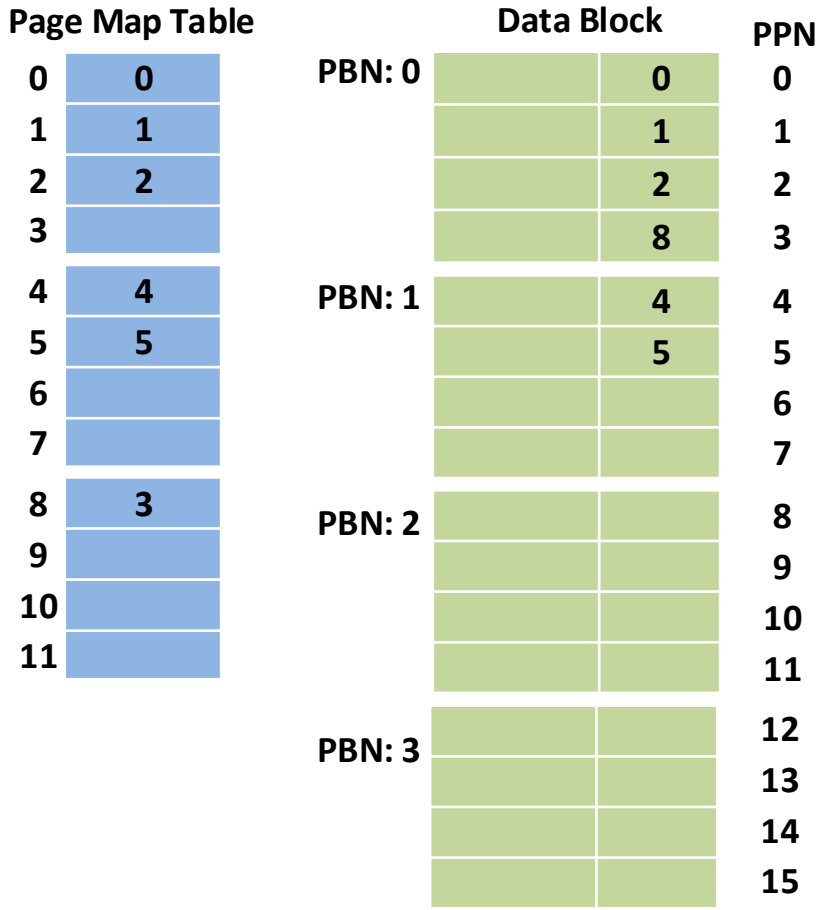

Example: Page Mapping

- Flash configuration
  - Page size: 4KB
  - # of pages / block = 4
- Current state
  - Written to page 0, 1, 2, 8, 4, 5
- New requests (in order)
  - Write to page 9
  - Write to page 3
  - Write to page 5

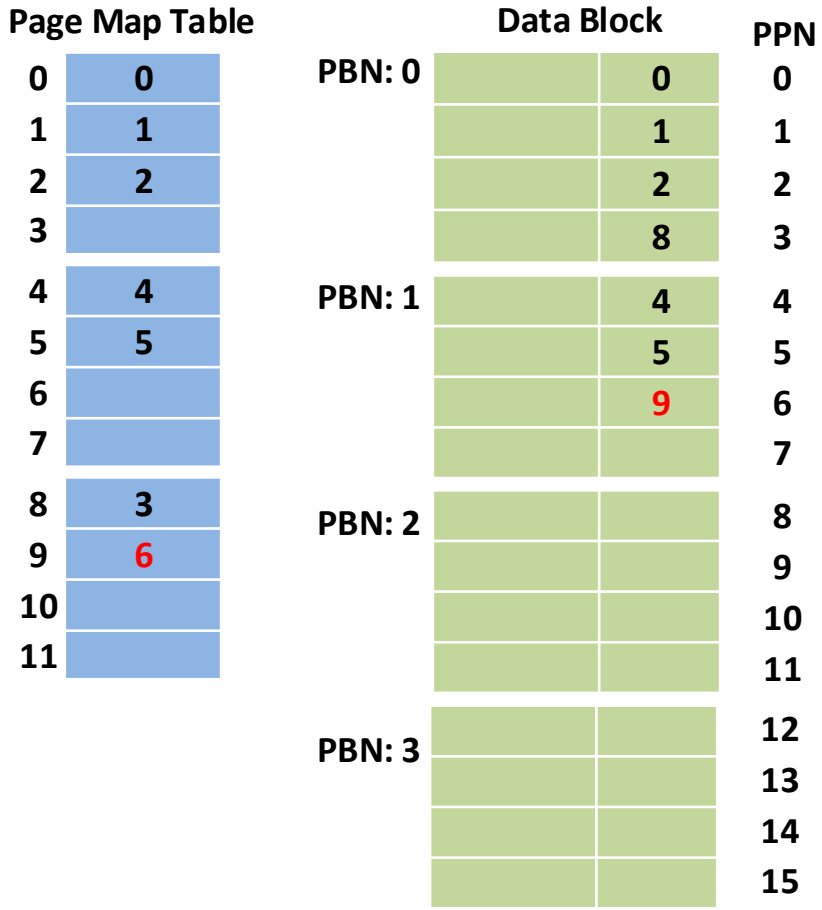

Example: Page Mapping

- Flash configuration
  - Page size: 4KB
  - # of pages / block = 4
- Current state
  - Written to page 0, 1, 2, 8, 4, 5
- New requests (in order)
  - Write to page 9
  - Write to page 3
  - Write to page 5

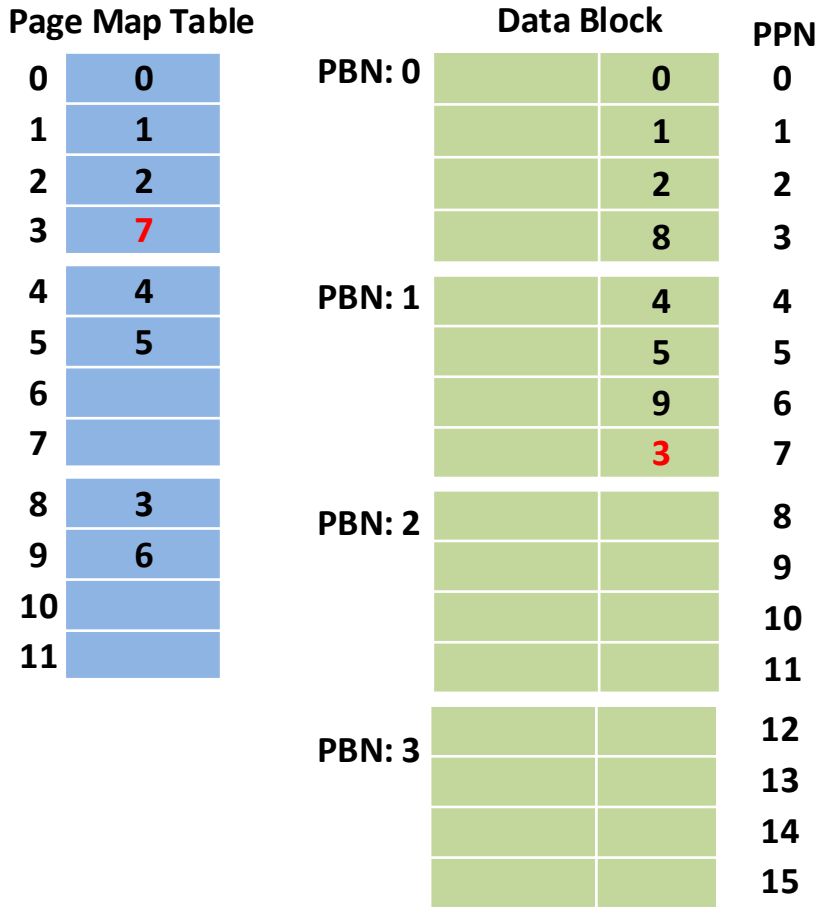

Example: Page Mapping

- Flash configuration
  - Page size: 4KB
  - # of pages / block = 4
- Current state
  - Written to page 0, 1, 2, 8, 4, 5
- New requests (in order)
  - Write to page 9
  - Write to page 3
  - Write to page 5

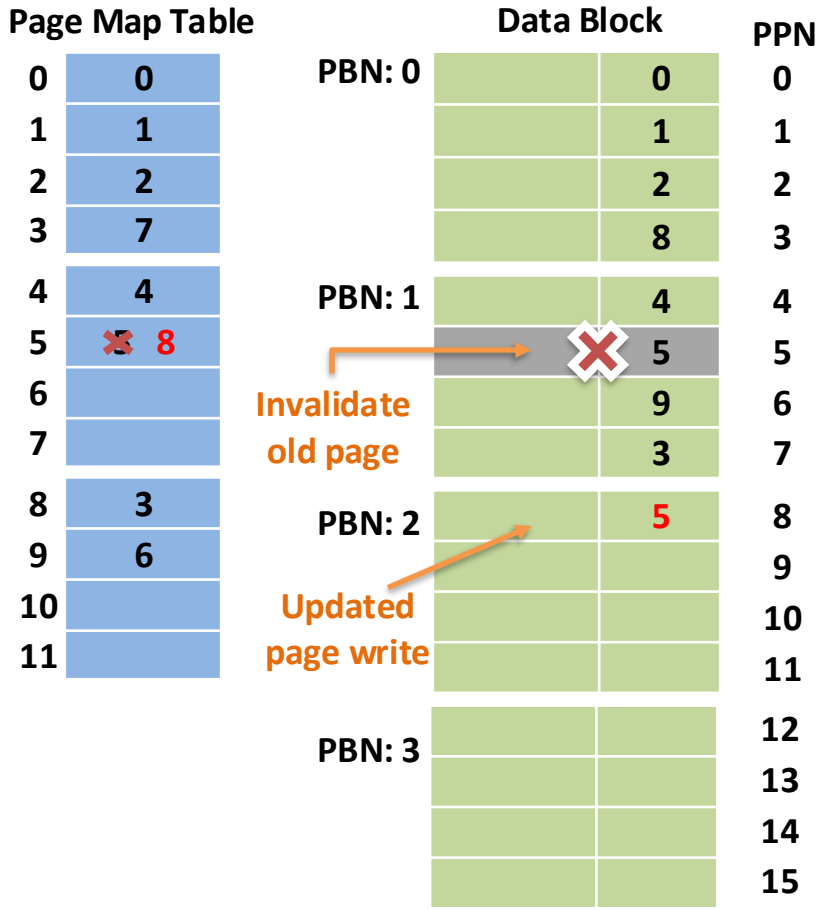

Garbage Collection

- Garbage collection (GC)
  - Eventually, FTL will run out of blocks to write to
  - GC must be performed to reclaim free space
  - Actual GC procedure depends on the mapping scheme

- GC in page-mapping FTL
  - Select victim block(s)
  - Copy all valid pages of victim block(s) to free block
  - Erase victim block(s)
  - Note: At least one free block should be reserved for GC

Example: GC in Page Mapping

- Current state
  - Written to page 0, 1, 2, 8, 4, 5
  - Written to page 9, 3, 5
- New requests (in order)
  - Write to page 8
  - Write to page 9
  - Write to page 3
  - Write to page 1
  - Write to page 4

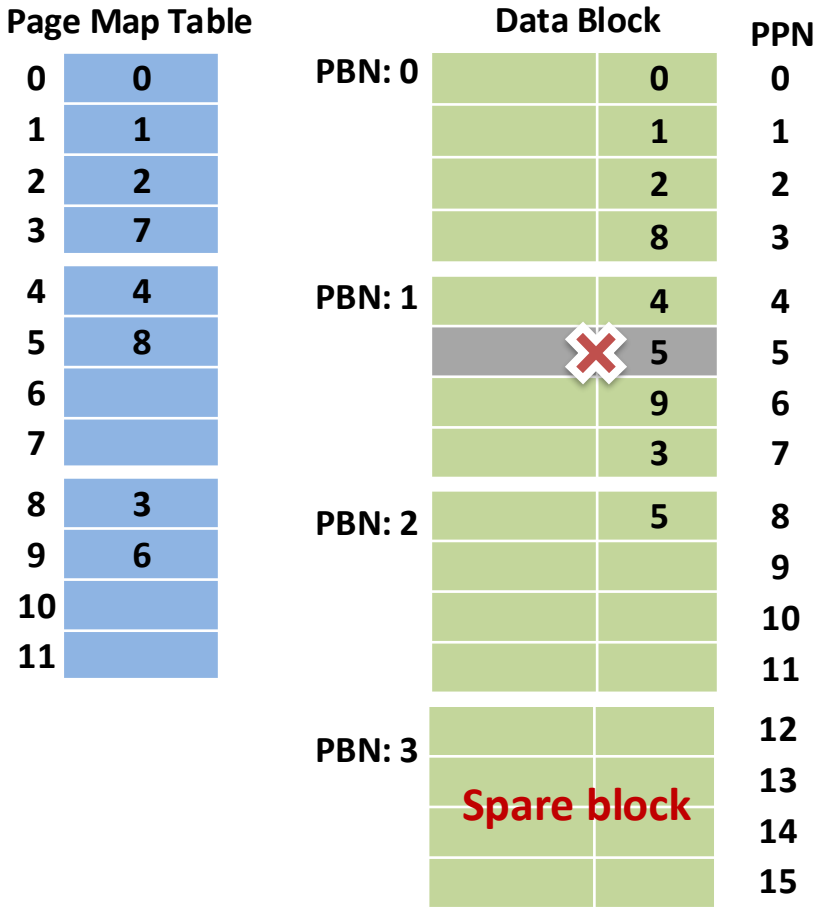

Example: GC in Page Mapping

- Current state
  - Written to page 0, 1, 2, 8, 4, 5
  - Written to page 9, 3, 5
- New requests (in order)
  - Write to page 8
  - Write to page 9
  - Write to page 3
  - Write to page 1
  - Write to page 4

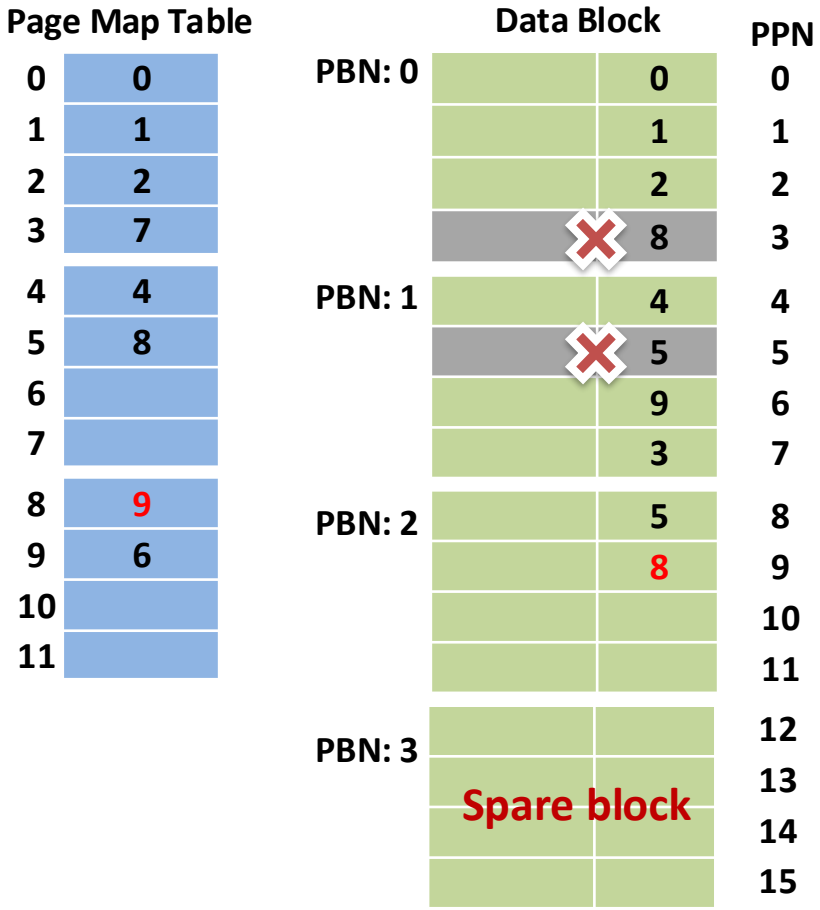

Example: GC in Page Mapping

- Current state
  - Written to page 0, 1, 2, 8, 4, 5
  - Written to page 9, 3, 5
- New requests (in order)
  - Write to page 8
  - Write to page 9
  - Write to page 3
  - Write to page 1
  - Write to page 4

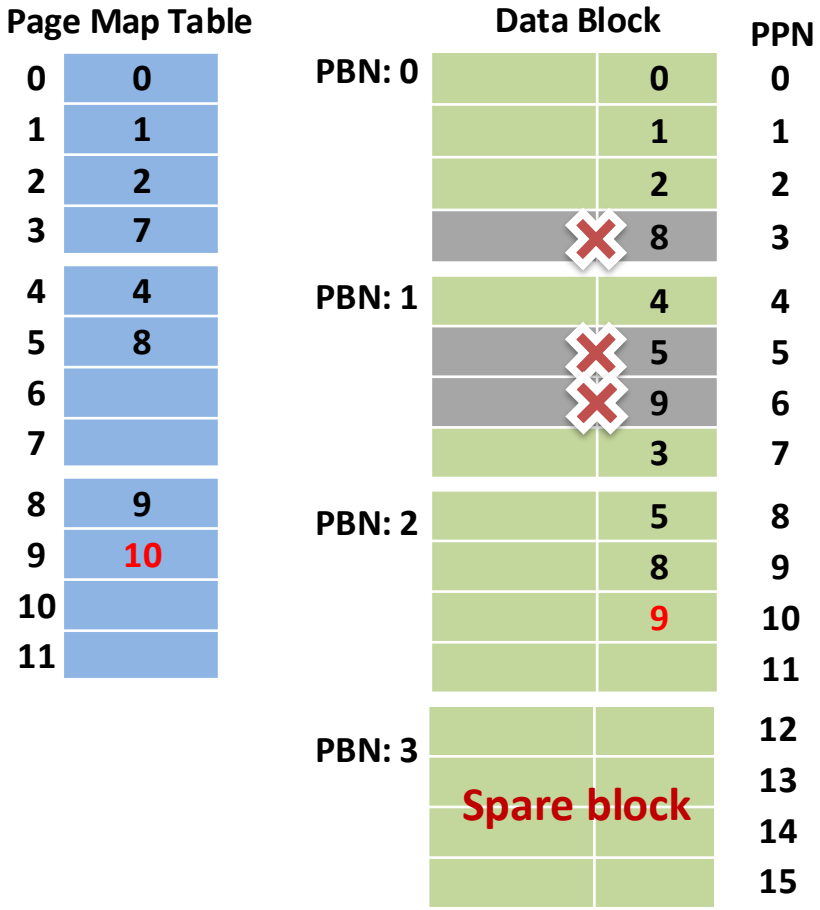

Example: GC in Page Mapping

- Current state
  - Written to page 0, 1, 2, 8, 4, 5
  - Written to page 9, 3, 5
- New requests (in order)
  - Write to page 8
  - Write to page 9
  - Write to page 3
  - Write to page 1
  - Write to page 4

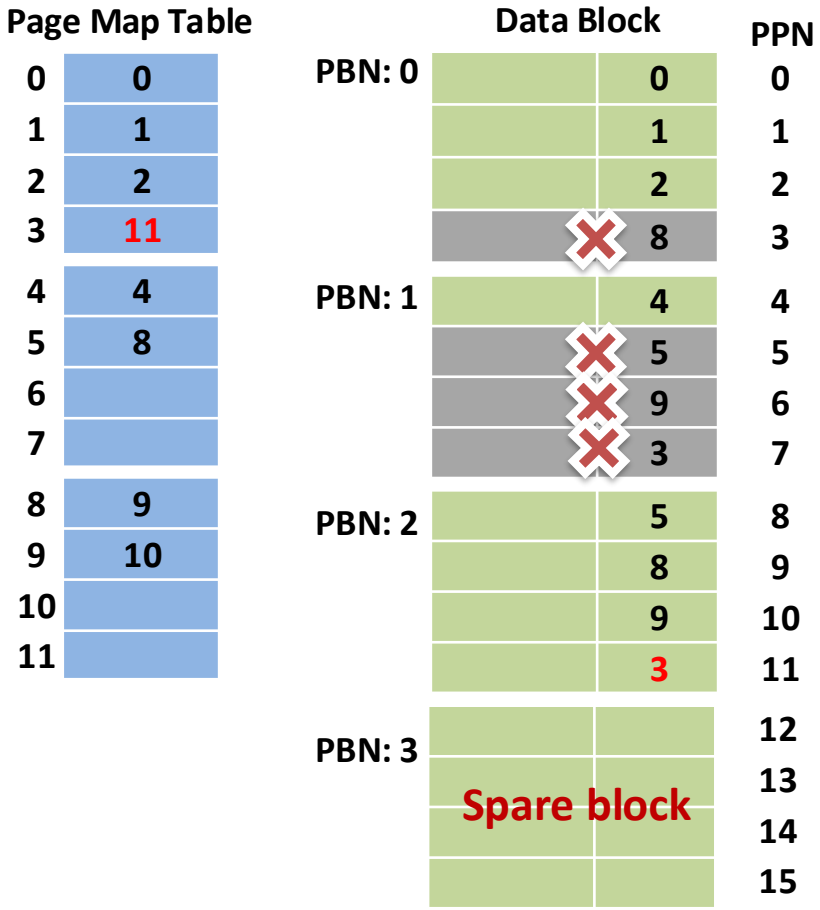

Example: GC in Page Mapping

- Current state
  - Written to page 0, 1, 2, 8, 4, 5
  - Written to page 9, 3, 5
- New requests (in order)
  - Write to page 8
  - Write to page 9
  - Write to page 3
  - Write to page 1
  - Write to page 4

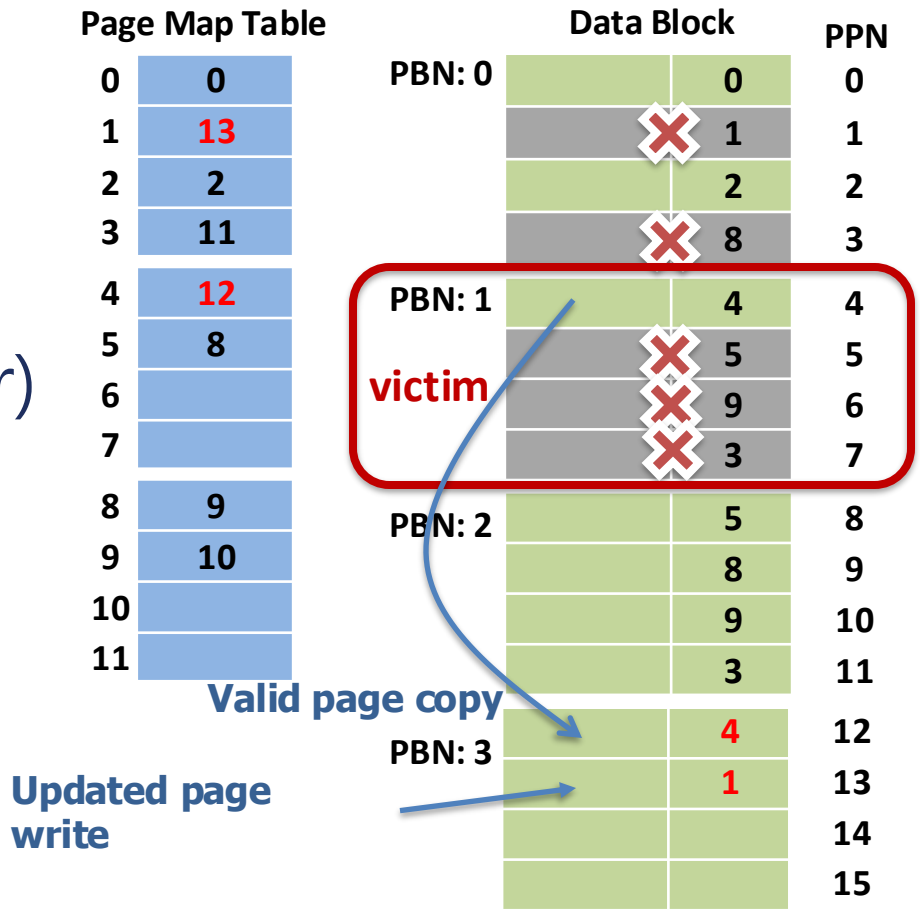

Example: GC in Page Mapping

- Current state
  - Written to page 0, 1, 2, 8, 4, 5
  - Written to page 9, 3, 5
- New requests (in order)
  - Write to page 8
  - Write to page 9
  - Write to page 3
  - Write to page 1
  - Write to page 4

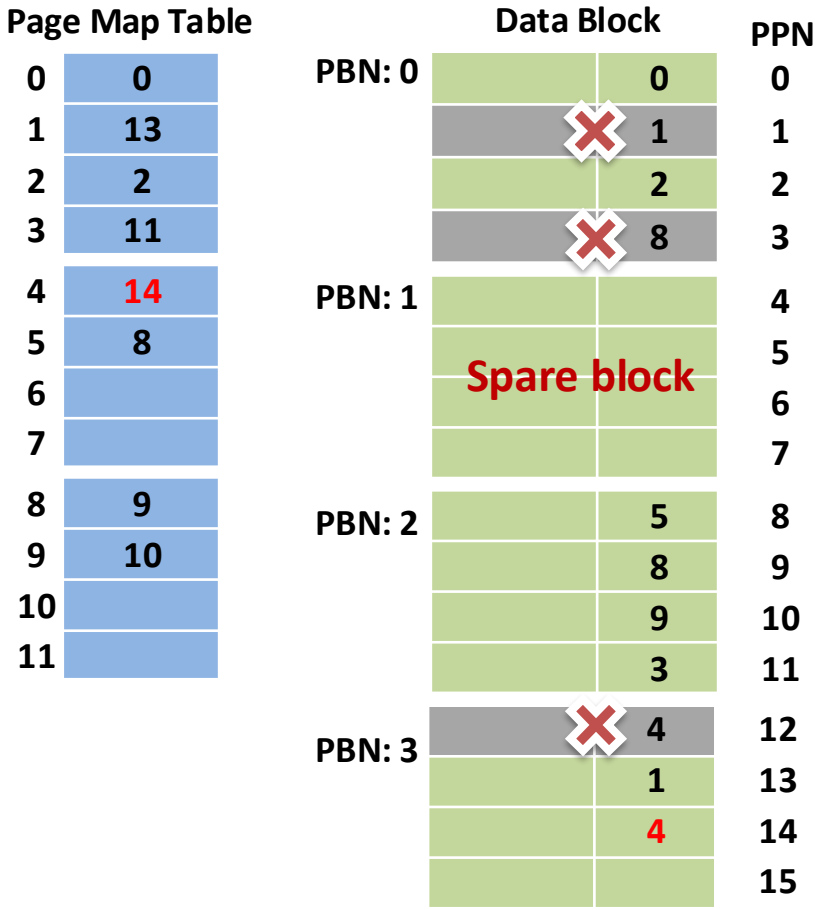

OS Implications

- NAND flash has different characteristics compared to disks
  - No seek time
  - Asymmetric read/write access times
  - No in-place-update
  - Good sequential read/write and random read performance, but bad random write performance
  - Wear-leveling
  - …
  - Traditional operating systems have been optimized for disks. What should be changed?

SSD Support in OS

- Turn off "defragmentation" for SSDs
- New "TRIM" command
  - Remove-on-delete
- Simpler I/O scheduler
- Align file system partition with SSD layout
- Flash-aware file systems (e.g. F2FS in Linux)
- Larger block size (4KB)

Beauty and the Beast

- NAND Flash memory is a beauty
  - Small, light-weight, robust, low-cost, low-power non-volatile device
- NAND Flash memory is a beast
  - Much slower program/erase operations
  - No in-place-update
  - Erase unit > write unit
  - Limited lifetime
  - Bit errors, bad blocks, …
- Software support is essential for performance and reliability!